

# To a Billion Messages and Beyond

---

János Csorvási

---

# János Csorvási

Infrastructure Engineer at Meltwater

Engineering Enablement Mission

# What is Engineering Enablement?

# What is Meltwater?



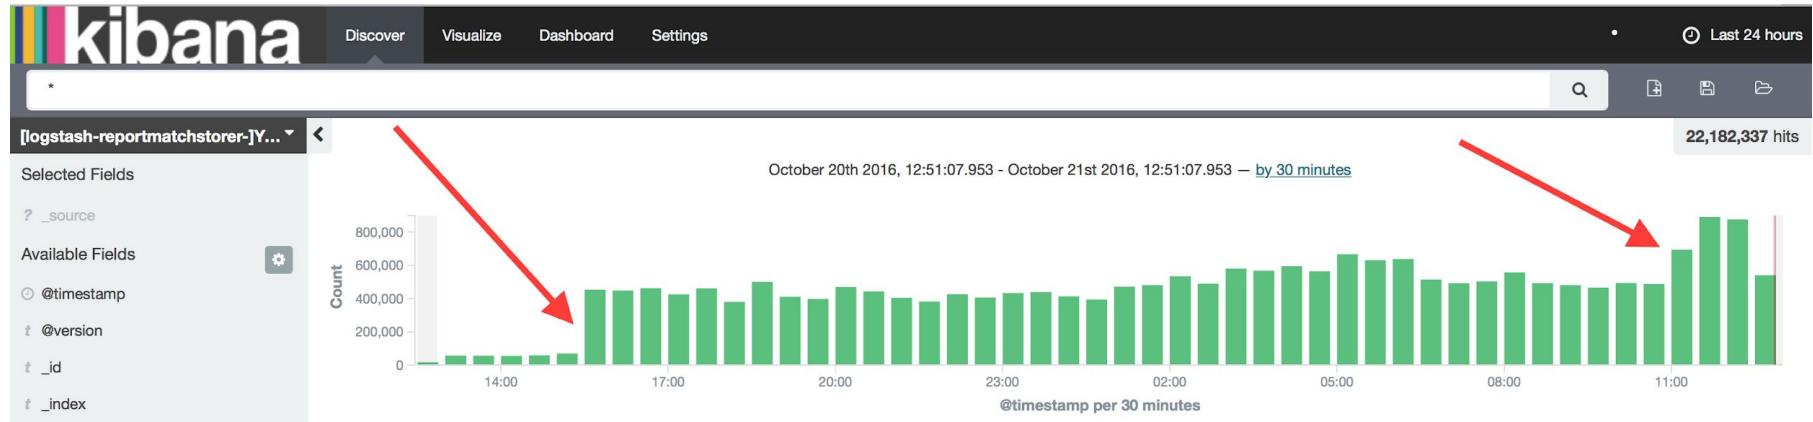
# What is this talk about?

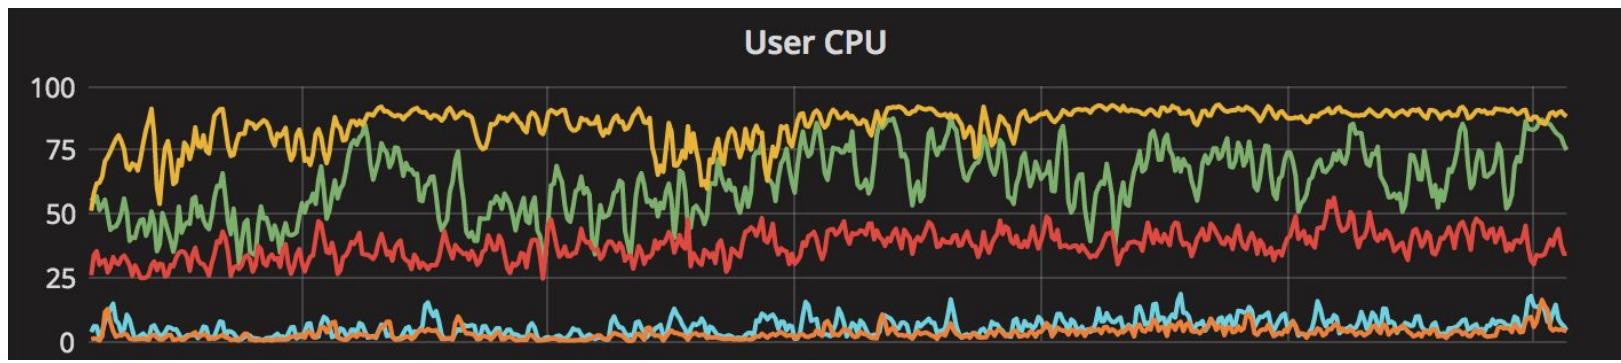




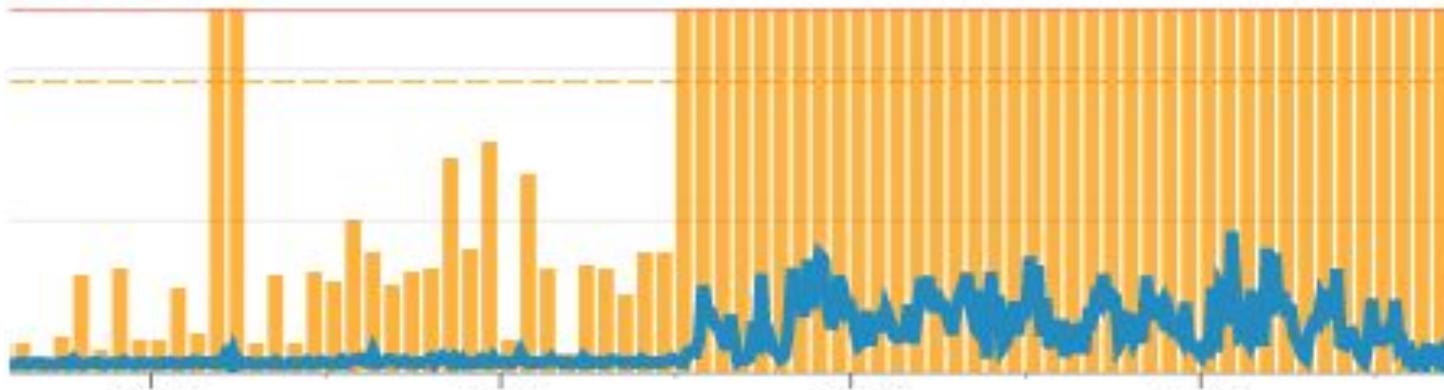






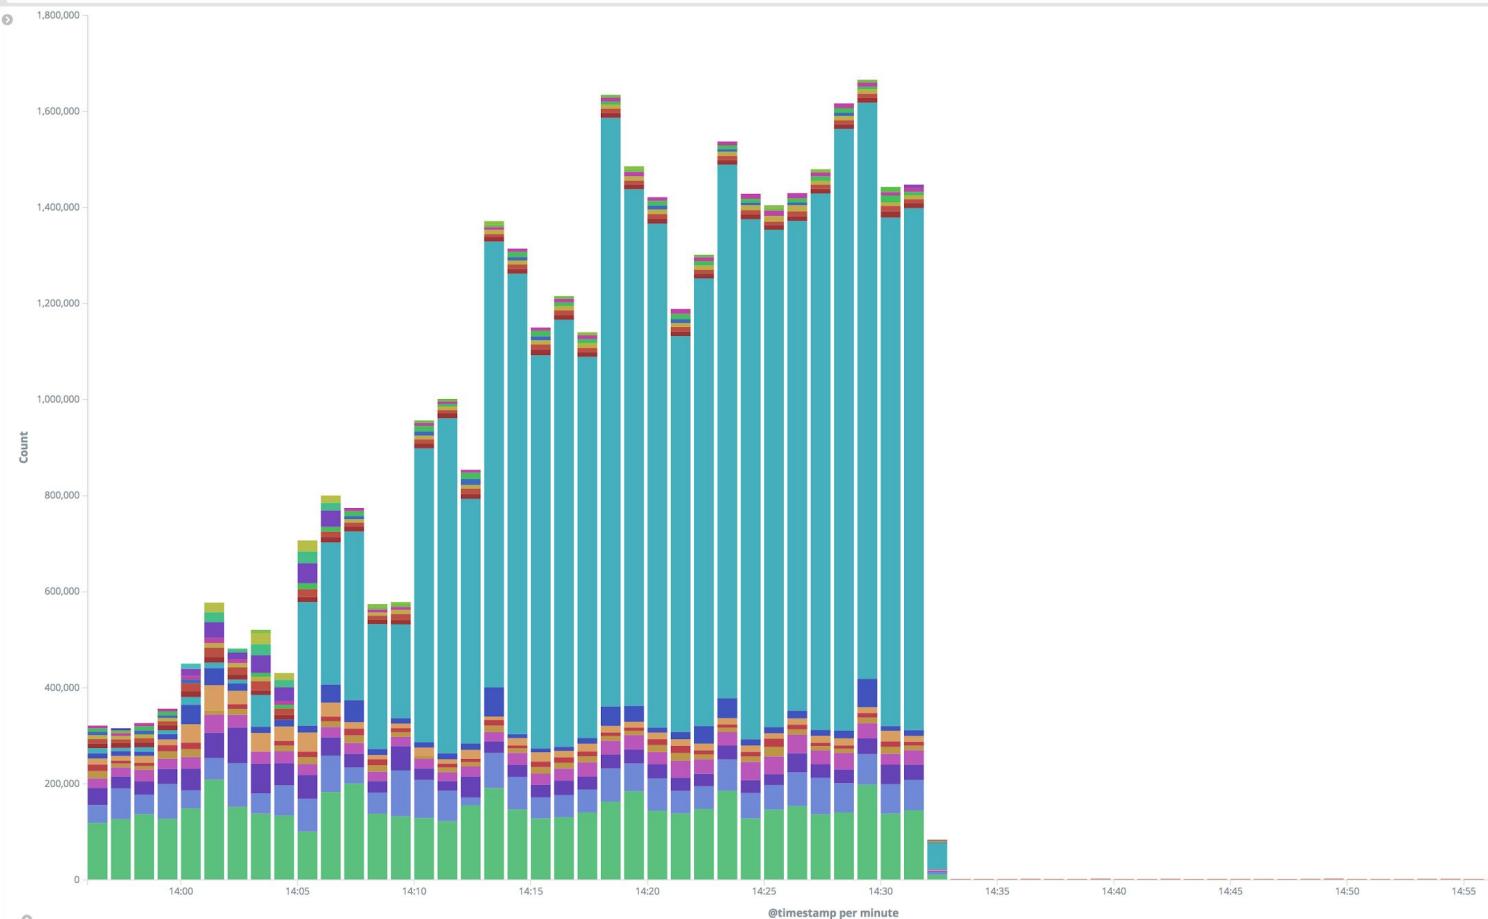


timeout



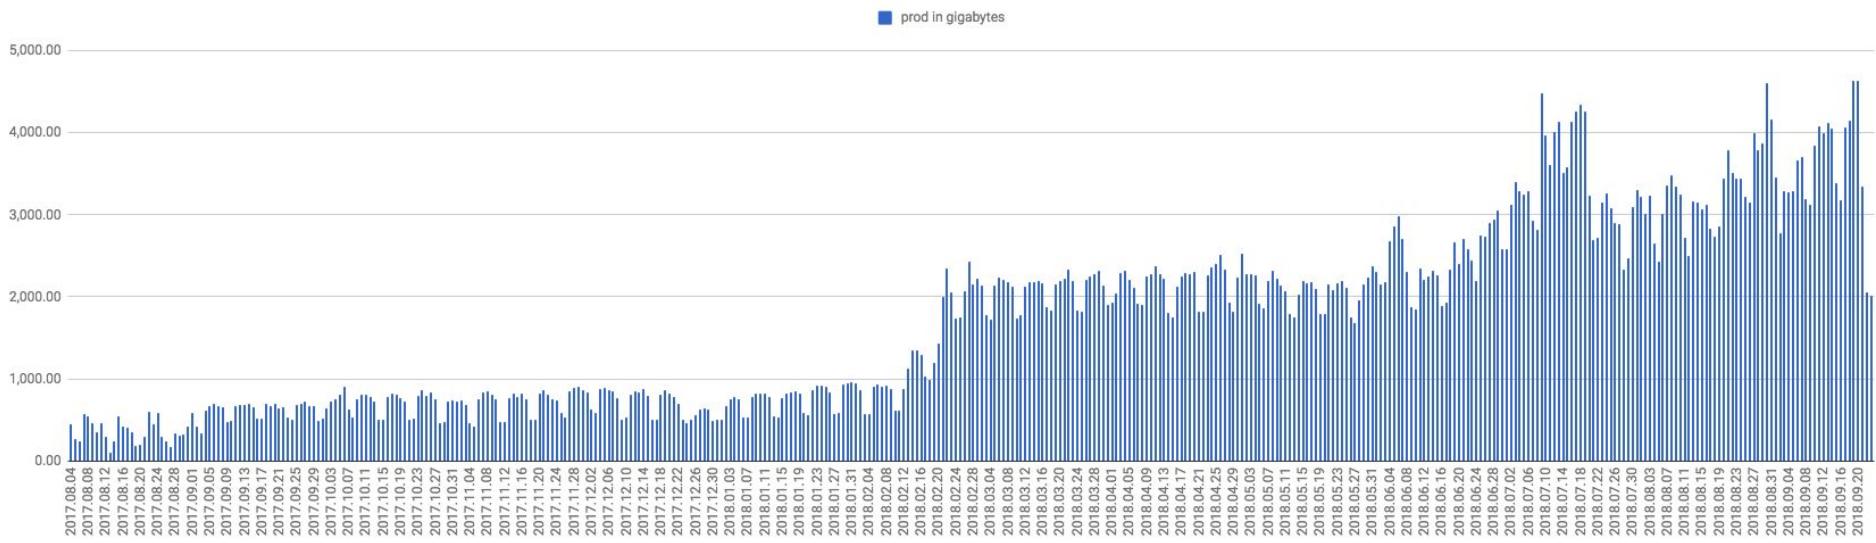
```
-      <root level="INFO">
+      <root level="TRACE">
          <appender-ref ref="LOGSTASH"/>
      </root>
```

```
for (Item item : items) {  
    log.trace("Starting to process item {}", item);  
    process(item);  
    log.trace("Done processing item {}", item);  
}
```









3 800 GB / day  
(with replicas)

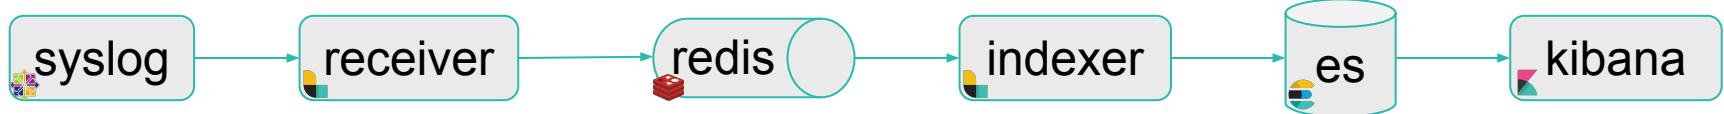
1 700 000 000 msg / day

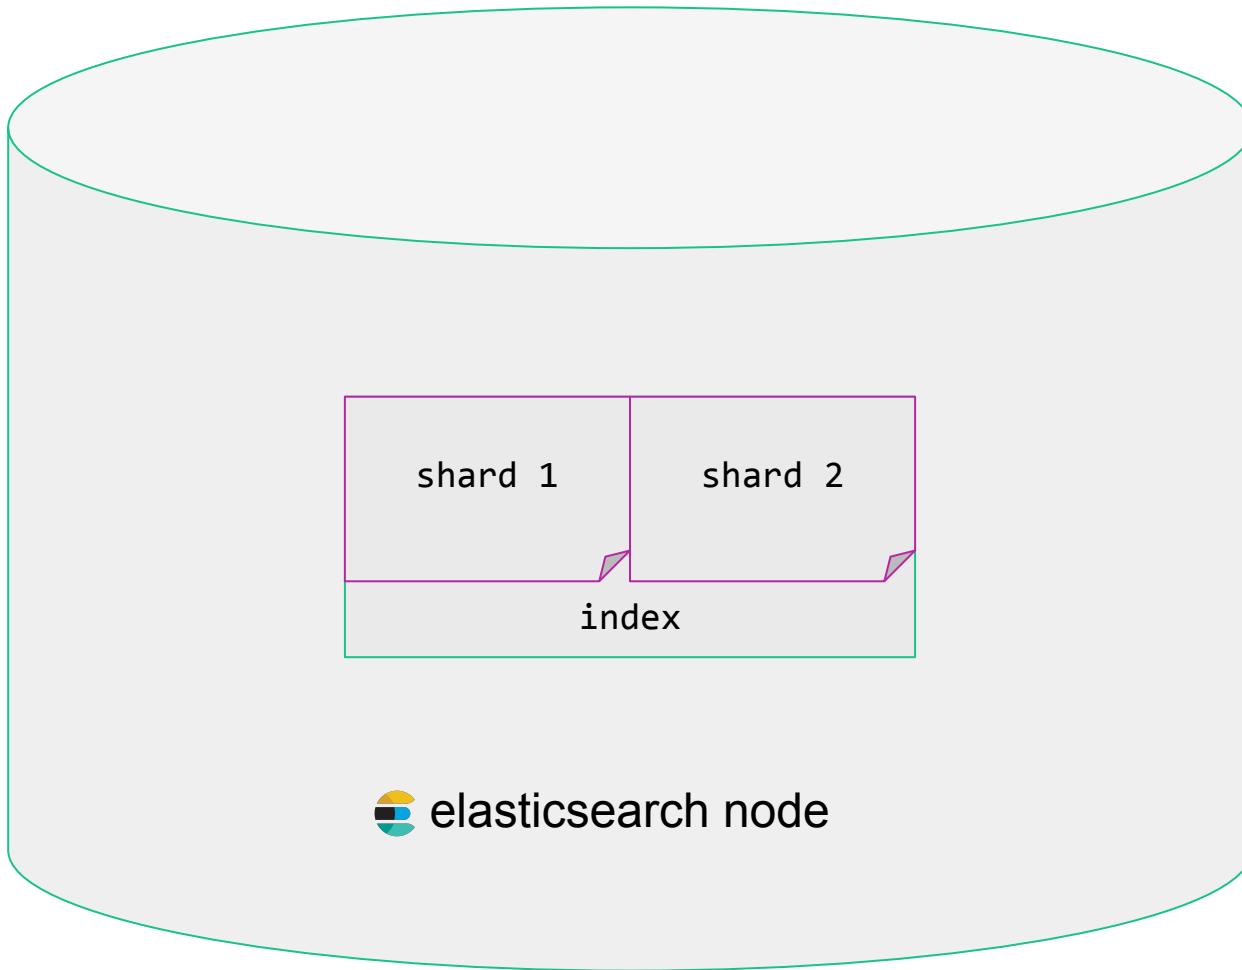
40 000 msg / sec  
(with replicas)

32 day rolling retention

# How did we evolve?

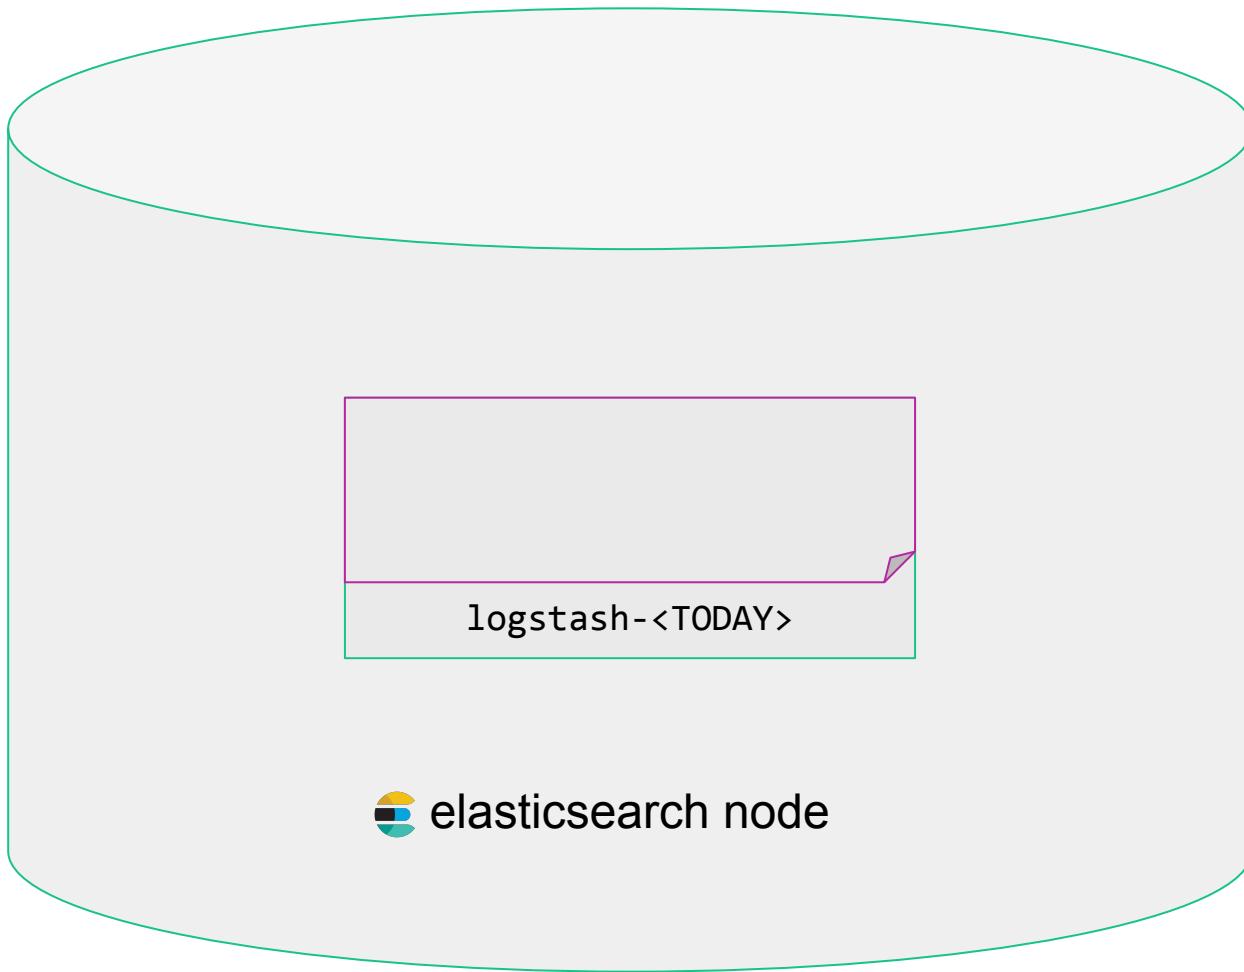
# Inherited system

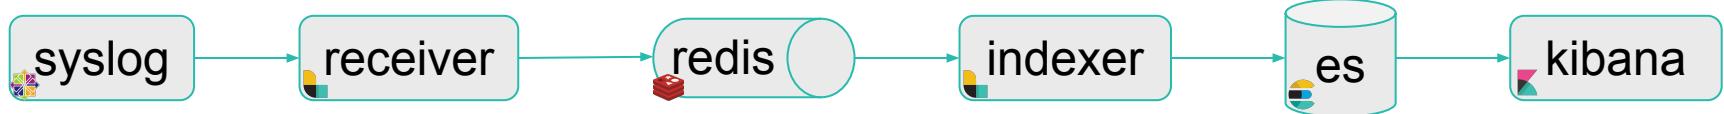


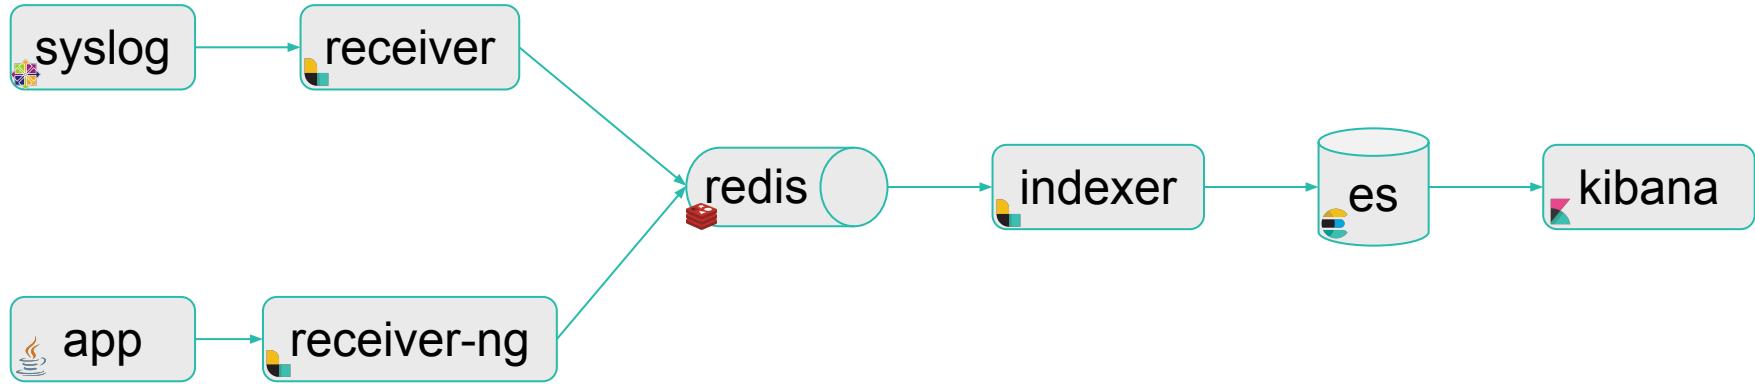




elasticsearch node







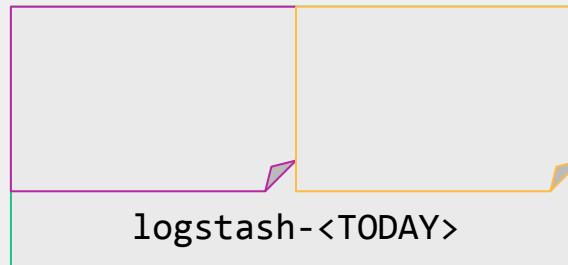


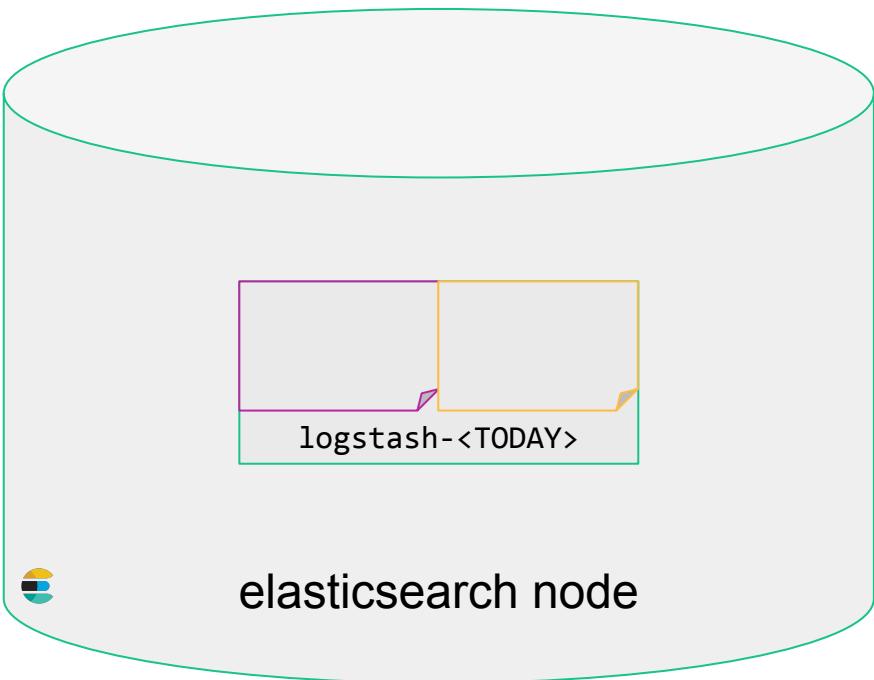
elasticsearch node



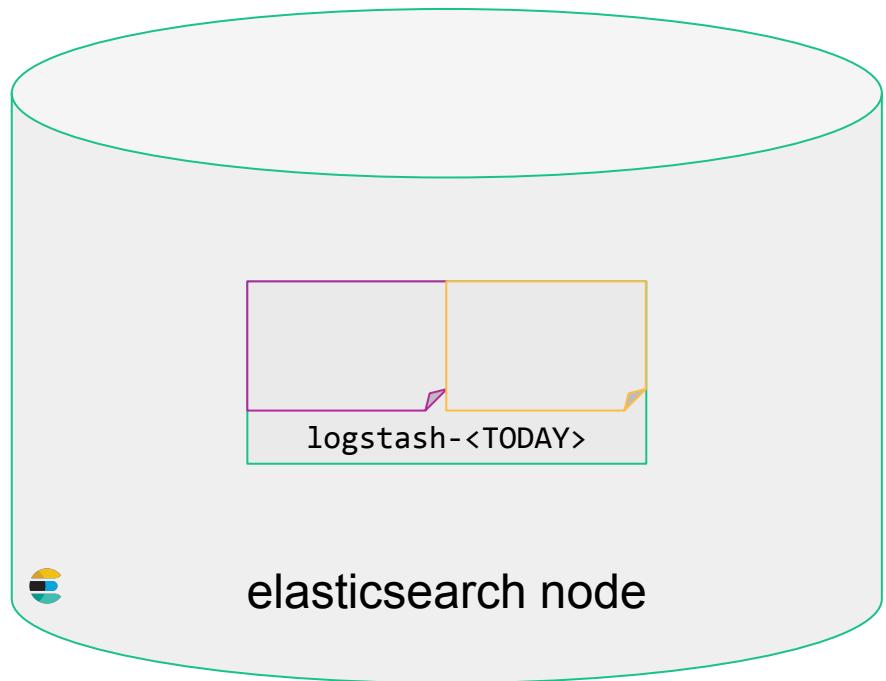


elasticsearch node

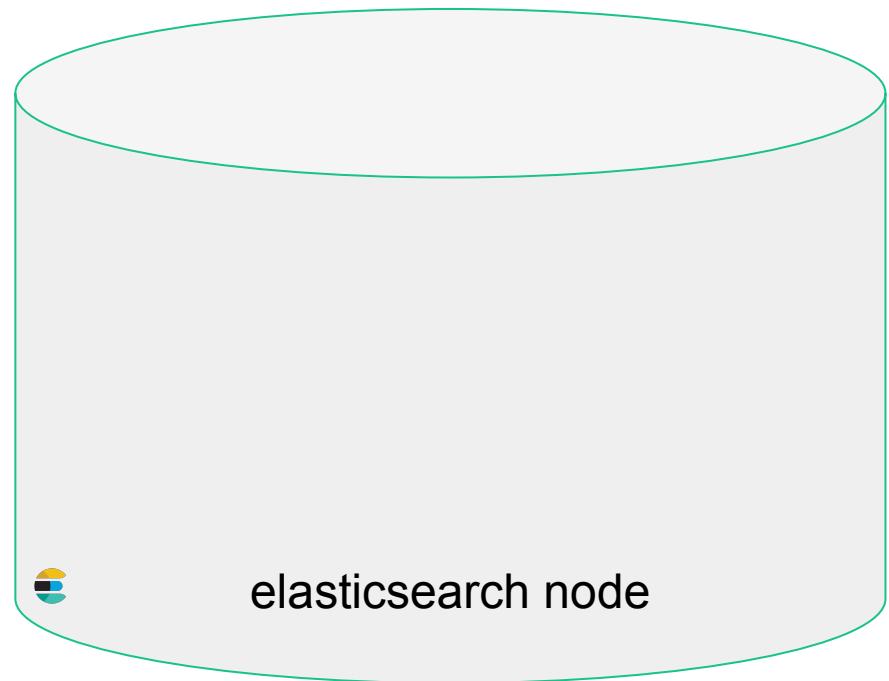




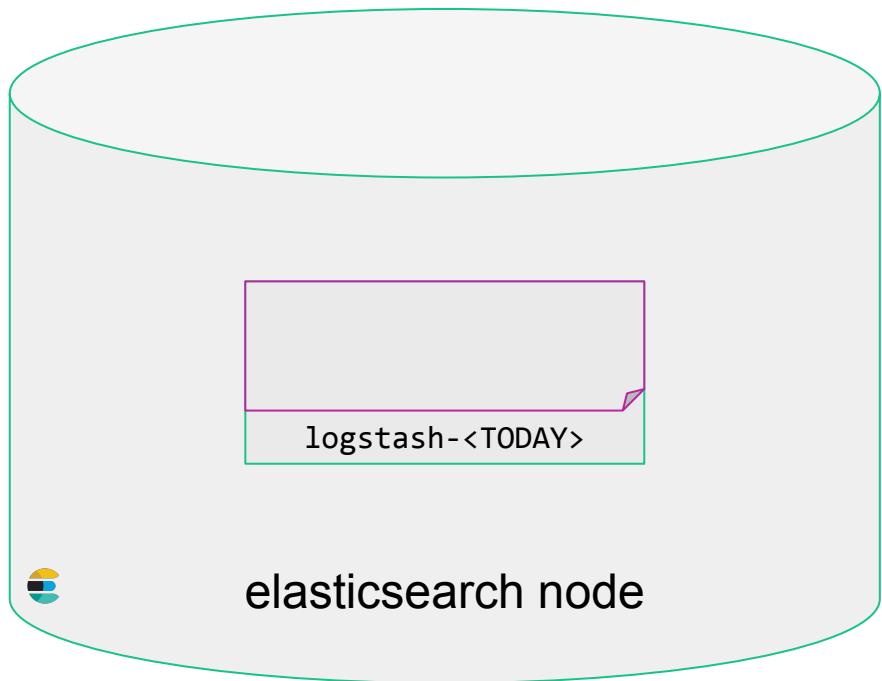
elasticsearch node



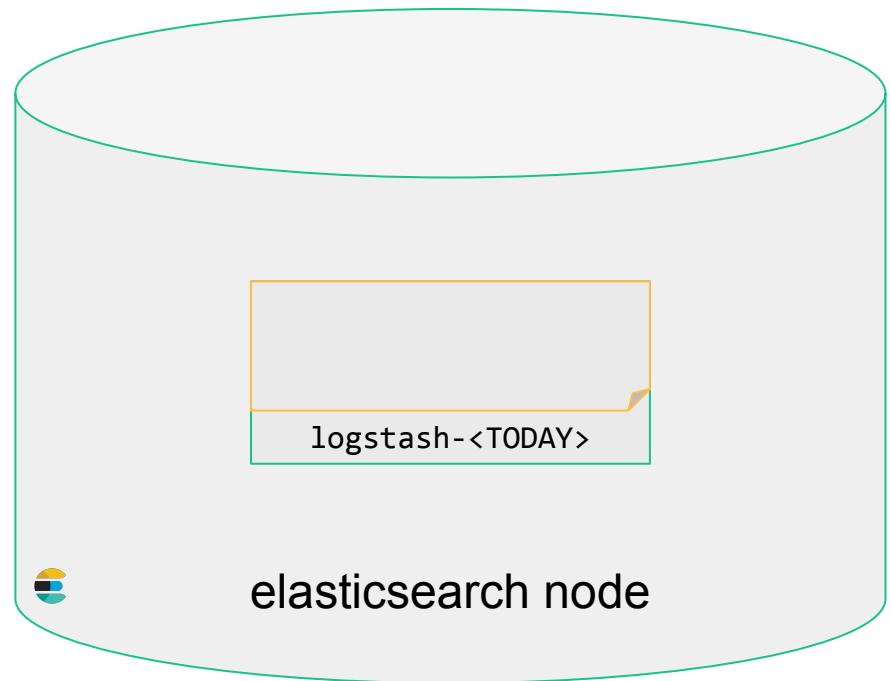
elasticsearch node



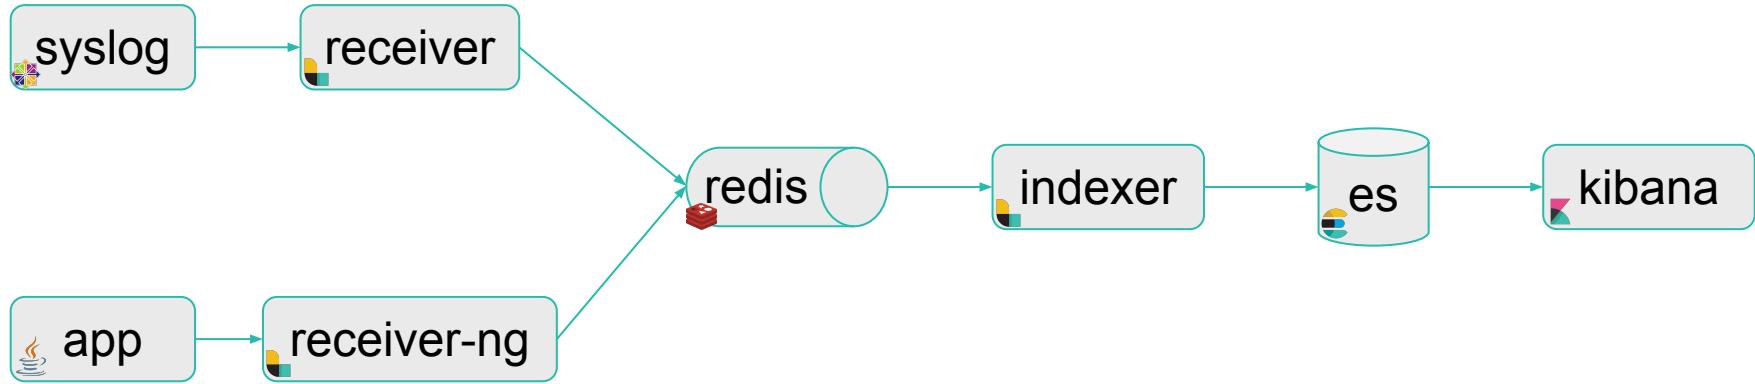
elasticsearch node

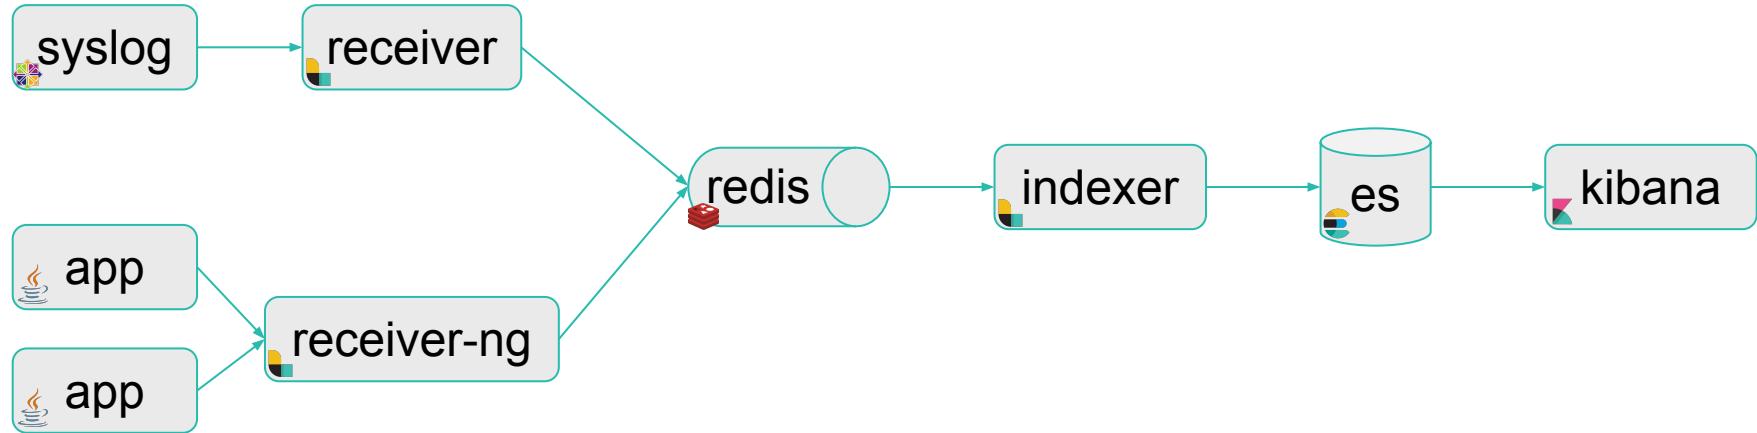


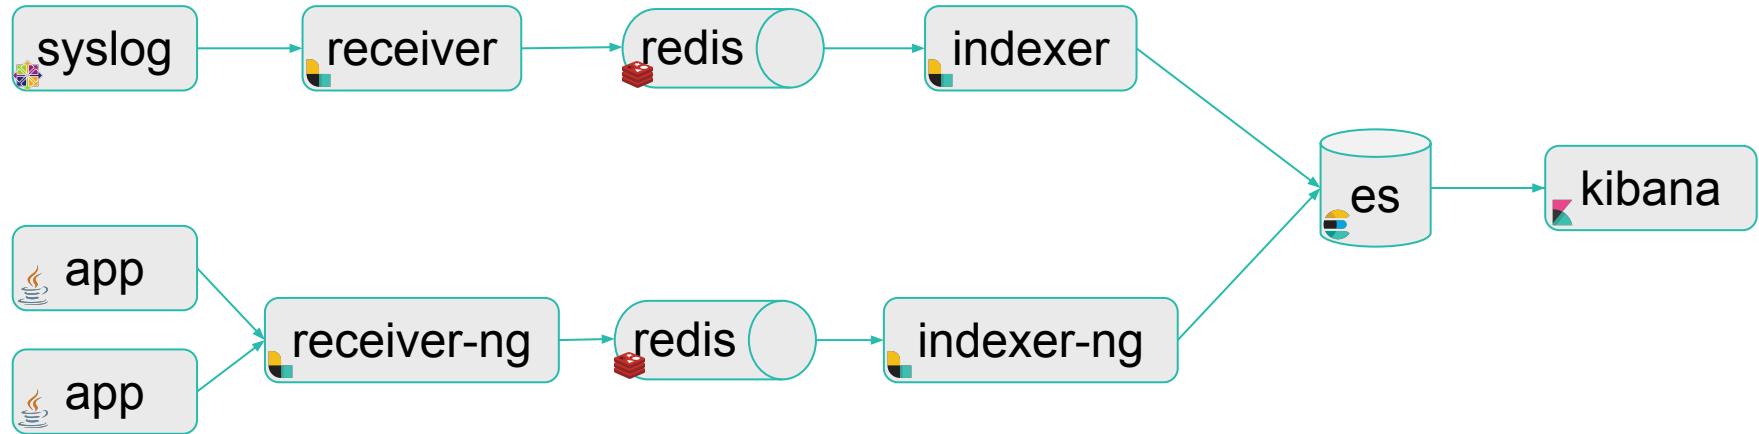
elasticsearch node

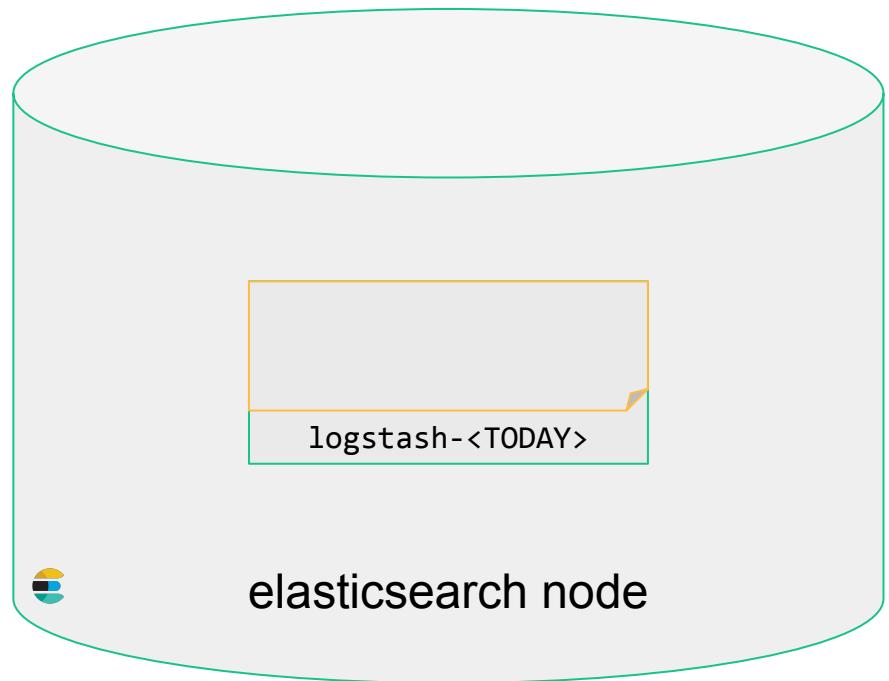
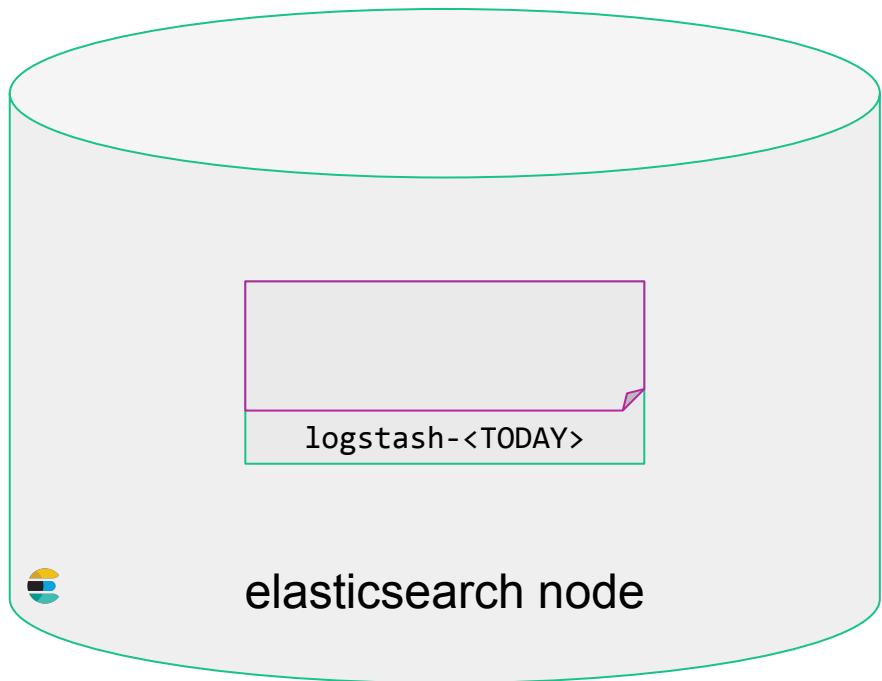


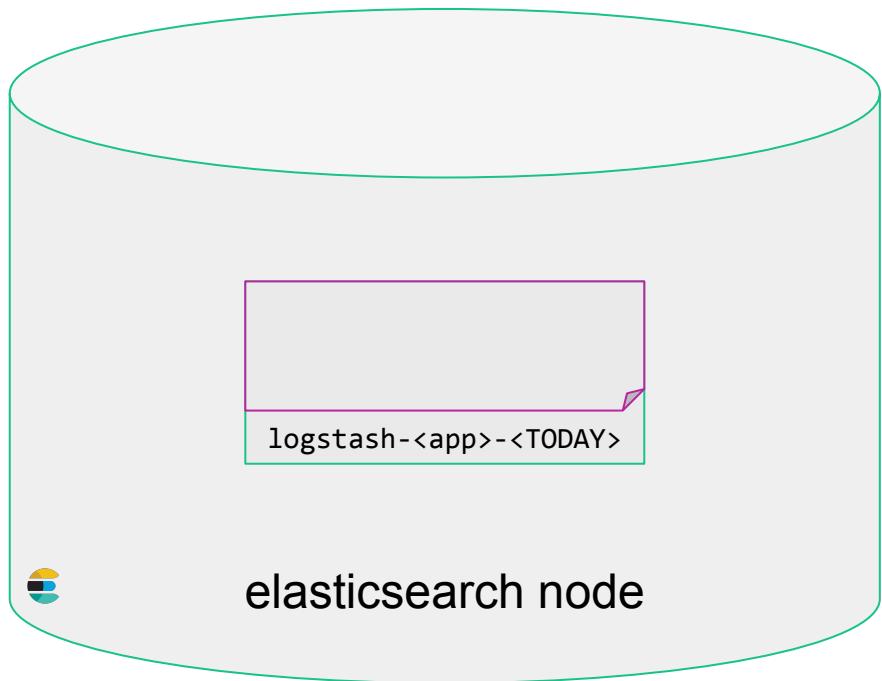
elasticsearch node



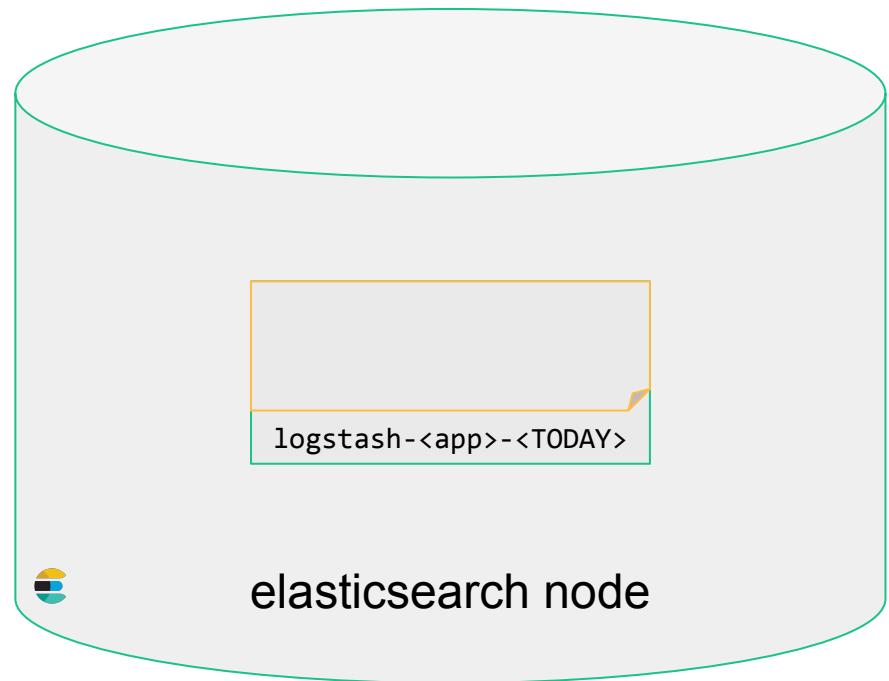




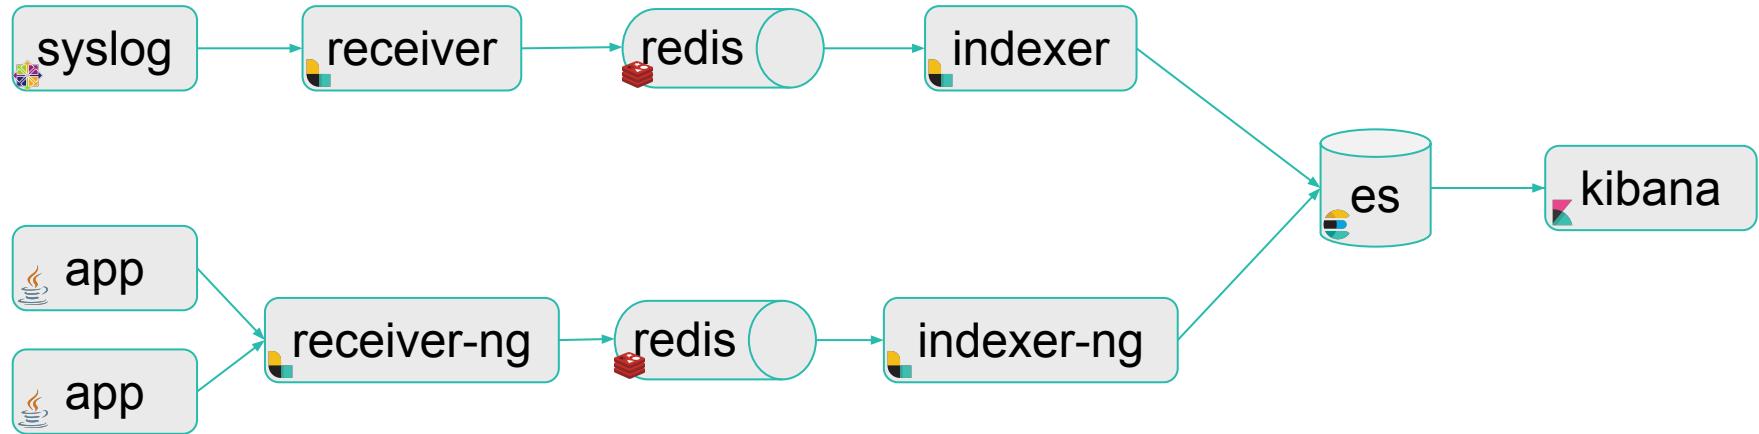


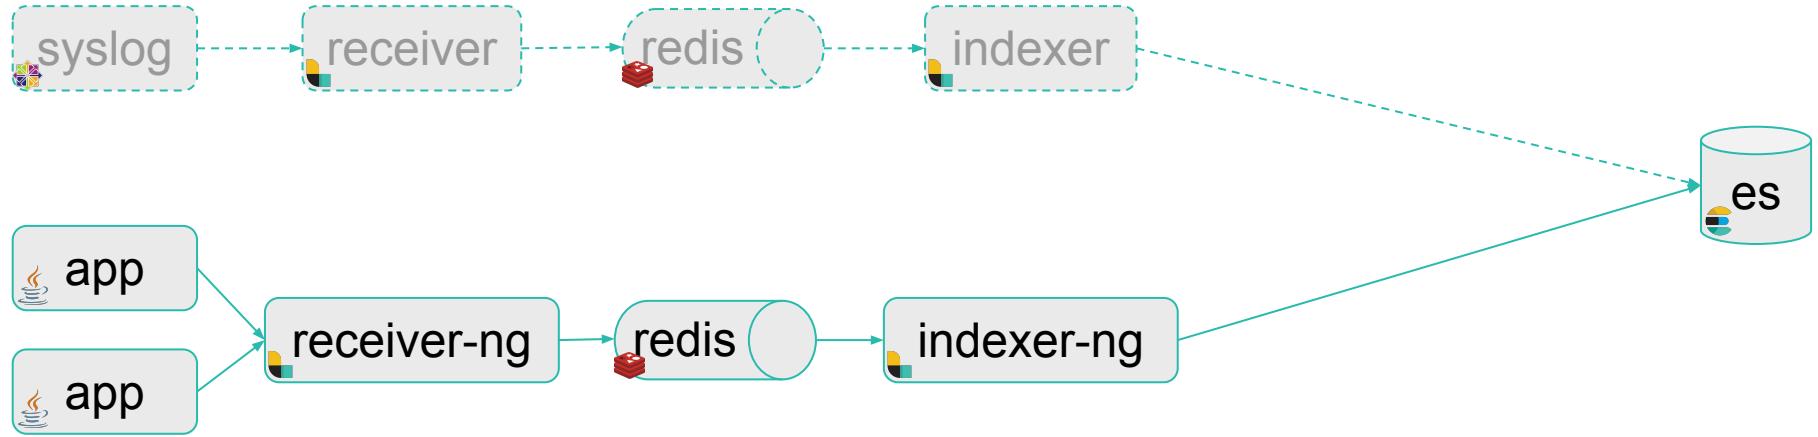


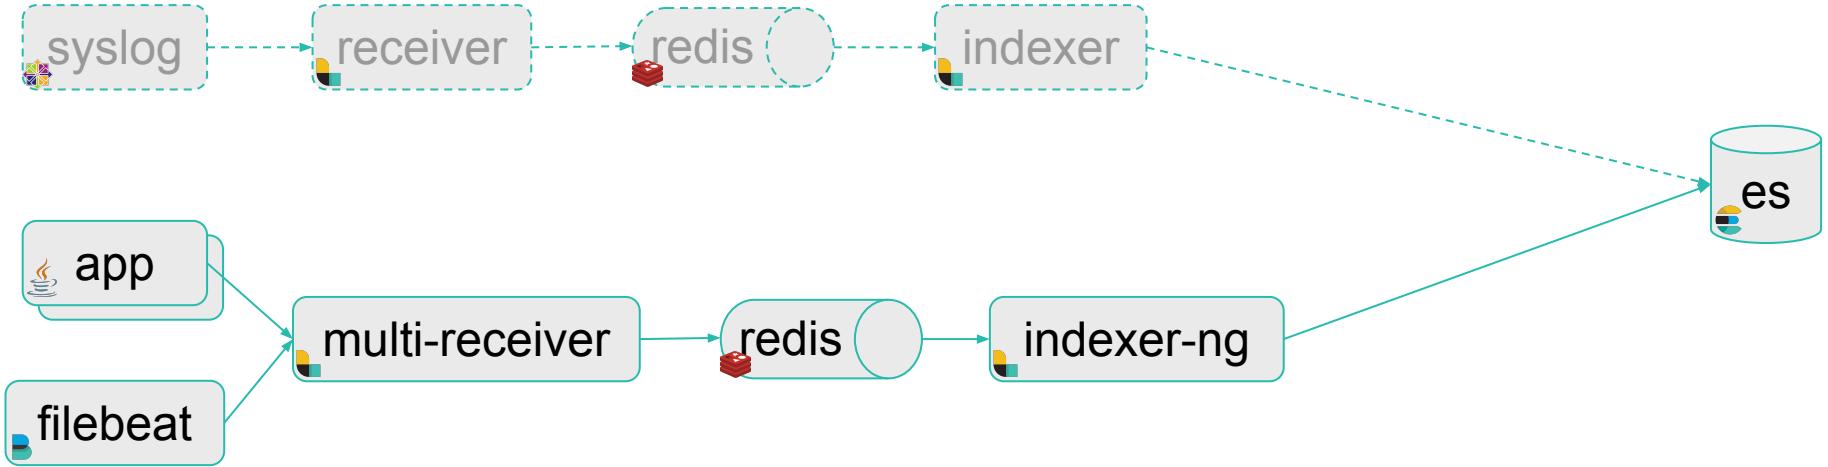
elasticsearch node

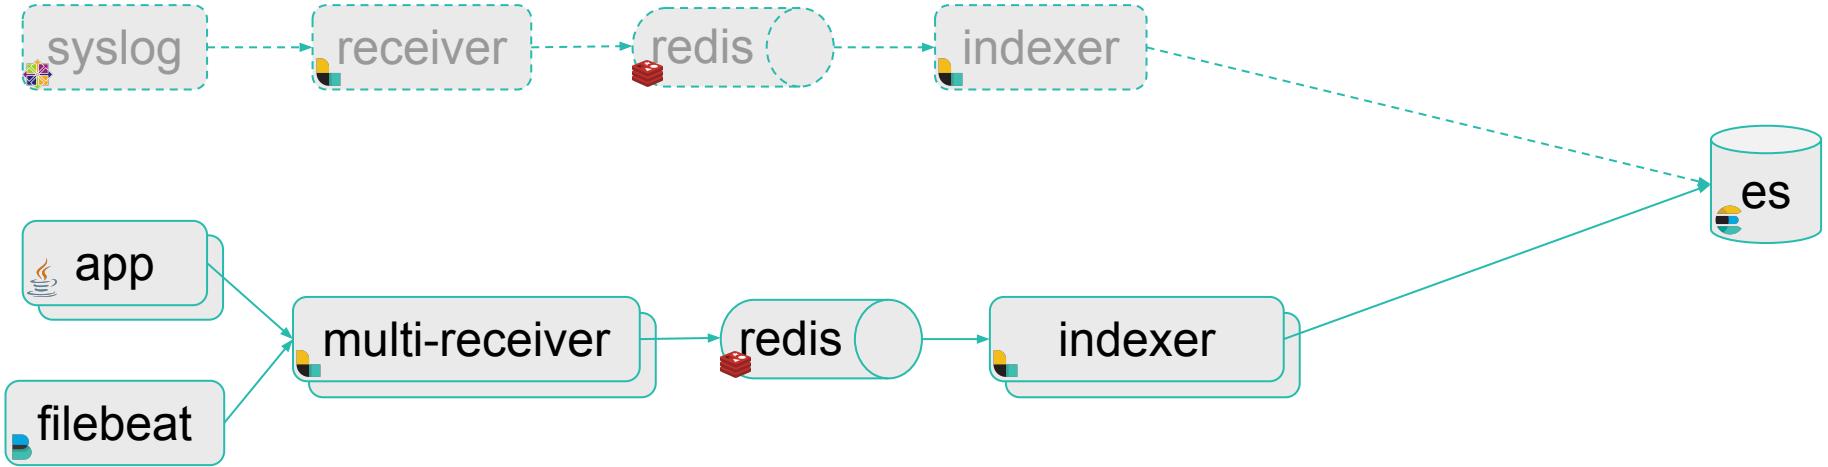


elasticsearch node



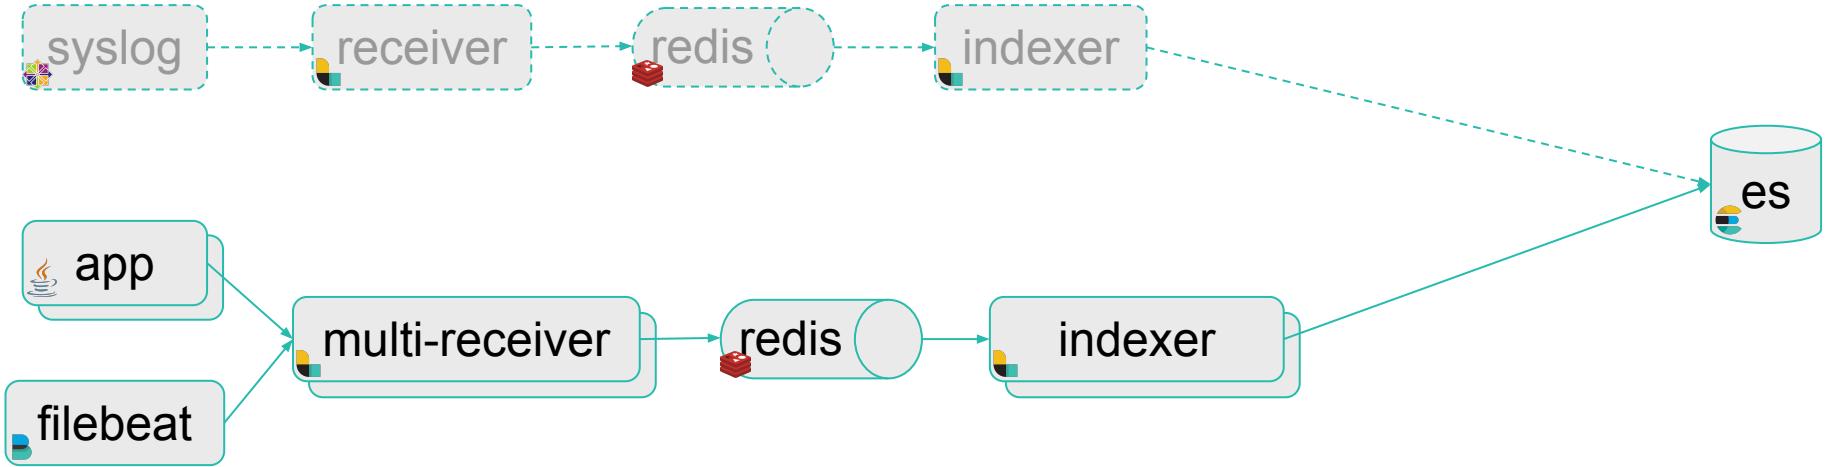


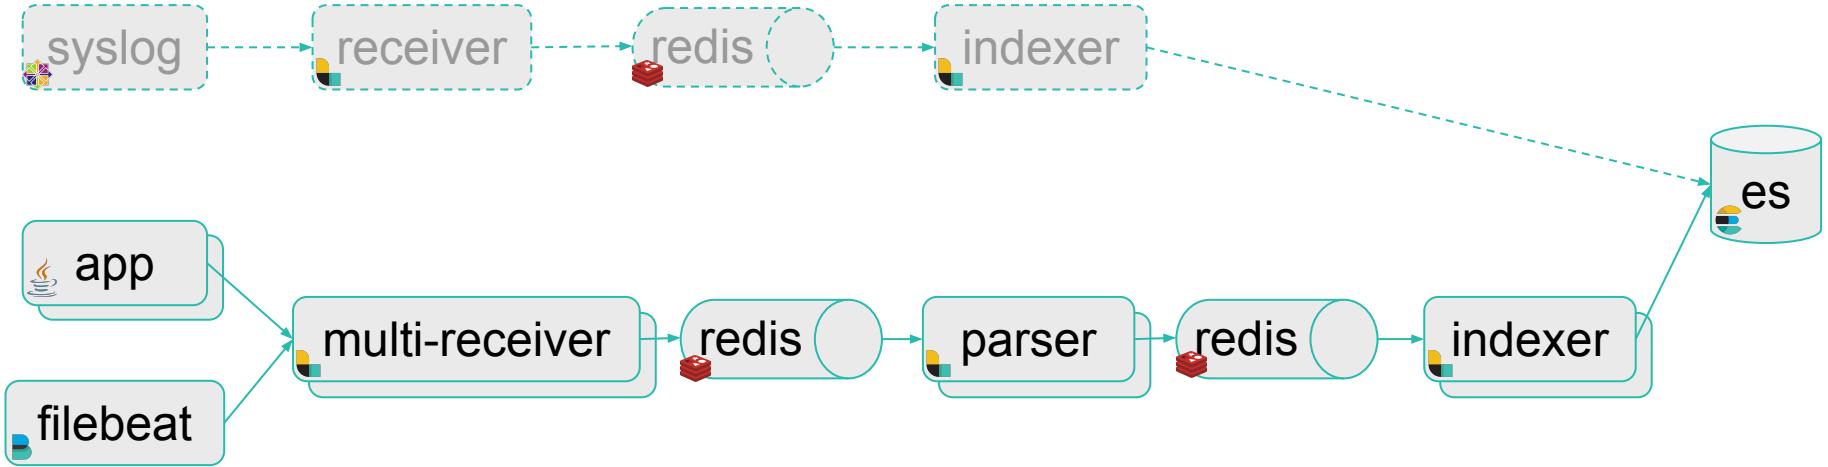




```
9.136.49.28 - - [27/Nov/2017:08:29:44 +0100] "GET /dashboard-legacy/v3/maps/dashboard-legacy.js.map HTTP/1.1"
200 1684812 "-" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_13_1) AppleWebKit/537.36 (KHTML, like Gecko)
Chrome/62.0.3202.94 Safari/537.36"
```

```
{
  "geoip": {
    "timezone": "Europe/Stockholm",
    "ip": "79.136.49.28",
    "latitude": 57.7167,
    "continent_code": "EU",
    "city_name": "Gothenburg",
    "country_name": "Sweden",
    "country_code2": "SE",
    "country_code3": "SE",
    "region_name": "Västra Götaland",
    "location": {
      "lon": 11.9667,
      "lat": 57.7167
    },
    "postal_code": "400 16",
    "region_code": "O",
    "longitude": 11.9667
  }
}
```





# Bug

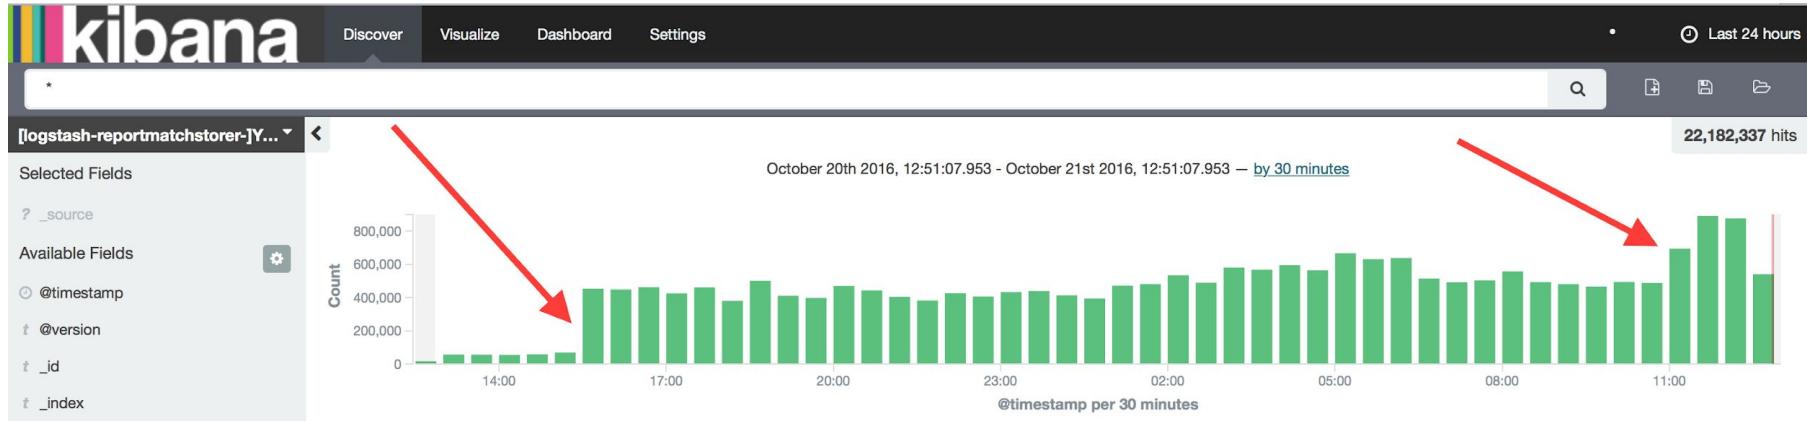


json\_lines codec

```
{"message": "Stuff happened"}\n{"message": "Stuff happened"}\n{"message": "Stuff happened"}\n{"message": "Stuff happened"}\nEOF
```

json codec

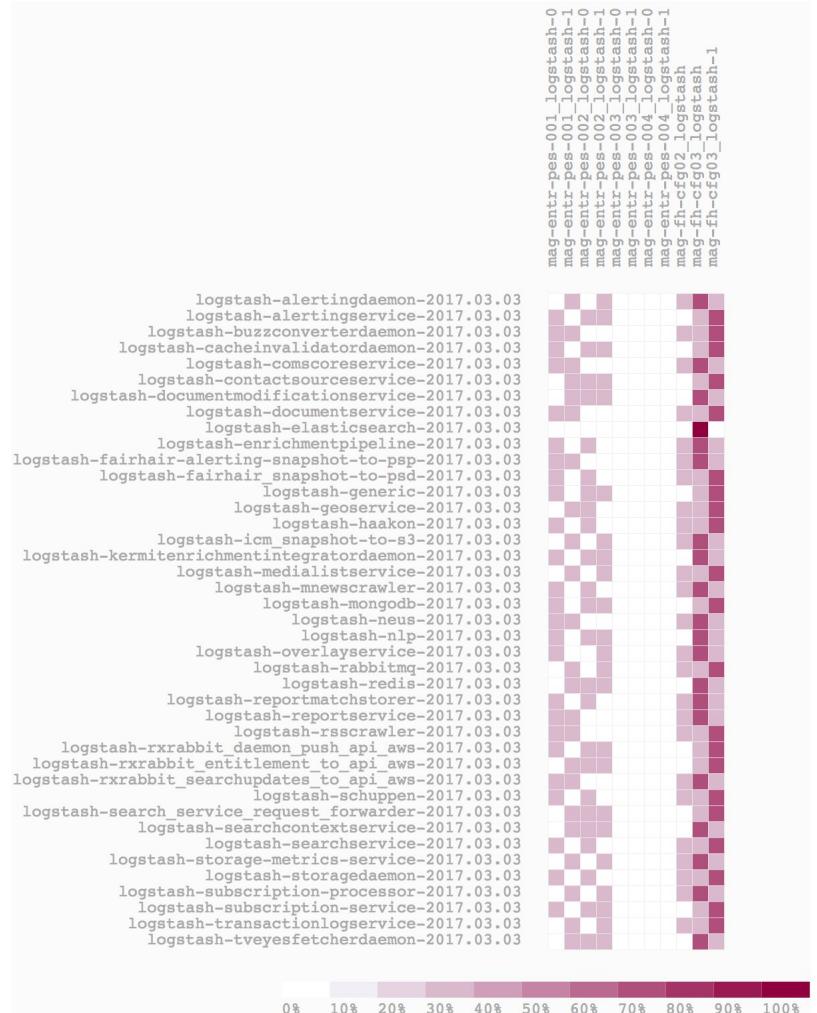
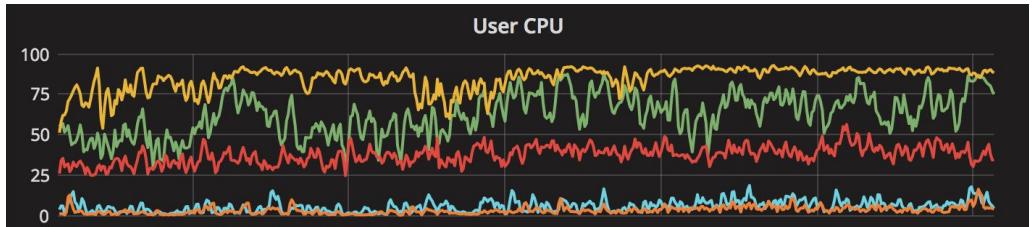
```
{"message": "Stuff happened"}\nEOF
```



# Challenge #1

# Challenge

- Massive elasticsearch scaling issues
- Capacity problems (no more)
- Uneven cluster
  - Running on spare hw
  - Allocation catch-22



# Challenge

- Massive elasticsearch scaling issues
- Capacity problems (no more)
- Uneven cluster
  - Running on spare hw
  - Allocation catch-22

# Solution

- Build new cluster (migrate to AWS)
  - Use opportunity to upgrade stack

Easy-peasy!

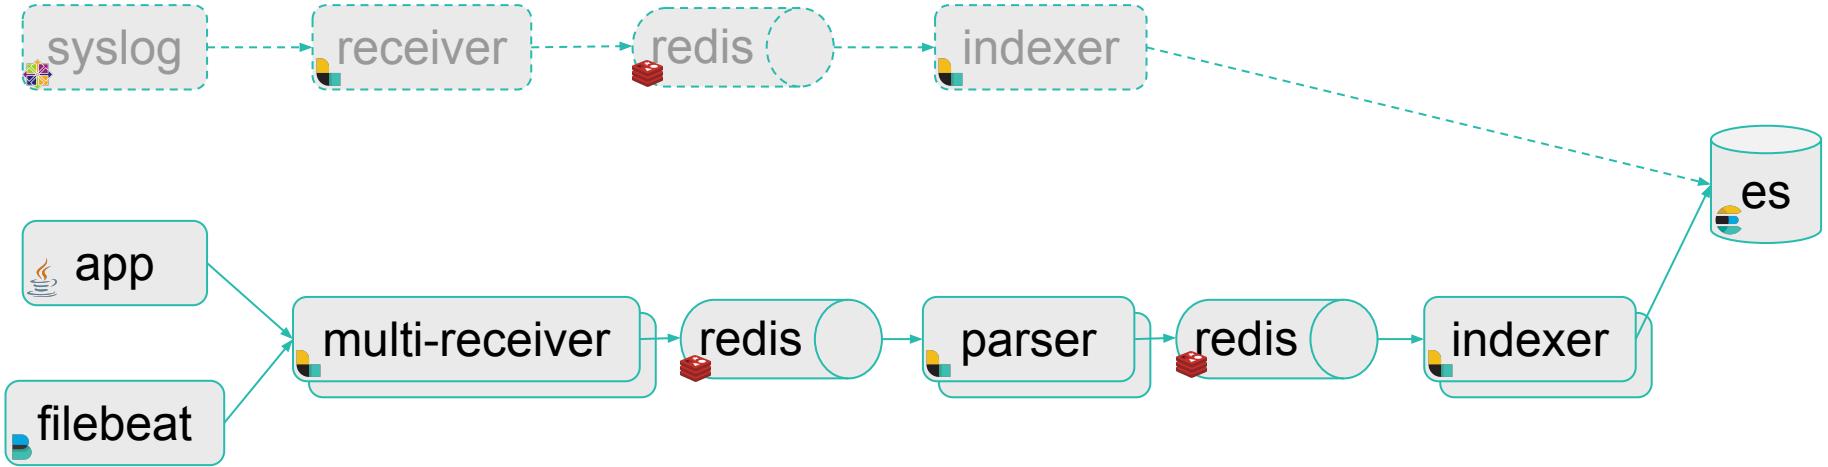
# Challenge #2

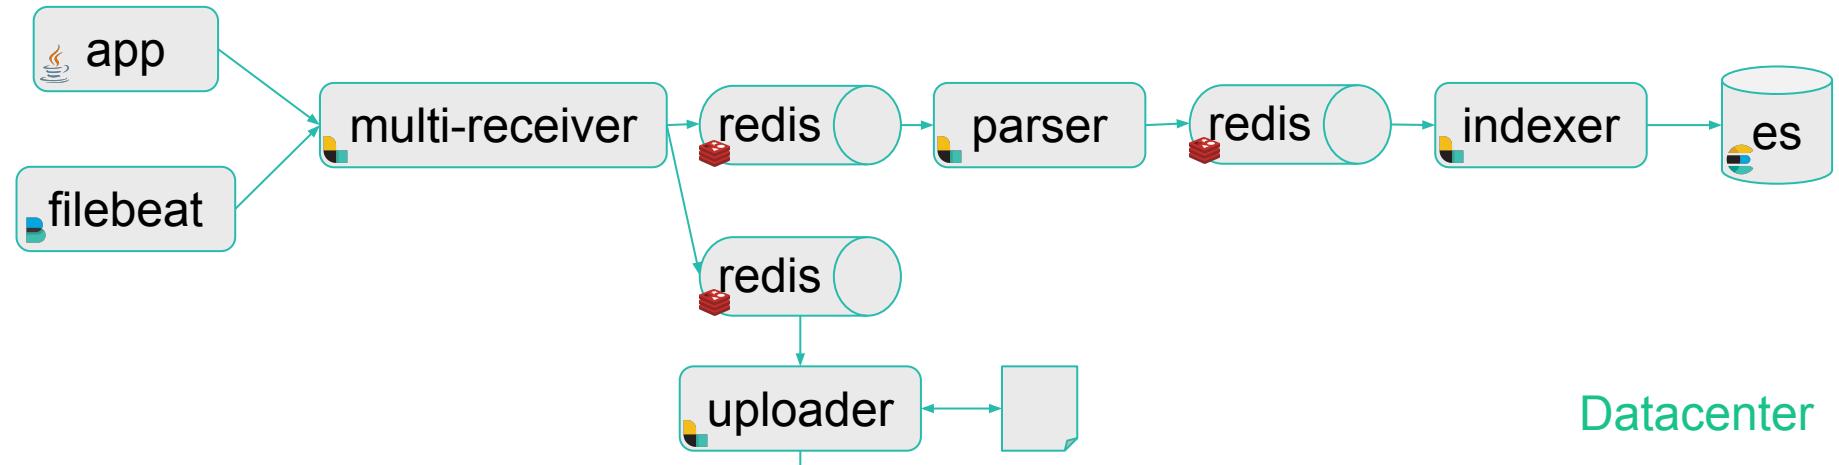
# Challenge

- Migrate to AWS...

# Solution

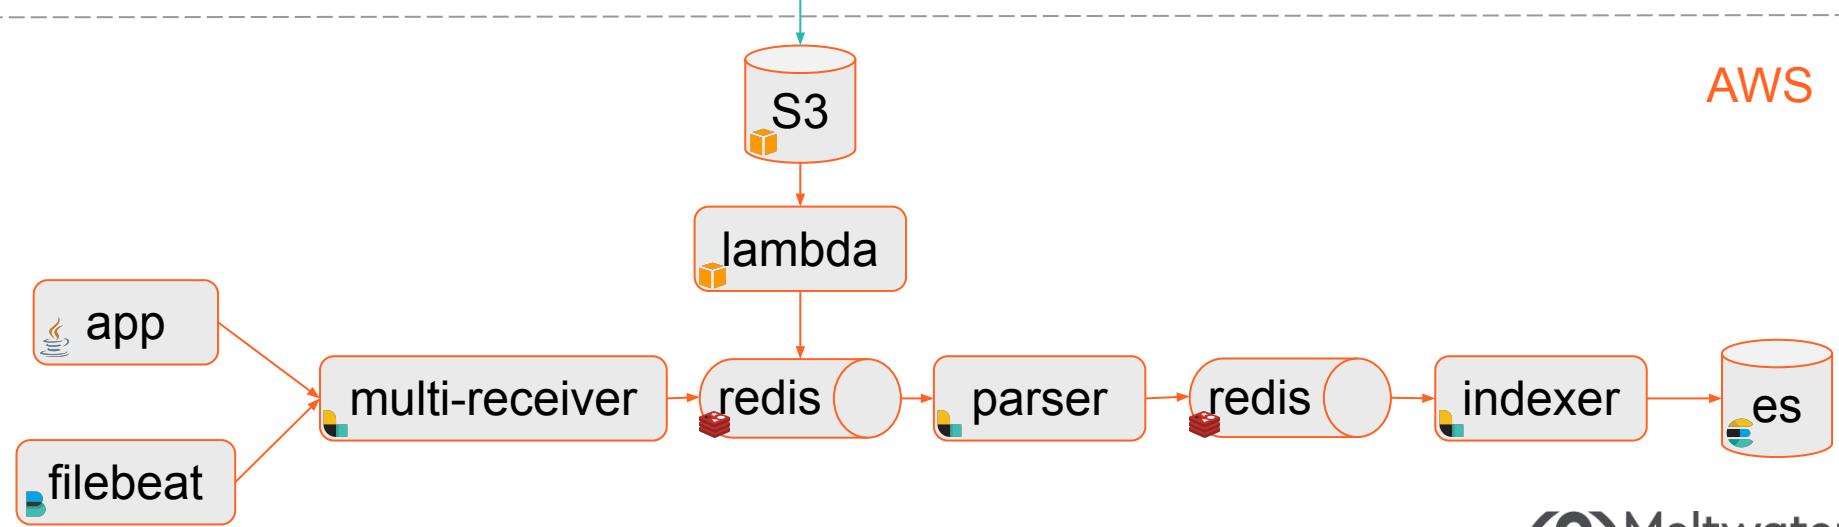
- Build new mesos cluster
- Build new and shiny elasticsearch cluster
- Send data to both places





Datacenter

AWS



s3-to-redis lambda invocations



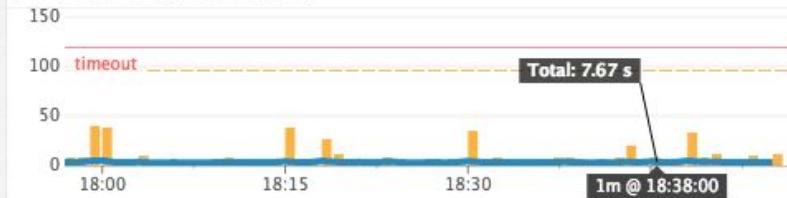
s3-to-redis log lines processed

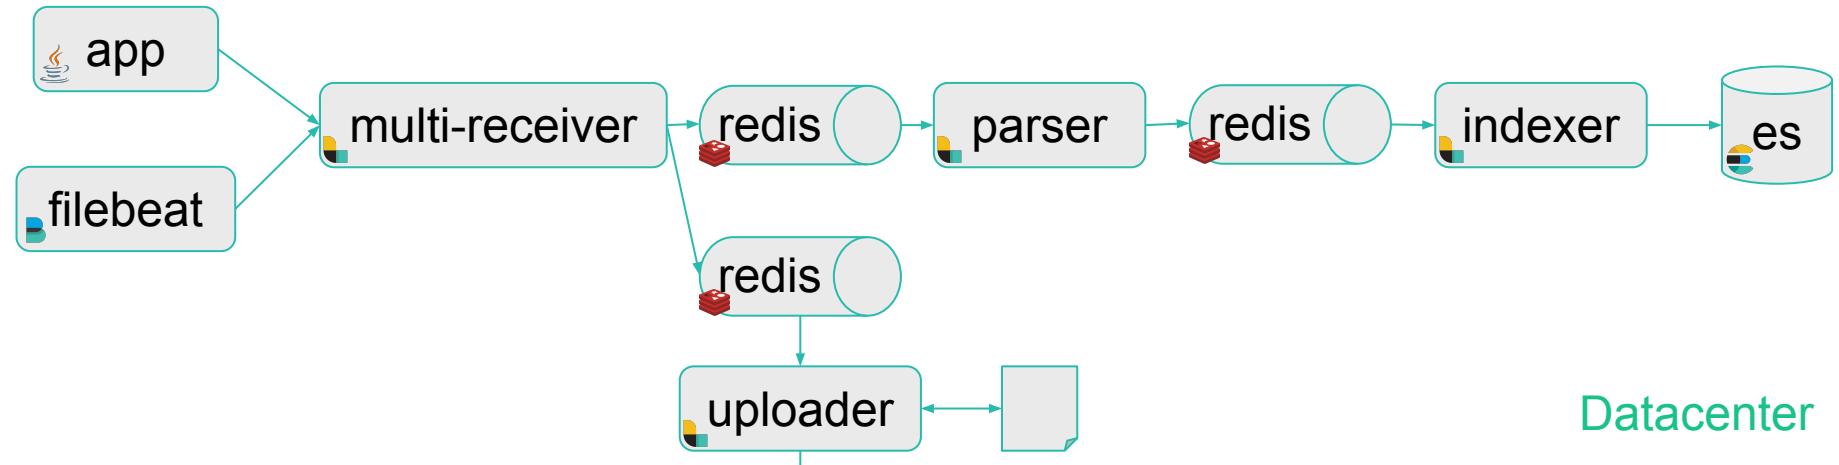


s3-to-redis lambda problems



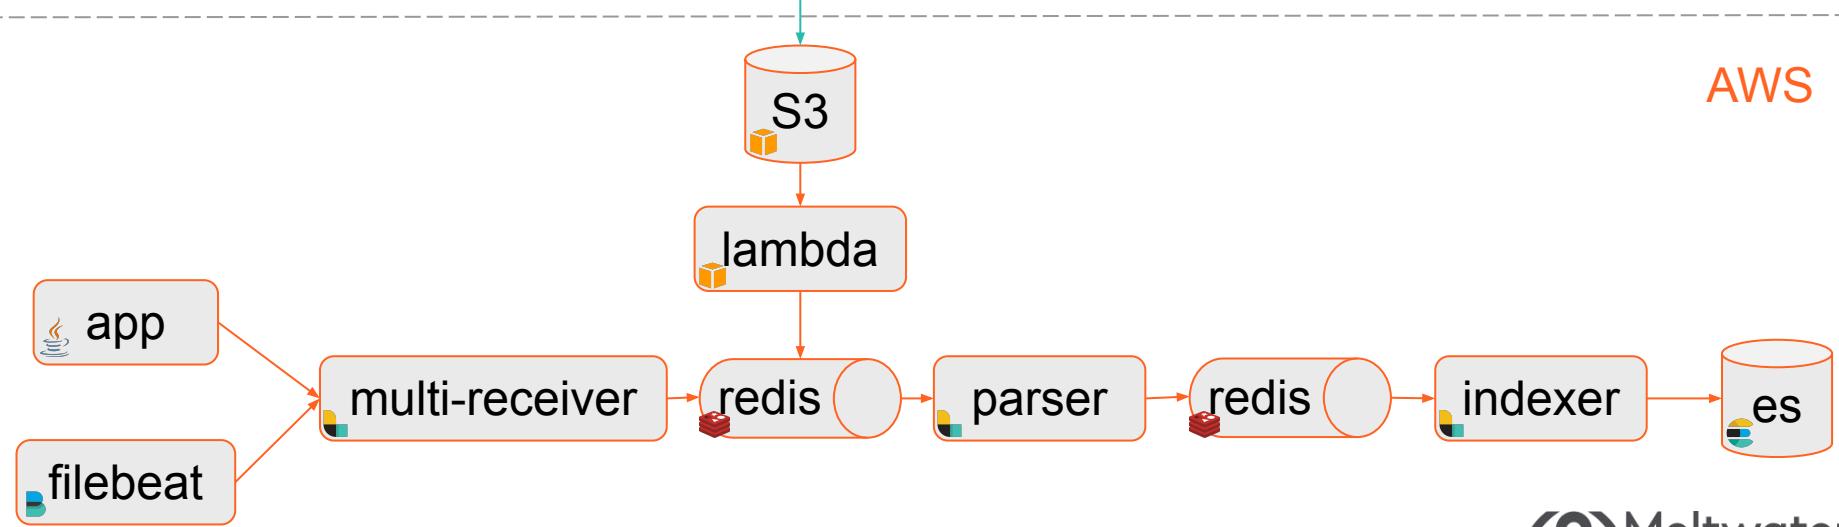
s3-to-redis lambda duration

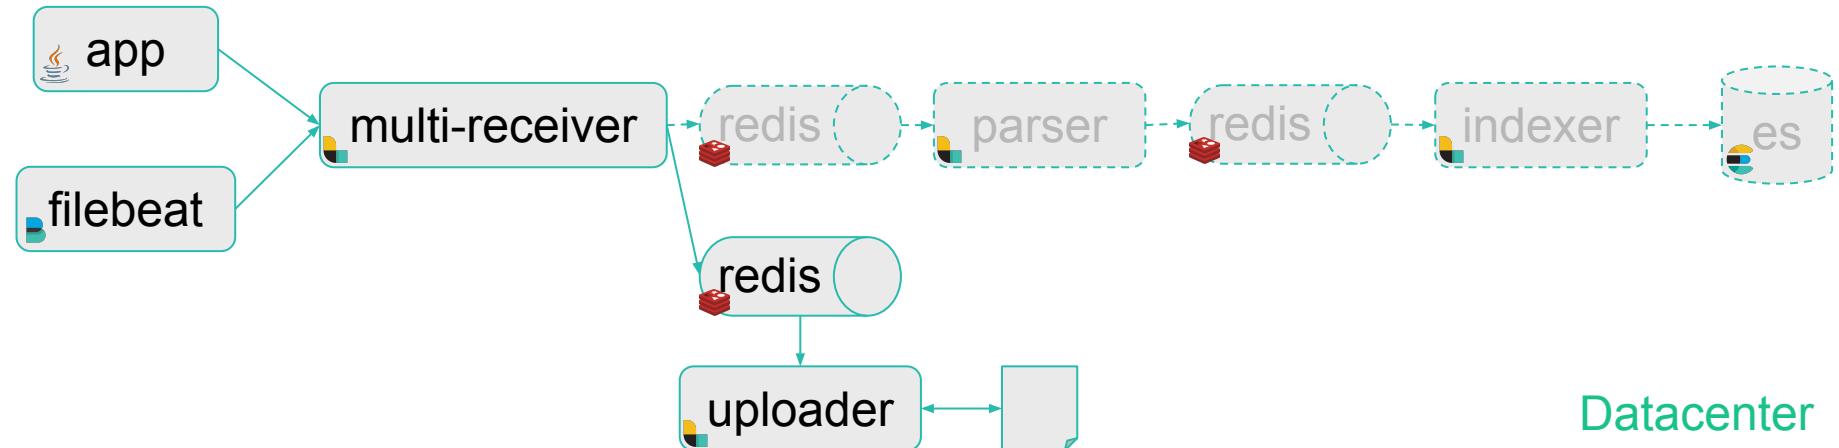




Datacenter

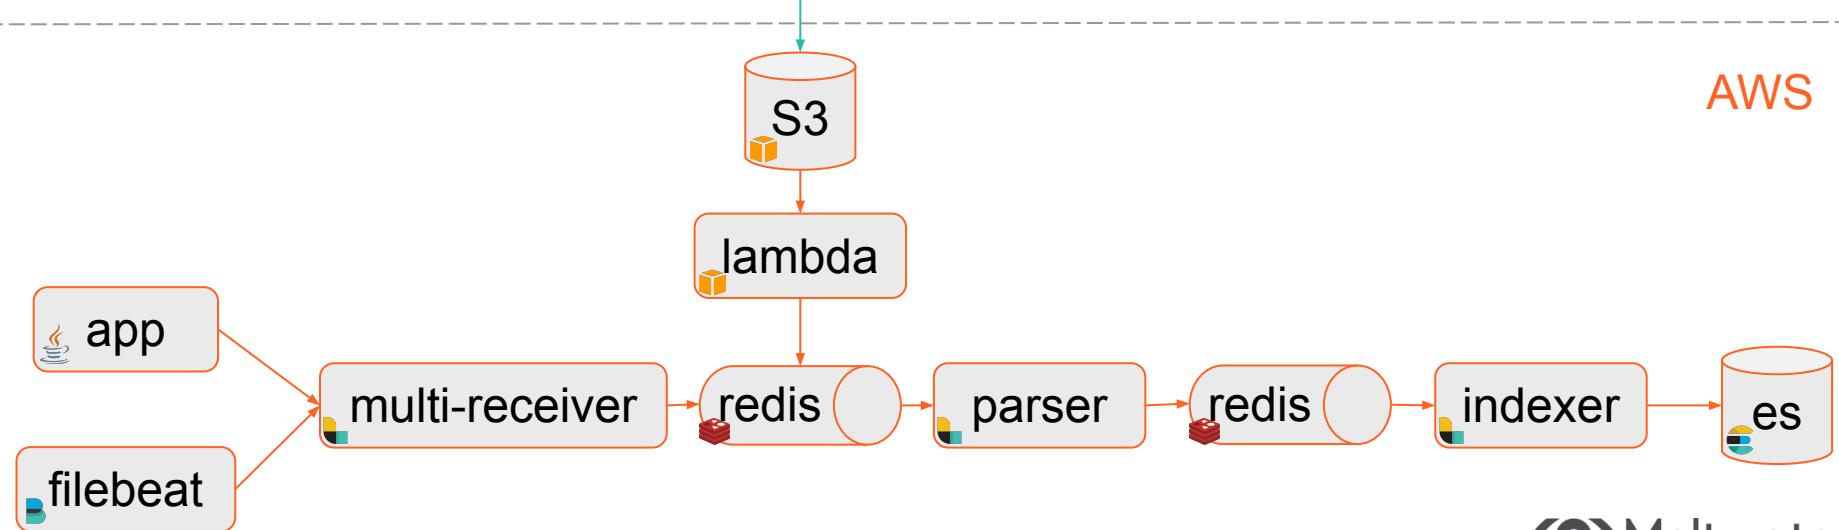
AWS





Datacenter

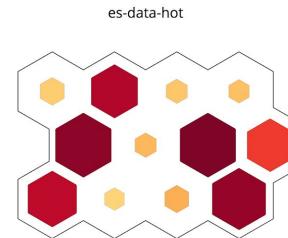
AWS



# Challenge #3

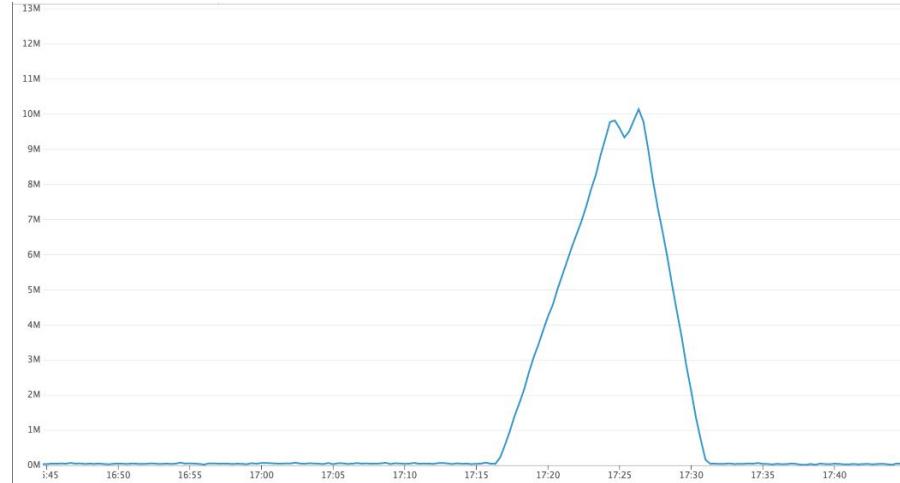
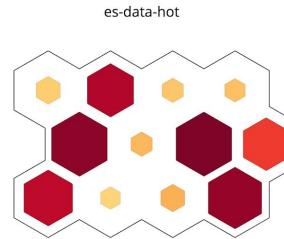
# Challenge

- Elasticsearch hotspots



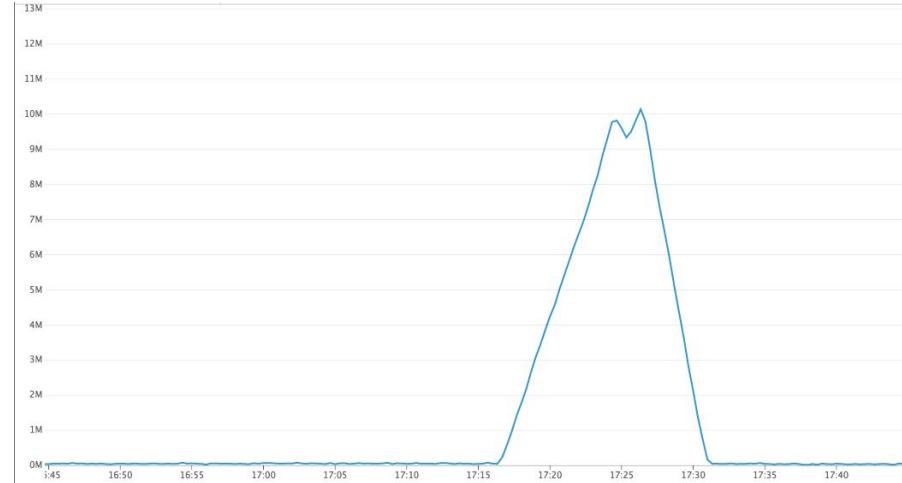
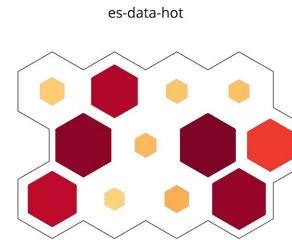
# Challenge

- Elasticsearch hotspots
- Hotspot causes an indexing backlog



# Challenge

- Elasticsearch hotspots
- Hotspot causes an indexing backlog
- Potential root causes
  - Uneven distribution of indices to nodes
  - Data surge

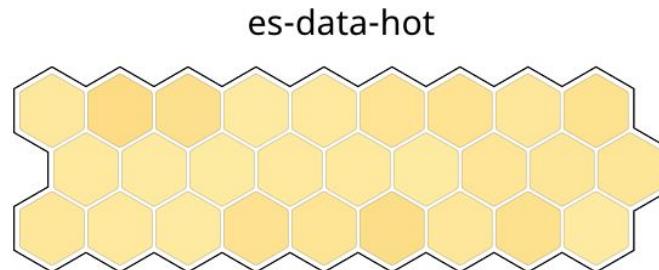


# Challenge

- Elasticsearch hotspots
- Hotspot causes an indexing backlog
- Potential root causes
  - Uneven distribution of indices to nodes
  - Data surge

# Solution

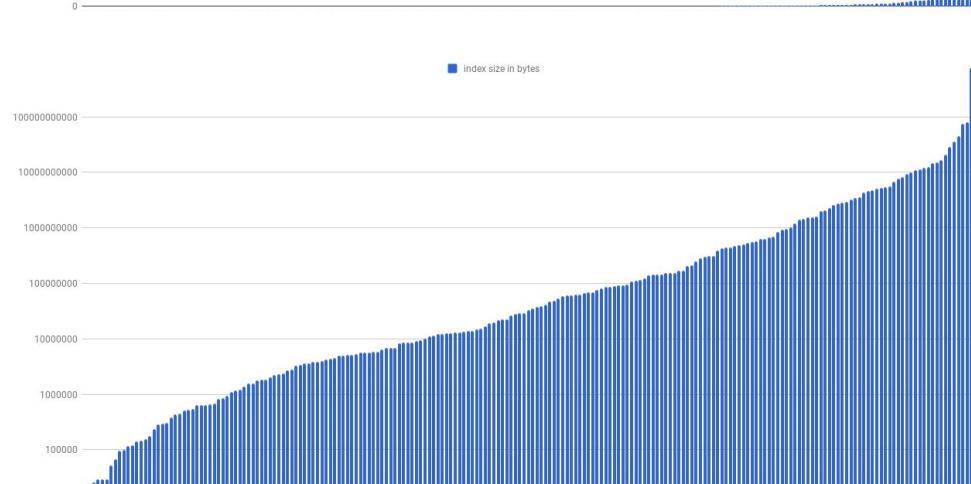
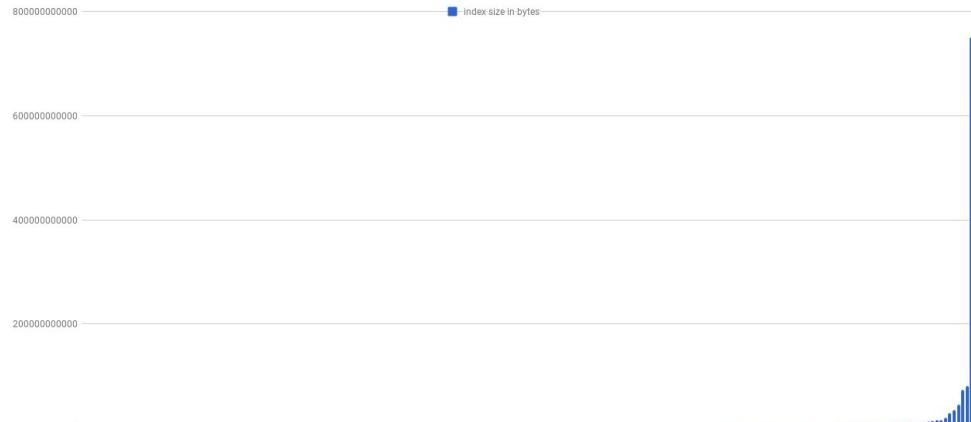
- ~~Manually solve distribution problem~~
  - ~~Move indices around~~
- Don't run the system at 85% capacity
  - Achieve better distribution



# Challenge #4

# Challenge

- Elasticsearch hotspots
- Hotspot causes an indexing backlog
- Potential root causes
  - Uneven distribution of index sizes
  - Too much data per index shard
  - Number of shards per index hardcoded

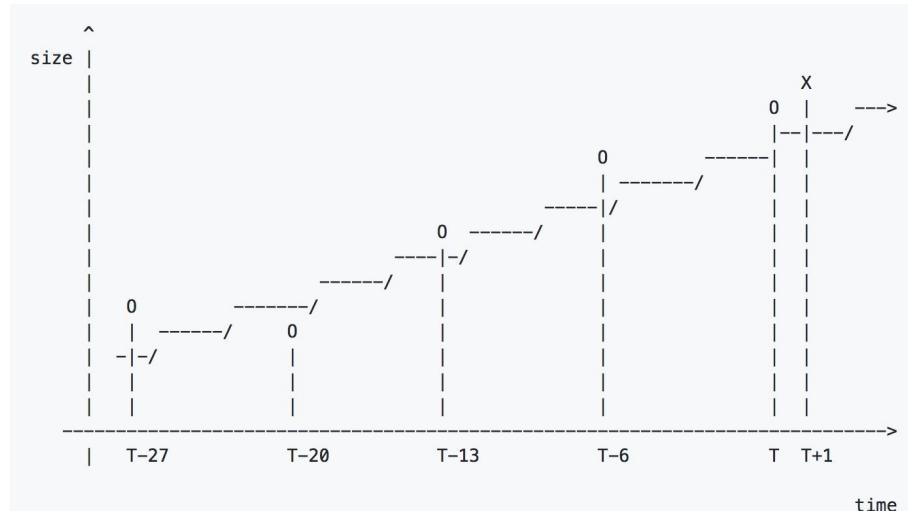


# Challenge

- Elasticsearch hotspots
- Hotspot causes an indexing backlog
- Potential root causes
  - Uneven distribution of index sizes
  - Too much data per index shard
  - Number of shards per index hardcoded

# Solution

- Try to right-size indices
  - Appropriate amount of shards
  - Based on historical data

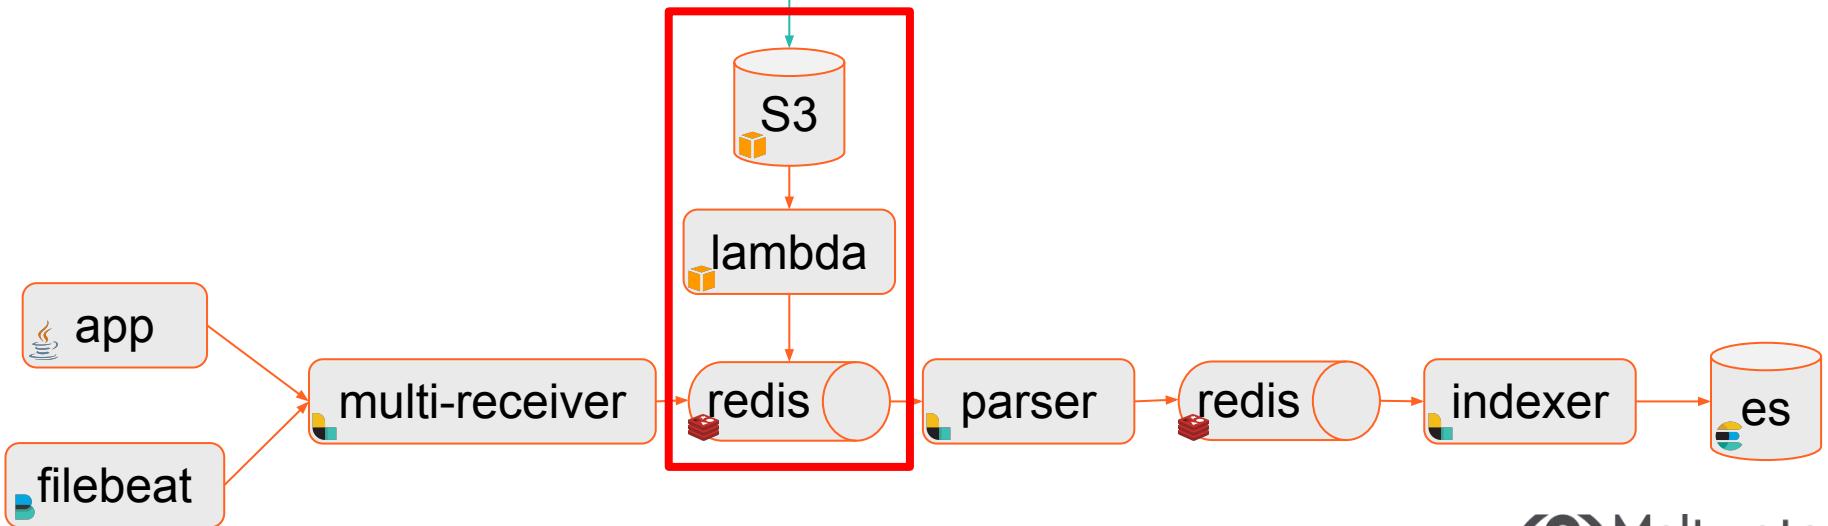
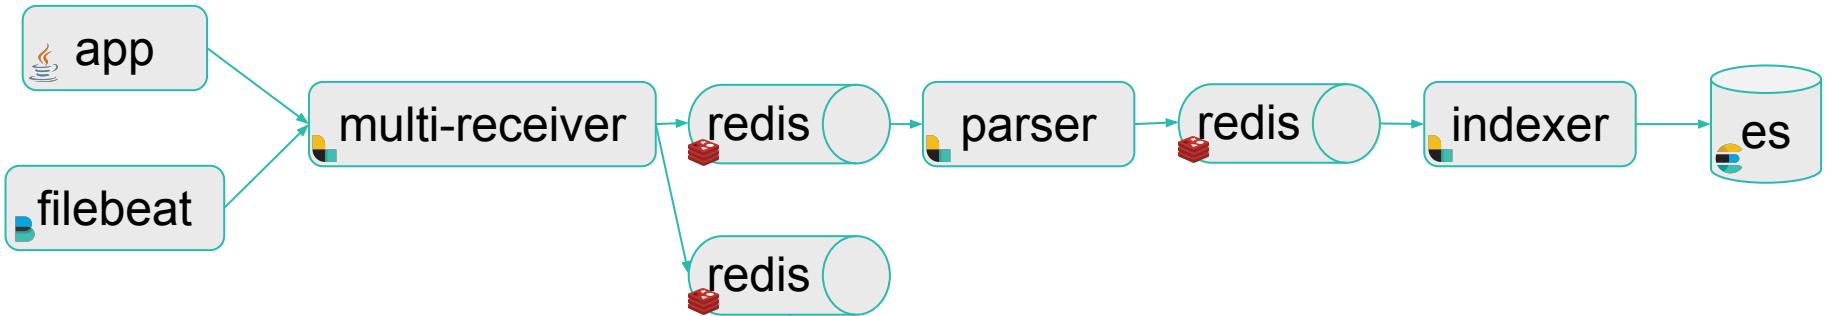




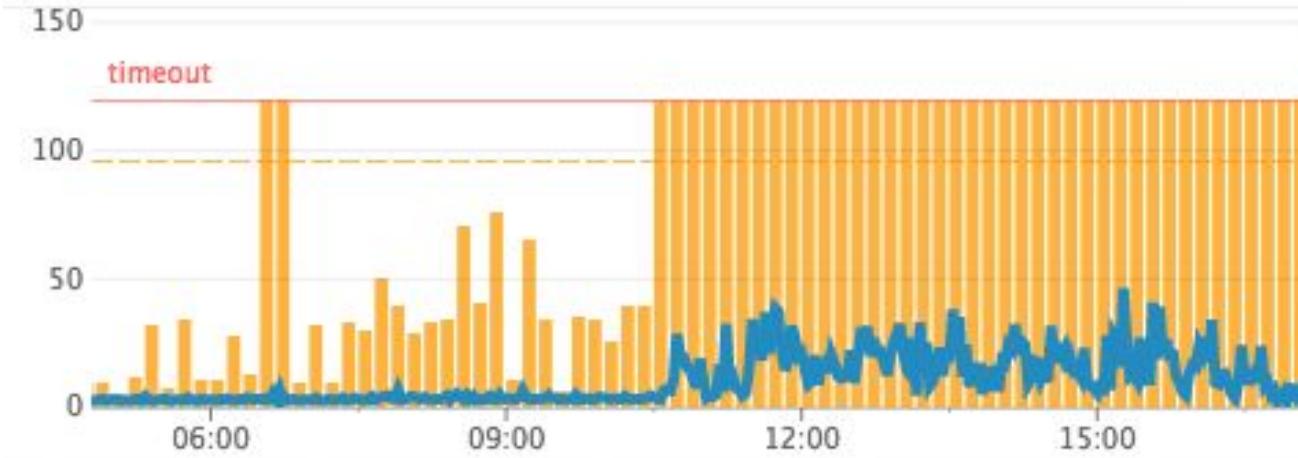
# Challenge #5

# Challenge

- Redis not appropriate anymore
  - In-memory
  - Single-threaded
  - Not a queue (data duplication)
  - Not distributed
  - Hard to scale



## s3-to-redis lambda duration

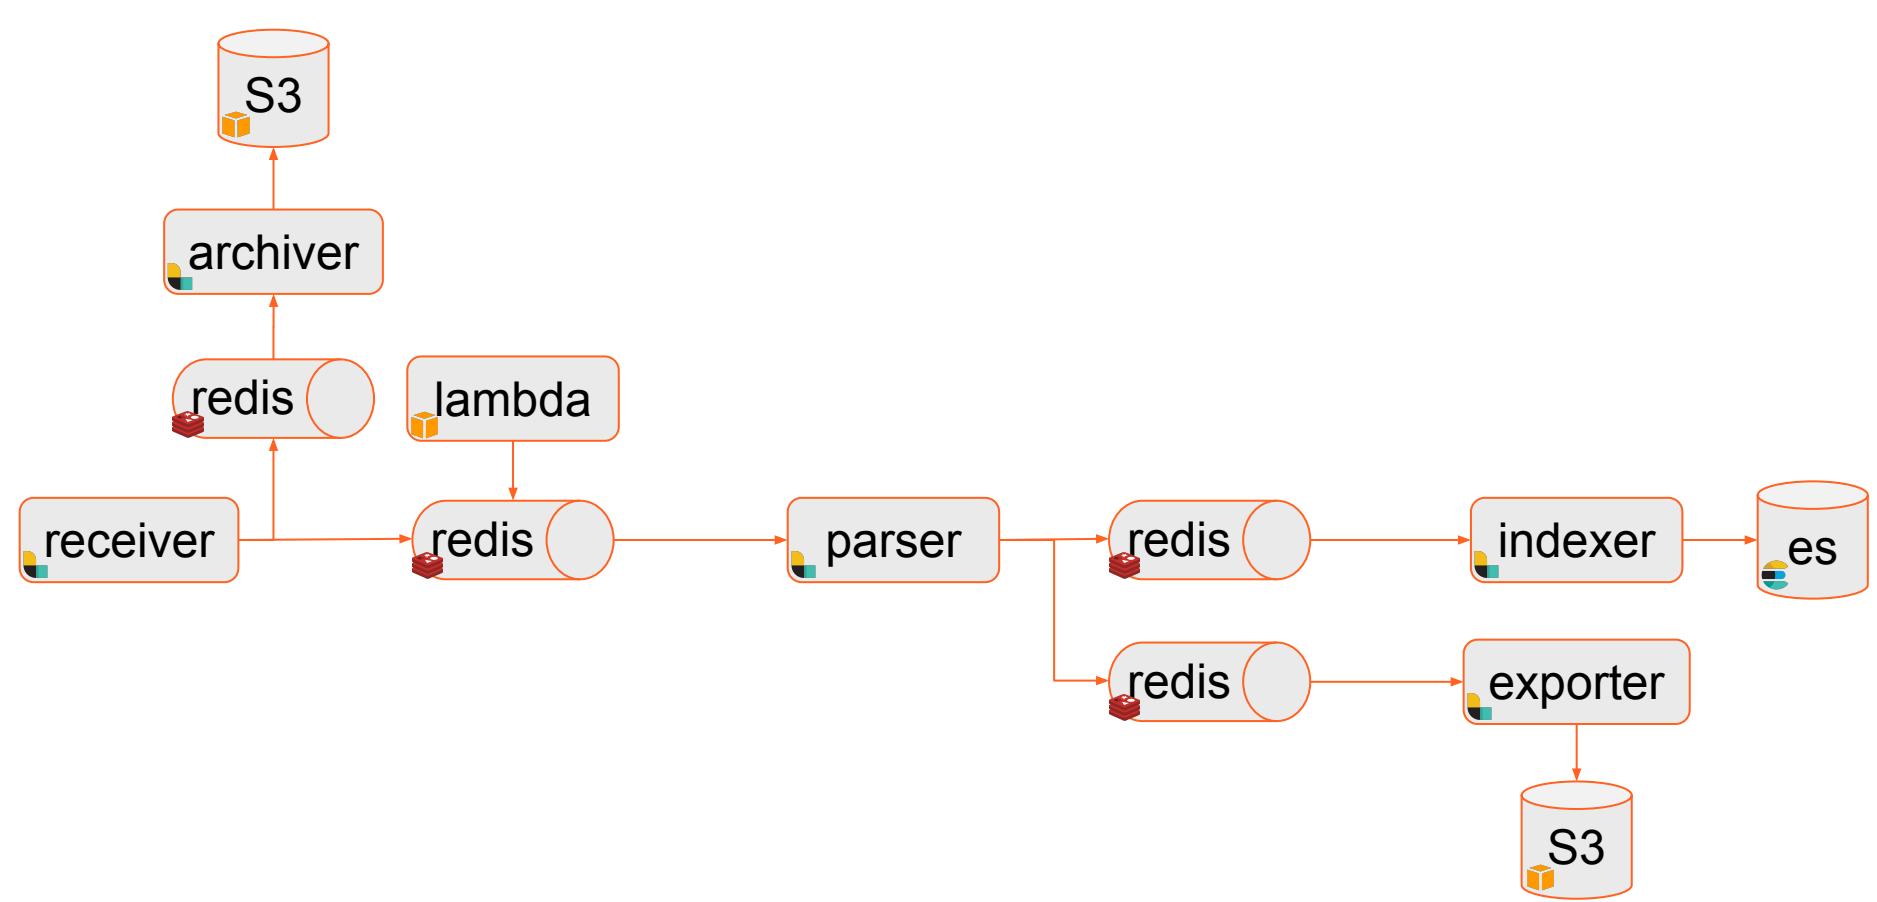


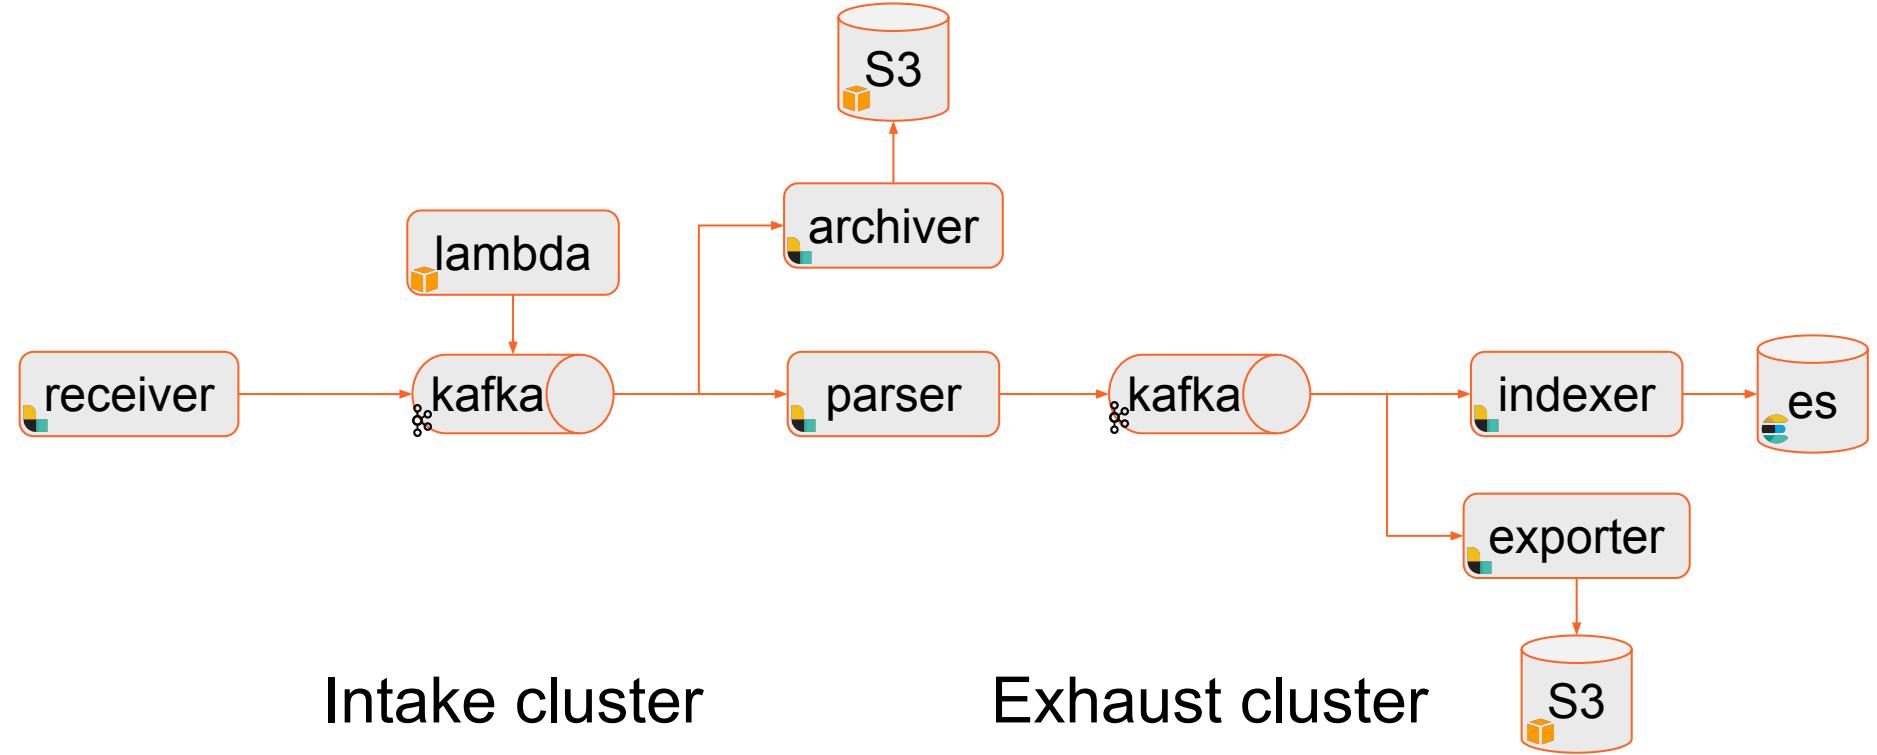
# Challenge

- Redis not appropriate anymore
  - In-memory
  - Single-threaded
  - Not a queue (data duplication)
  - Not distributed
  - Hard to scale

# Solution

- Replace Redis with Kafka
  - Disk-based
  - Multi-threaded
  - Append-only log
  - Distributed
  - Horizontally scalable





# Intake cluster

- Main focus: fast and safe
  - Data not backed up yet
  - Consumers CPU-focused

# Exhaust cluster

- Main focus: huge buffer
  - All data can be replayed
  - Consumers I/O-focused

# Intake cluster

- Main focus: fast and safe
  - Data not backed up yet
  - Consumers CPU-focused
- More nodes

# Exhaust cluster

- Main focus: huge buffer
  - All data can be replayed
  - Consumers I/O-focused
- Fewer nodes

# Intake cluster

- Main focus: fast and safe
  - Data not backed up yet
  - Consumers CPU-focused
- More nodes
- Small, fast disks

# Exhaust cluster

- Main focus: huge buffer
  - All data can be replayed
  - Consumers I/O-focused
- Fewer nodes
- Huge, slower disks

# Intake cluster

- Main focus: fast and safe
  - Data not backed up yet
  - Consumers CPU-focused
- More nodes
- Small, fast disks
- Higher replication

# Exhaust cluster

- Main focus: huge buffer
  - All data can be replayed
  - Consumers I/O-focused
- Fewer nodes
- Huge, slower disks
- Lower replication

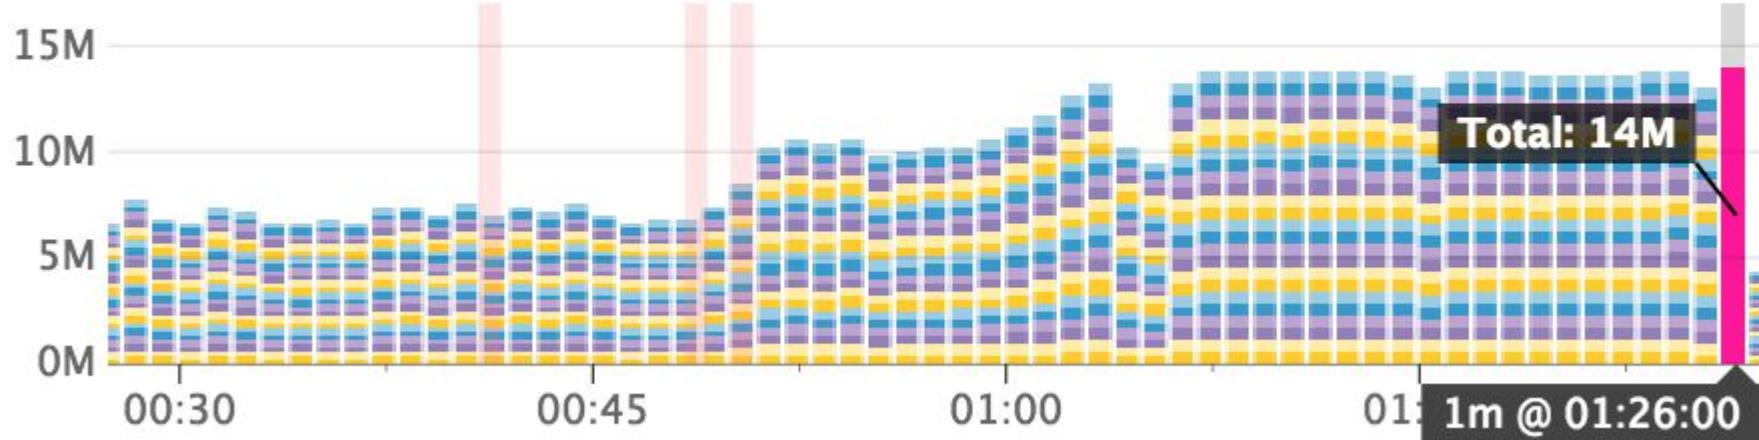
# Intake cluster

- EC2 type: i3.xlarge
  - 4 CPUs
  - 30GB RAM
- Disk: NVMe SSD
  - 0.95 TB
- Count: 3
- Kafka replication: 3
- Partitions per topic: 48

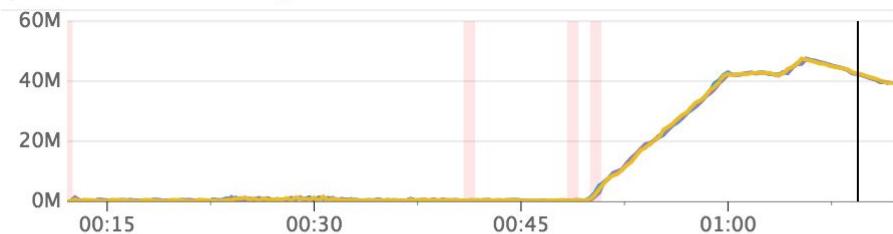
# Exhaust cluster

- EC2 type: r4.xlarge
  - 4 CPUs
  - 30GB RAM
- Disk: st1 EBS
  - 6 TB
- Count: 2
- Kafka replication: 2
- Partitions per topic: 48

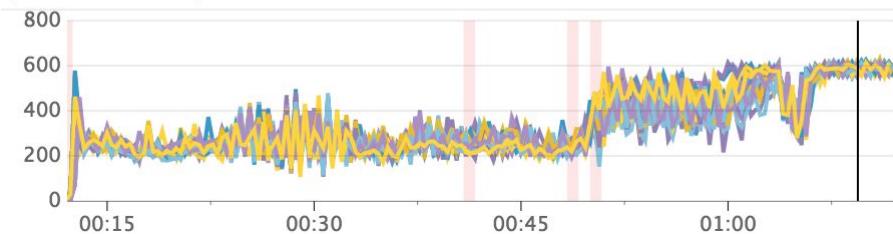
## parser - log events parsed



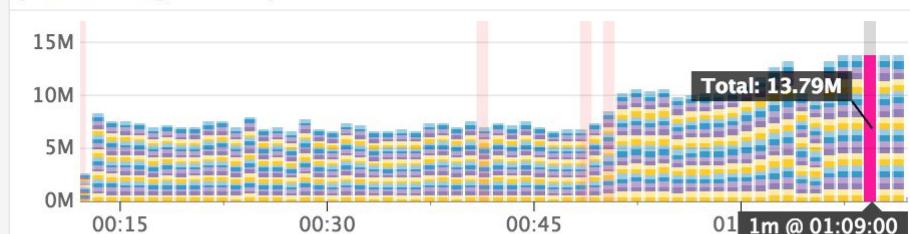
parser - consumer lag



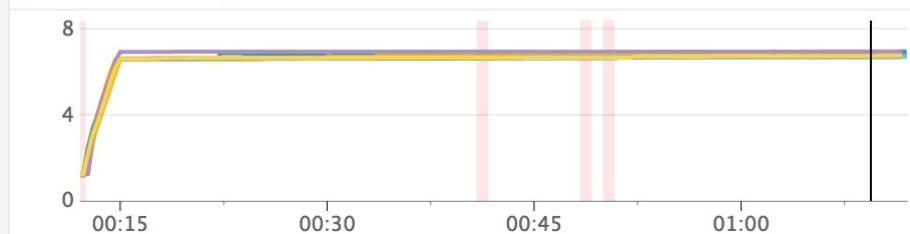
parser - user cpu



parser - log events parsed



parser - memory rss



$$14M / \text{min} = 233\ 333 / \text{s}$$

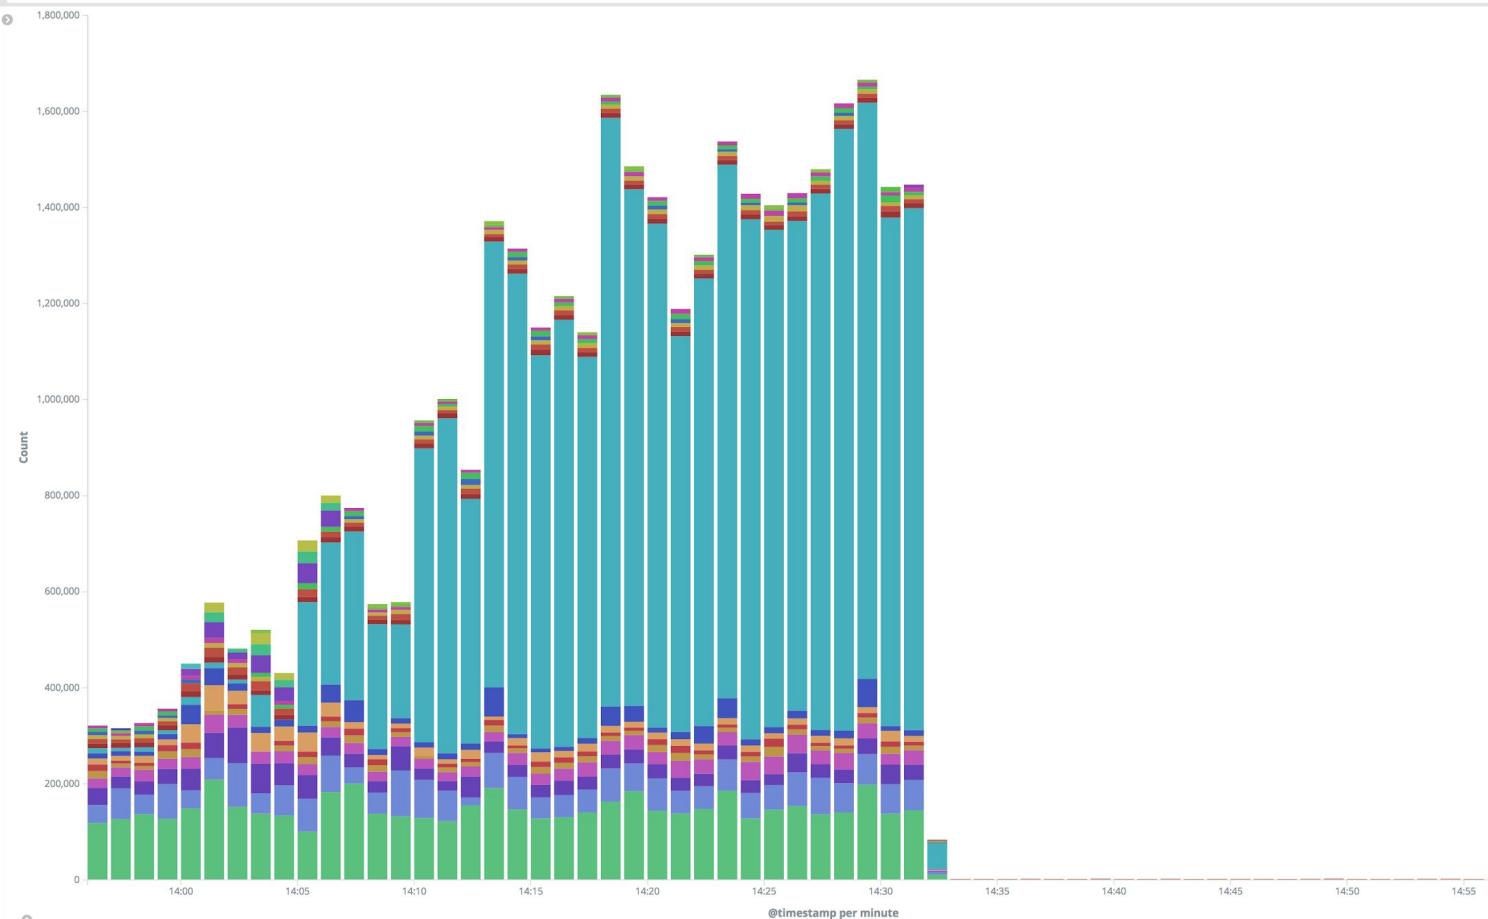
Still bottlenecked on consumer cpu

Kafka as is can do much more

And can still scale

# What next?

```
-      <root level="INFO">
+      <root level="TRACE">
          <appender-ref ref="LOGSTASH"/>
      </root>
```



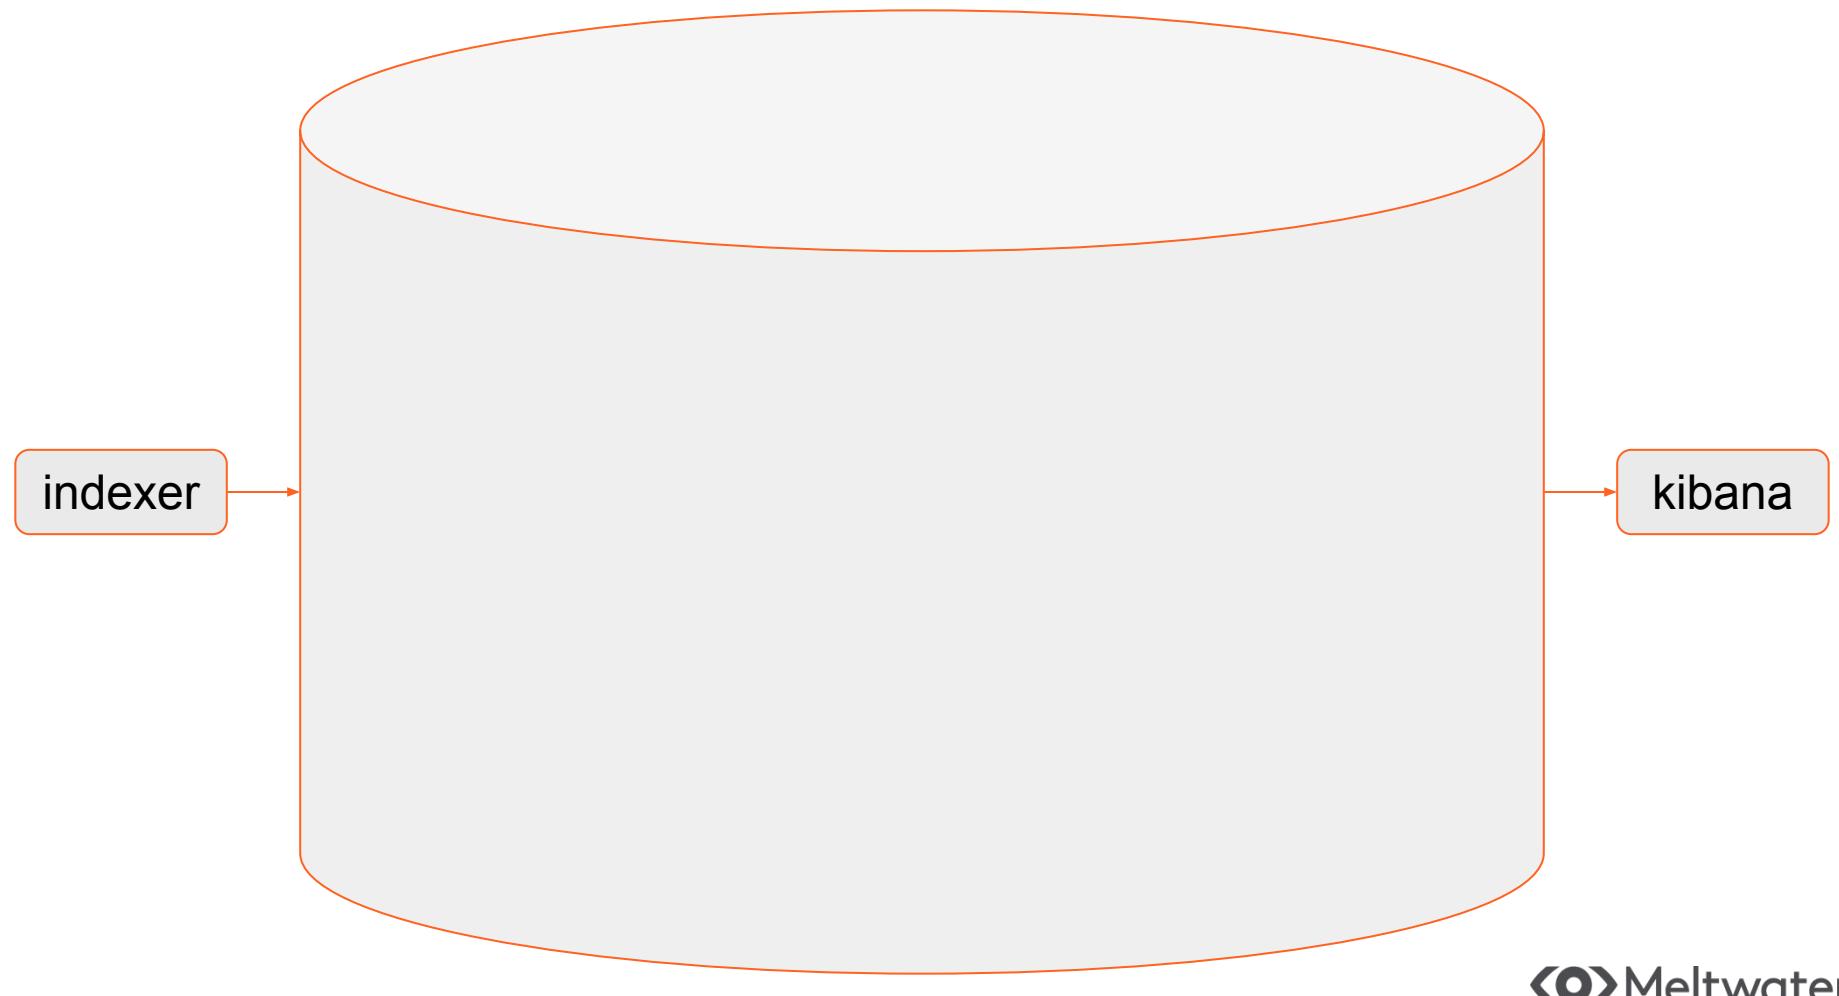
# Challenge

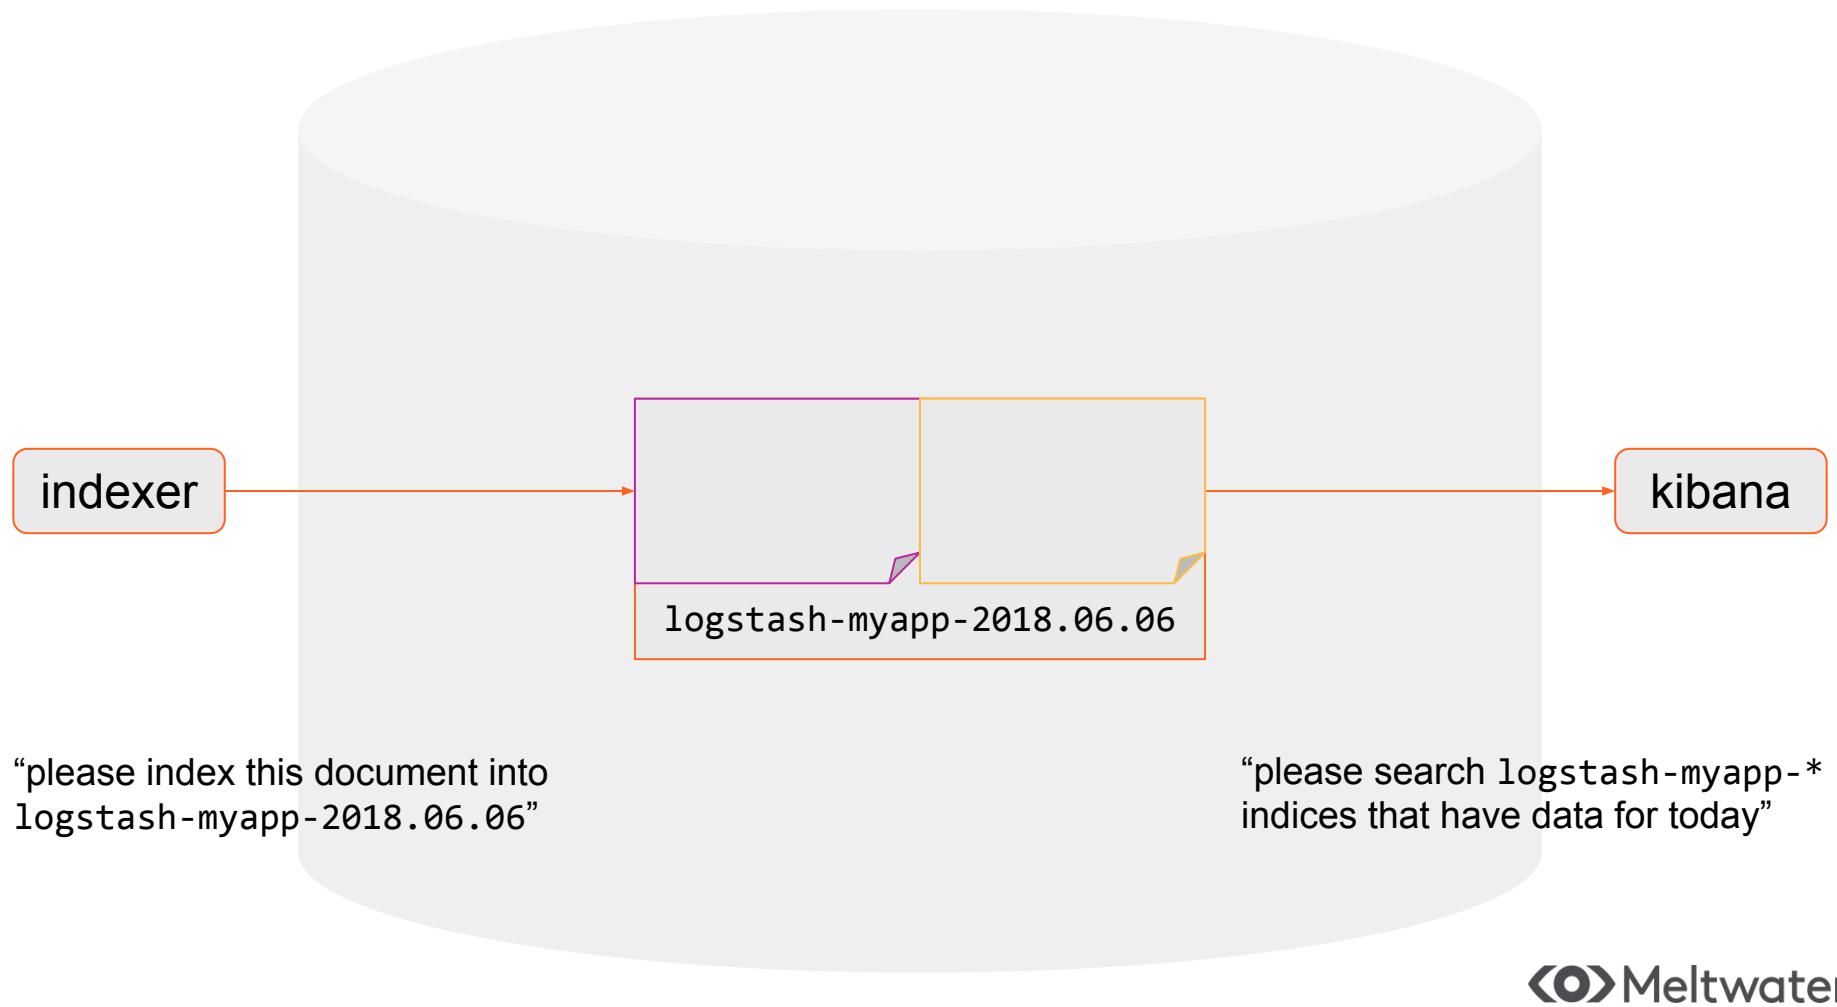
- Hard to react to data spikes
  - Indexing falls behind

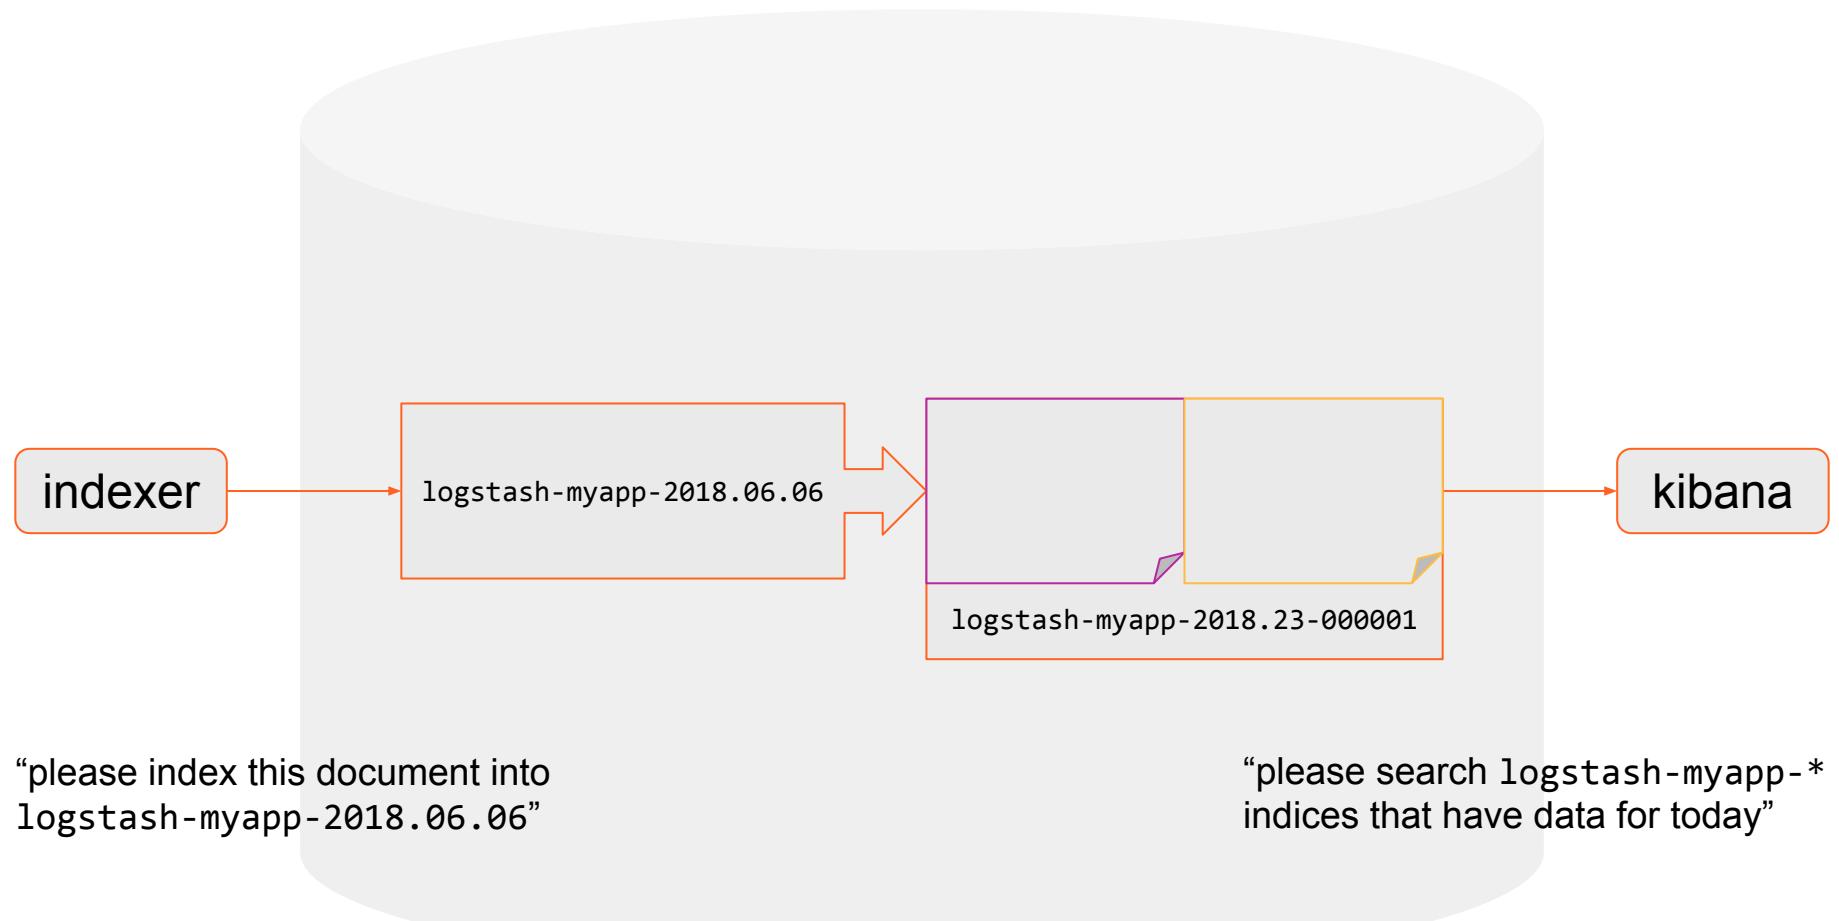
# Solution

- Use Elasticsearch index aliases
- Use Elasticsearch rollover API



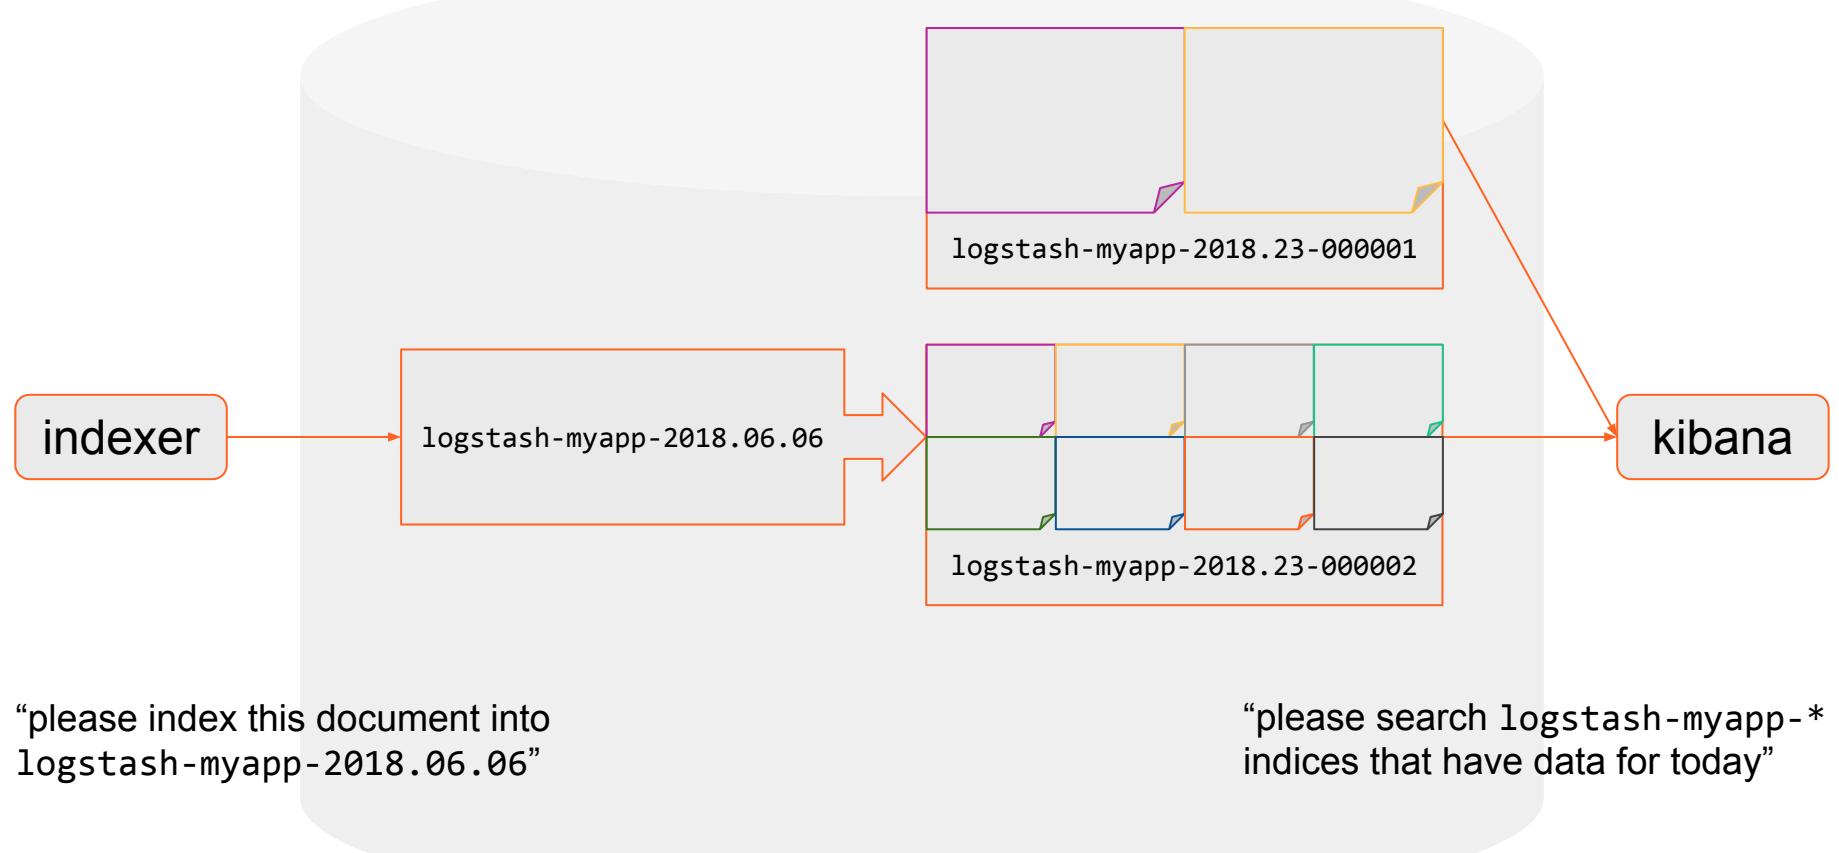




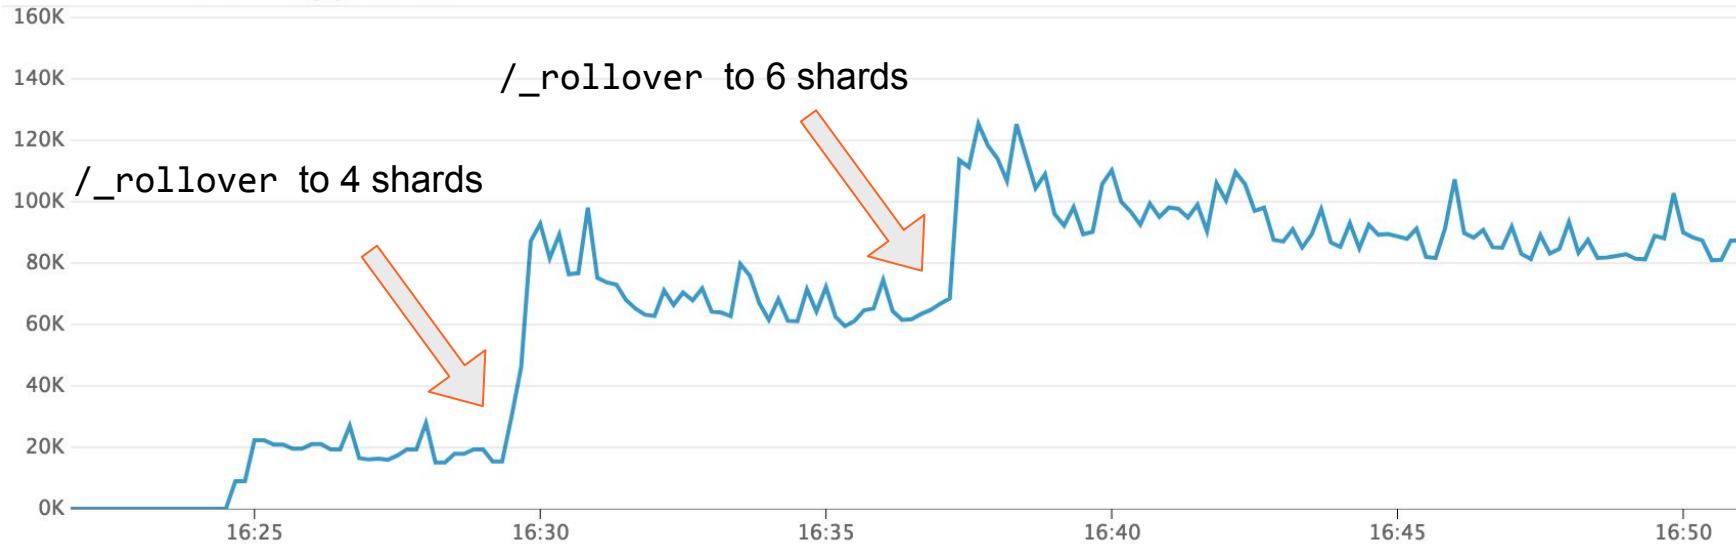


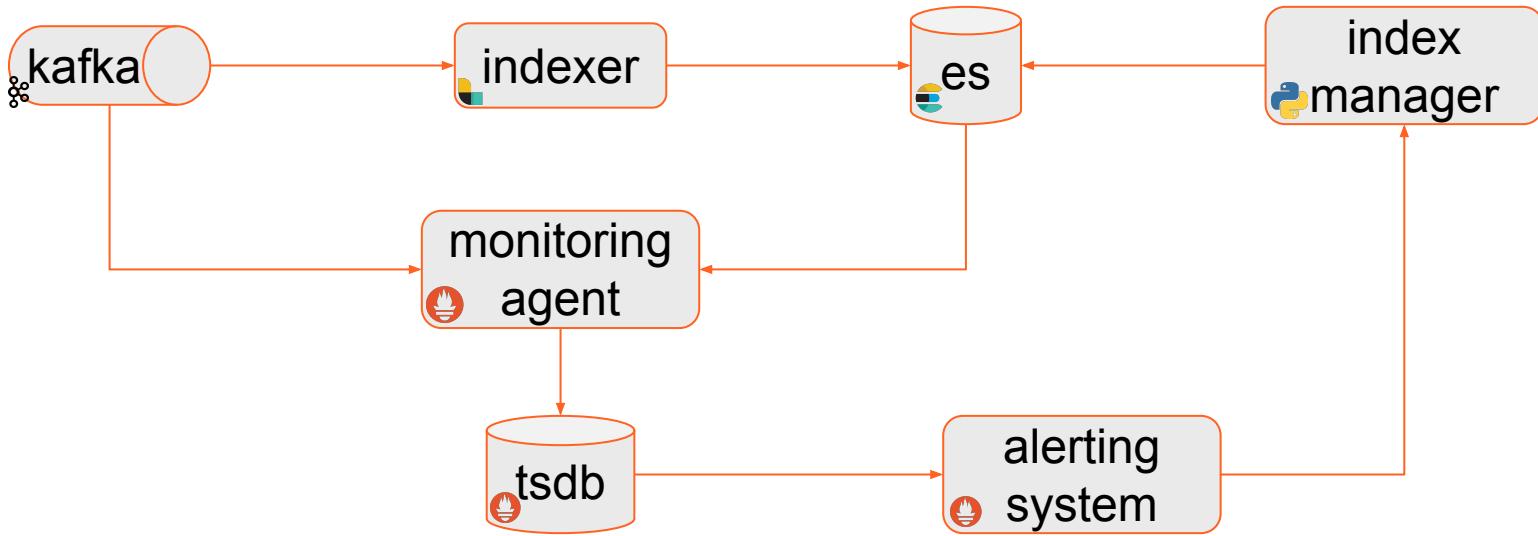
# POST /logstash-myapp-2018.06.06/\_rollover

```
{  
    "settings": {  
        "index.number_of_shards": 8  
    }  
}
```



elasticsearch indexing per second





To a billion messages  
and (well) beyond

# Thanks!

János Csorvási

[janos.csorvasi@meltwater.com](mailto:janos.csorvasi@meltwater.com)

[linkedin.com/in/jcsorvasi](https://linkedin.com/in/jcsorvasi)

Meltwater

[underthehood.meltwater.com](http://underthehood.meltwater.com)

