
Zero downtime Kubernetes cluster upgrades

Simone Sciarrati






Meltwater

@meltwatereng

Mission



FOUNDATION

"We enable teams to accelerate without thinking too much about the infrastructure, allowing a rapid path from ideation to prototype to providing business value."



kops



dex



CNI



Thanos



 Meltwater



What could possibly....

- AWS limits
- CNI plugin, CNI plugin, CNI plugin....
- Kops
- Cross account IAM
- Network connectivity / DNS
- Much more...
- ... and Kubernetes

kubelet fails to heartbeat with API server with stuck TCP connections #48638



Closed

derekwaynecarr opened this issue on 7 Jul 2017 · 33 comments



derekwaynecarr commented on 7 Jul 2017 • edited

Member

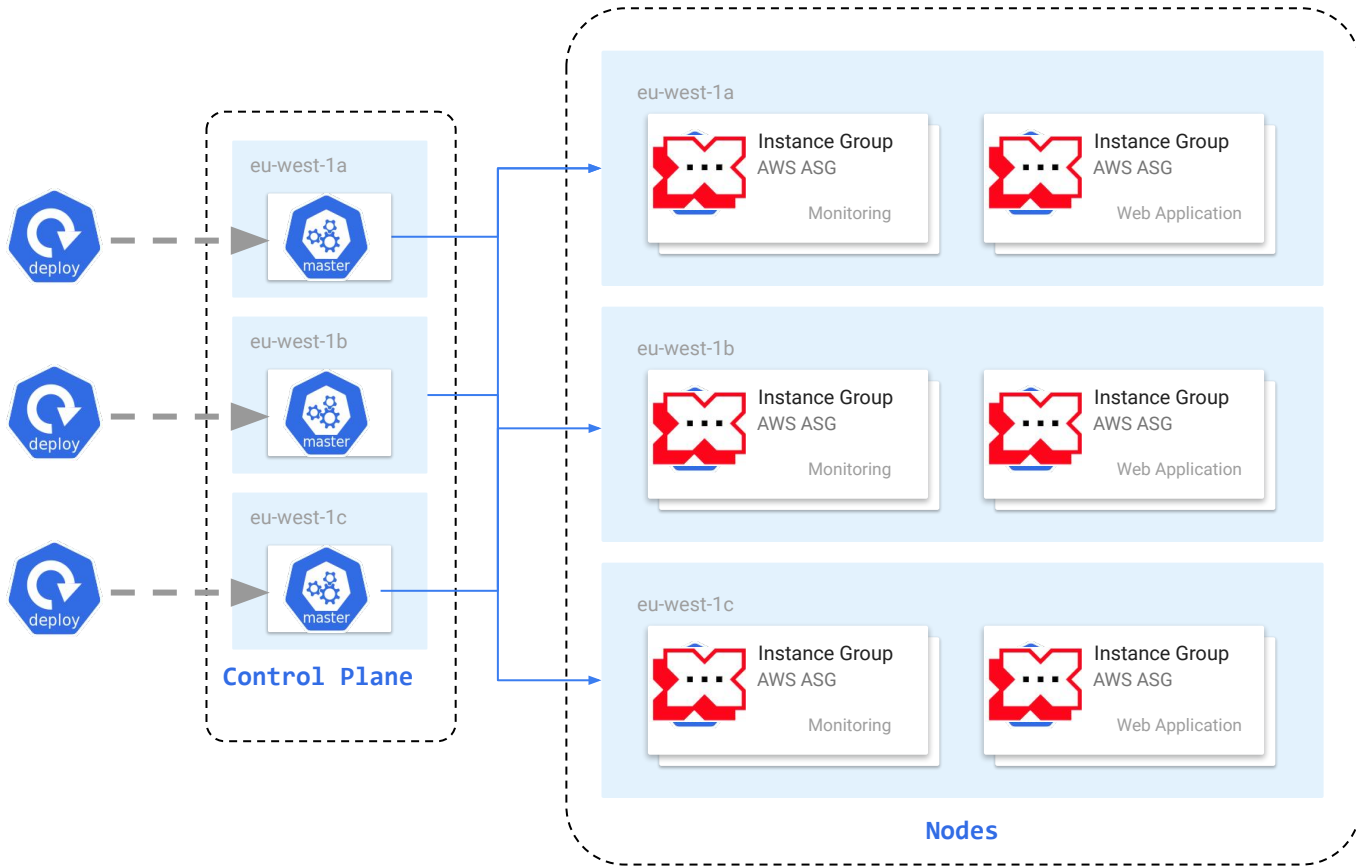


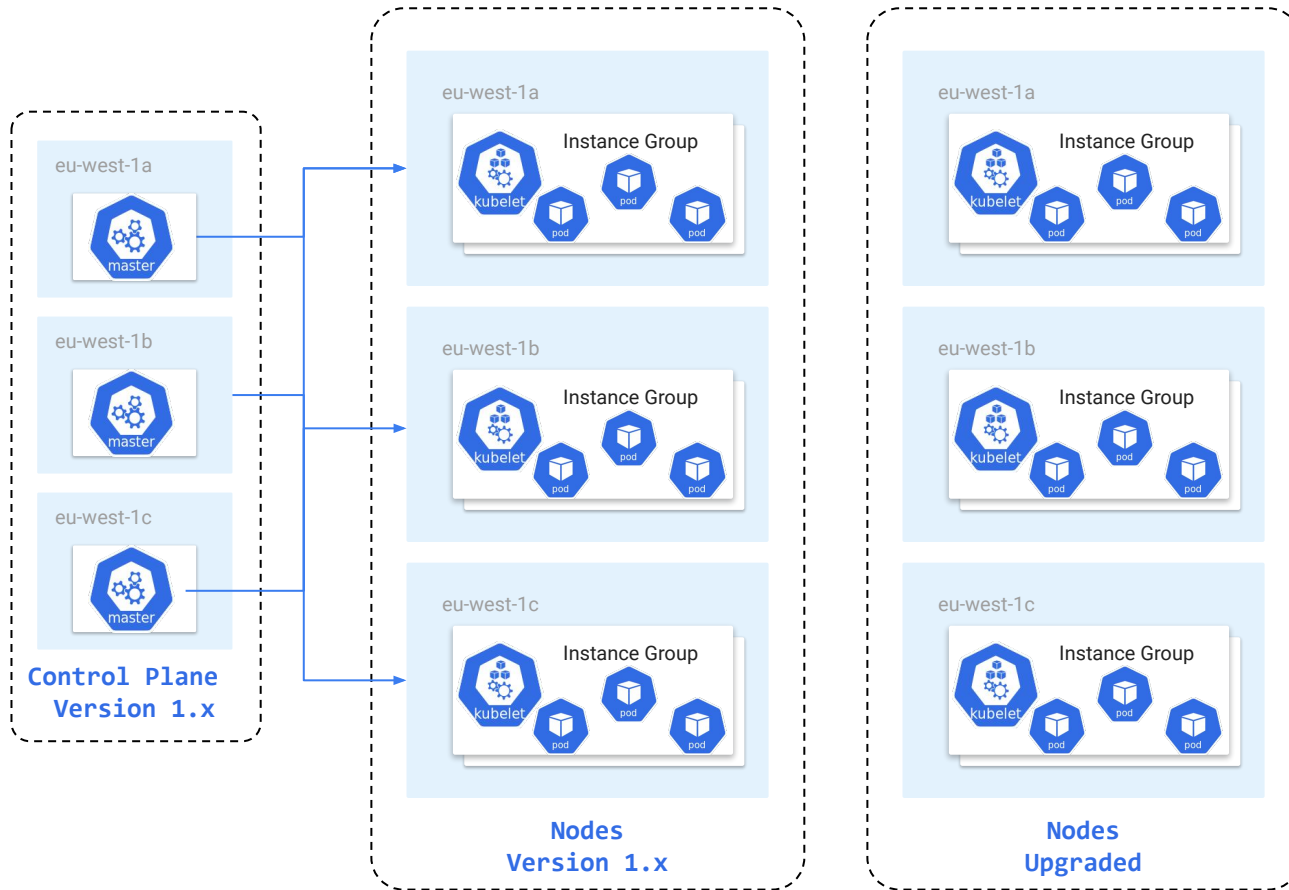
Is this a BUG REPORT or FEATURE REQUEST?:

/kind bug

What happened:

operator is running an HA master setup with a LB in front. kubelet attempts to update node status, but `tryUpdateNodeStatus` wedges. based on the goroutine dump, the wedge happens when it attempts to GET the latest state of the node from the master. operator observed 15 minute intervals between attempts to update node status when kubelet could not contact master. assume this is when the LB ultimately closes the connection. the impact is that node controller then marked node as lost, and workload was evicted.

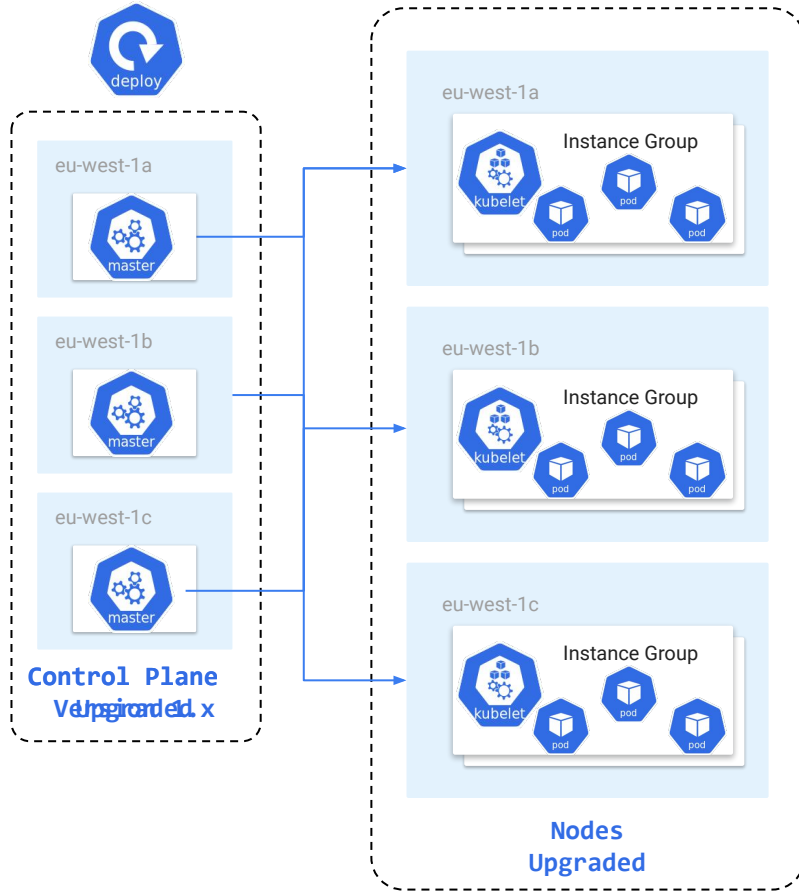




`kubectl cordon node`

`kubectl drain node`

*PodDisruptionBudget
Lifecycle Hooks
help minimising the
impact of draining*




Definition of “working”

-
- Kubernetes API and components
 - Dex
 - Deployments
 - Ingress/Services
 - StatefulSets
 - External-DNS
 - CNI
 - Metrics
 - Logging
 - etc


Pre-flight checks:


- **Kops:** changes since last release, default flags after upgrade, breaking changes etc
- **Kubernetes:** API deprecations, changes to the components behaviour etc
- **AWS:** limits (monitored via a prometheus exporter and aws trusted advisor)
- **Components:** Changes or upgrades (*-dns, kube-proxy, CA, CNI plugin, etc)
- Identify changes in behaviour or deprecations that need to be communicated to the users


Validate the new configuration on a small scale

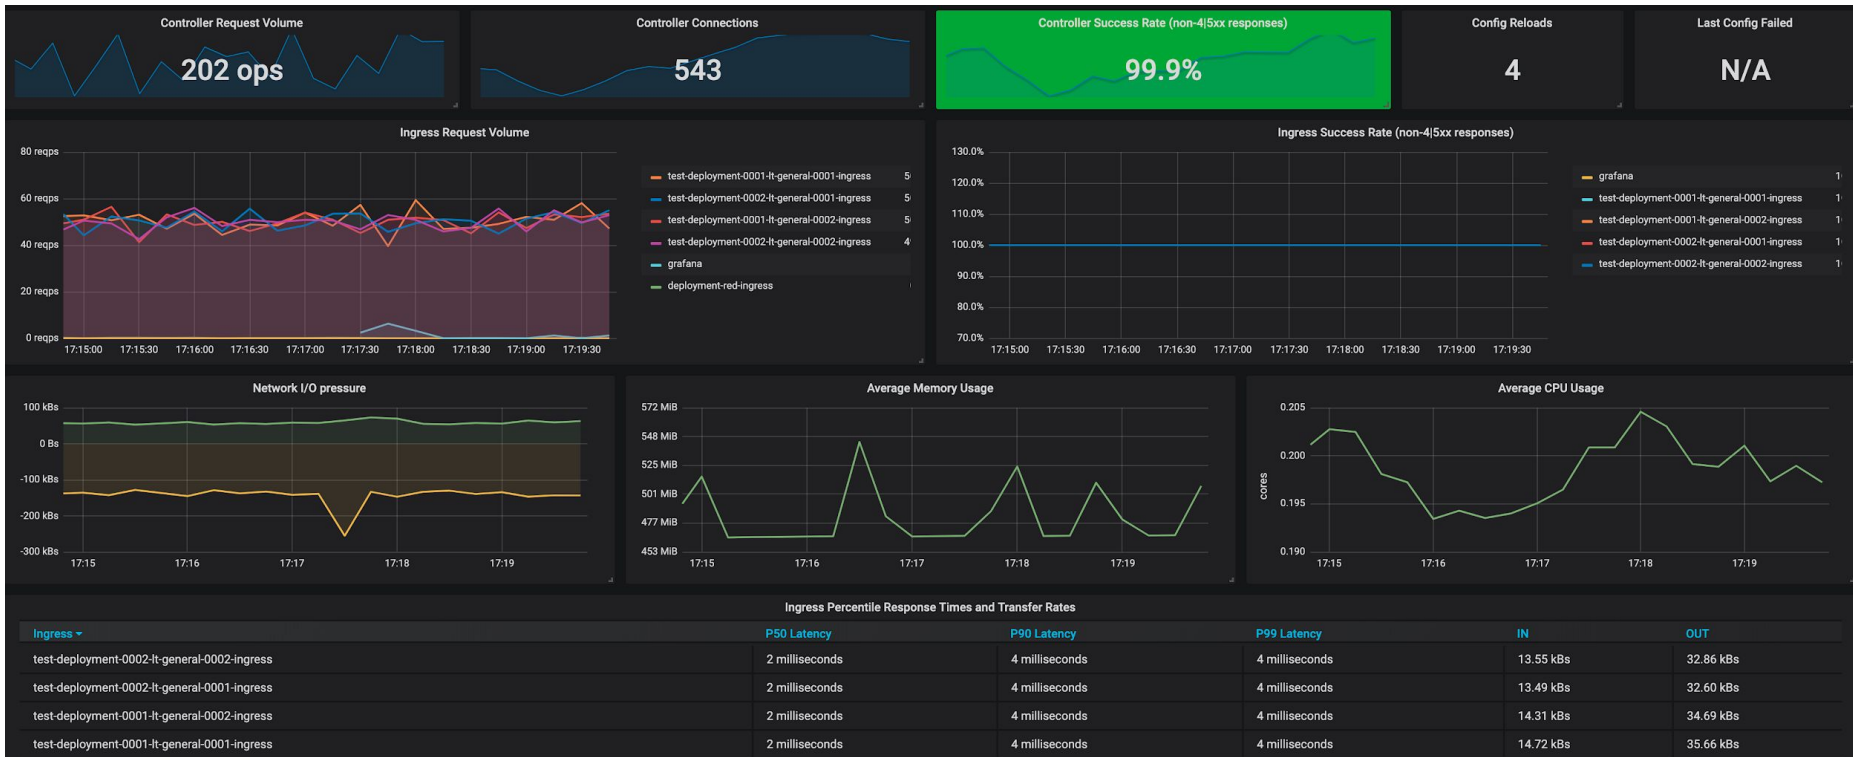










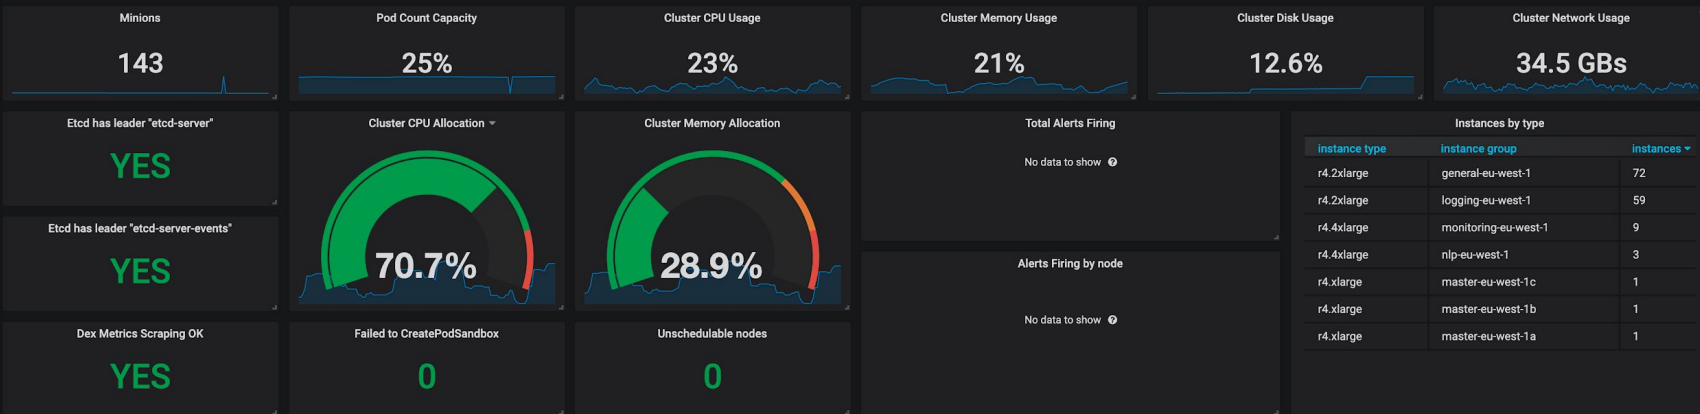




> AWS CNI metrics (9 panels)

> Configured Prometheus Alerts (15 panels)

~ System Cluster Wide



underthehood.meltwater.com
 [@MeltwaterEng](https://twitter.com/MeltwaterEng)



dezmodue



@dezmodue