
Scalability testing of a production Kubernetes cluster

Federico Hernandez & Simone Sciarrati



30000+ customers

1500+ employees

55 global offices

Founded 2001 in Oslo

```
1 //fires the appear event when appropriate
2 var check = function() {
3     //is the element hidden?
4     if (!t.is(':visible')) {
5         //it became hidden
6         t.appeared = false;
7         return;
8     }
9
10    //is the element inside the visible window?
11    var a = w.scrollLeft();
12    var b = w.scrollTop();
13    var o = t.offset();
14    var x = o.left;
15    var y = o.top;
16
17    var ax = settings.accX;
18    var ay = settings.accY;
19    var th = t.height();
20    var wh = w.height();
21    var tw = t.width();
22    var ww = w.width();
23
24    if (y + th + ay >= b &&
25        y <= b + wh + ay &&
26        x + tw + ax >= a &&
27        x <= a + ww + ax) {
28
29        //trigger the custom event
30        if (!t.appeared) t.trigger('appear', settings.data);
31
32    } else {
33
34        //it scrolled out of view
35        t.appeared = false;
36    }
37 };
38
39 //create a modified fn with some additional logic
40 var modifiedFn = function() {
41
42     //mark the element as visible
43     t.appeared = true;
44
45     //is this supposed to happen only once?
46     if (settings.one) {
47
48         //remove the check
49         w.unbind('scroll', check);
50         var i = $.inArray(check, $.fn.appear.checks);
51         if (i >= 0) $.fn.appear.checks.splice(i, 1);
52     }
53
54     //trigger the original fn
55     fn.apply(this, arguments);
56 }
57
58 //bind the modified fn to the element
59 t.one('appear', settings.data, modifiedFn);
60
61 //bind the one() to the element
62 t.one('appear', settings.data, modifiedFn);
```

9+ engineering offices (7+ countries)

350+ engineers in 50+ teams

AWS with 170+ accounts

December 2018

May 2019

October 2019



71

146

265



1251

2102

3954



283

561

944



150

435

818



25

77

103

Kubernetes Failure Stories

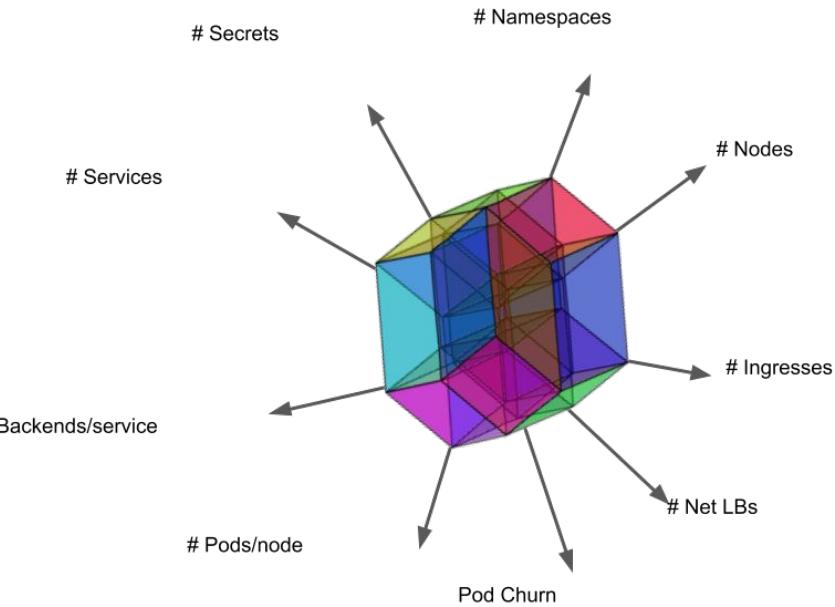
A compiled list of links to public failure stories related to Kubernetes. Most recent publications on top.

- A Kubernetes crime story - Prezi - blog post 2019
 - involved: AWS EKS, SNAT, conntrack, Amazon VPC CNI plugin
 - impact: delay of 1-3 seconds for outgoing TCP connections
- Postmortem: New K8s workers unable to join cluster - FREE NOW - postmortem 2019
 - involved: AWS spot instances, kops, CentOS, container-selinux
 - impact: insufficient cluster capacity in testing environments, failed deployments in production and staging environments
- How a simple admission webhook lead to a cluster outage - Jetstack - blog post 2019
 - involved: ValidatingWebhookConfiguration, GKE node auto-repair
 - impact: prolonged downtime of non-prod environment, nodes lost, failed master upgrade
- Post Mortem: Kubernetes Node OOM - Blue Matador - blog post 2019
 - involved: AWS, SystemOOM, EBS, fluentd-sumologic, no resource requests/limits
 - impact: unknown, Pods killed
- Kubernetes' dirty endpoint secret and Ingress - Ravelin - blog post 2019
 - involved: GKE, Ingress, replication controller, SIGTERM, "graceful shutdown"
 - impact: occasional 502 errors
- How a Production Outage Was Caused Using Kubernetes Pod Priorities - Grafana Labs 2019
 - involved: Pod priorities
 - impact: cascading Pod evictions
- Moving to Kubernetes: the Bad and the Ugly - Xing - ContainerDays EU 2019
 - involved: nginx Ingress, network interrupts, conntrack, frozen CronJob, PLEG, stuck controllers
 - impact: lost requests, response time jumps, not ready nodes
- Kubernetes Failure Stories, or: How to Crash Your Cluster - Zalando - ContainerDays EU 2019
 - involved: AWS IAM, Kubelet, `--kube-api-qps`, Skipper-Ingress, AWS, `OOMKill`, CronJob, CoreDNS, CPU throttling
 - impact: build errors, production outages
- Build Errors of Continuous Delivery Platform - Zalando - postmortem 2019



Kubernetes official thresholds

Quantity	Threshold: scope=namespace	Threshold: scope=cluster
#Nodes	n/a	5000
#Namespaces	n/a	10000
#Pods	3000	150000
#Pods per node	min(110, 10*#cores)	min(110, 10*#cores)
#Services	5000	10000
#All service endpoints	TBD	TBD
#Endpoints per service	250	n/a
#Secrets	TBD	TBD
#ConfigMaps	TBD	TBD
#Deployments	2000	TBD
#DaemonSets	TBD	TBD
#Jobs	TBD	TBD
#StatefulSets	TBD	TBD



Source of hypercube image: <http://www.greggman.net/APPLETS/29/29.html>

Kubernetes official SLOs

Latency of mutating API calls, 99th percentile: < 1s

Latency of non-streaming read-only API class, 99th percentile: < 1s

Startup latency of schedulable stateless pods, 99th percentile: < 5s

Startup latency of schedulable stateful pods, 99th percentile: < ?s

Latency of programming DNS, 99th percentile: < ?s

In-cluster network latency from a single pod, 99th percentile: < ?s

Under the condition that

Runs a single or more appropriately sized master machines

Events are stored in a separate etcd instance (or cluster)

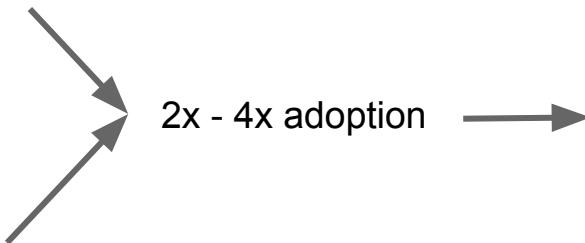
All etcd instances are running on master machine(s)

Cluster churn is ≤ 20 , where churn is defined as:

```
Churn = # (Pod spec creations/updates/deletions)
      + # (user originated requests) in a given second
```

Defining a real-world scalability test

- Current cluster size
- Architecture
- Utilization
- Thresholds
- SLOs



Nodes	NS	Dep/NS	Pods/Dep
2x	=	=	=
=	2x	2x	=
2x	2x	2x	=
4x	2x	2x	=

Known parameters

Hypothesis

Validation





Grafana

› **k8s objects and resources** (14 panels)

› **pod and node state** (2 panels)

› **aws cni** (32 panels)

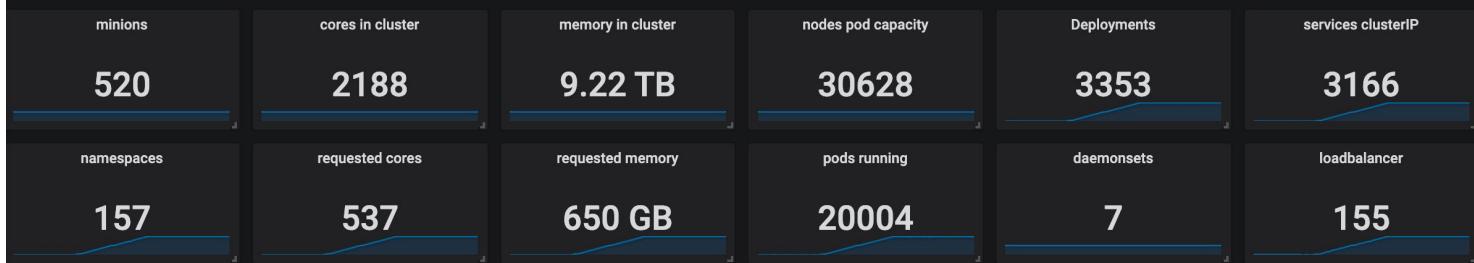
› **etcd** (16 panels)

› **kube-dns** (2 panels)

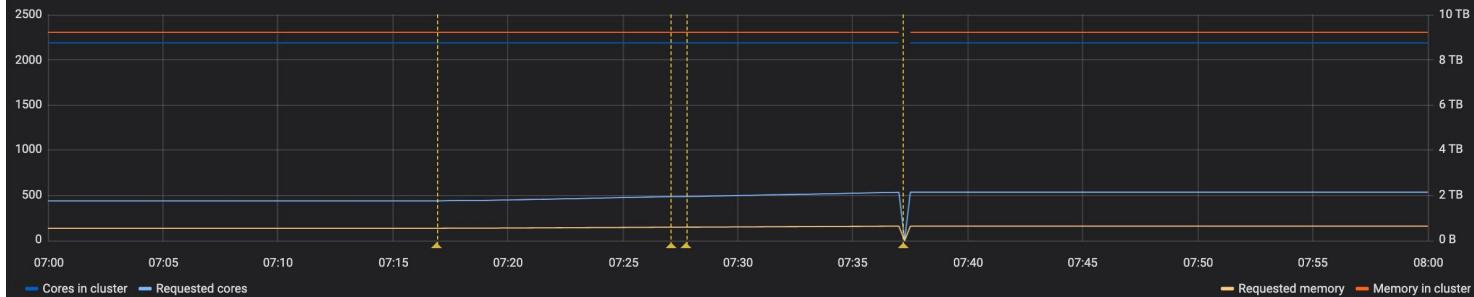
› **masters cpu memory** (6 panels)

› **k8s slo** (5 panels)

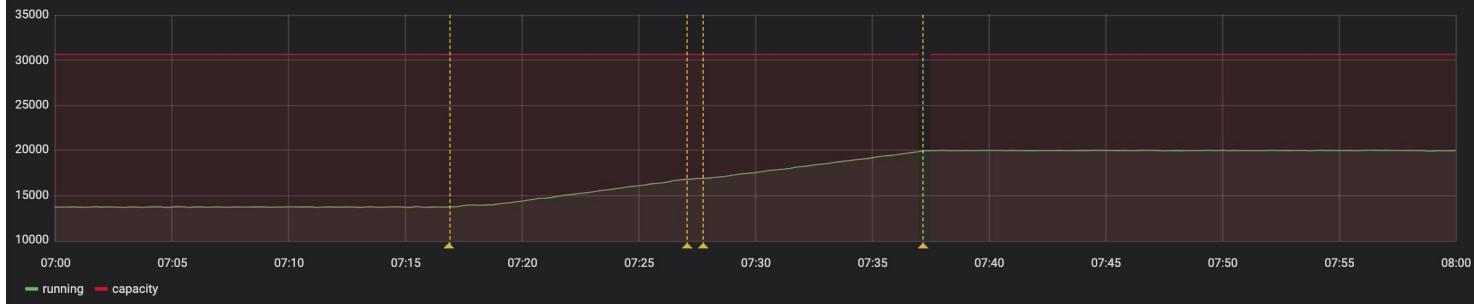
▼ k8s objects and resources



cluster cores and memory

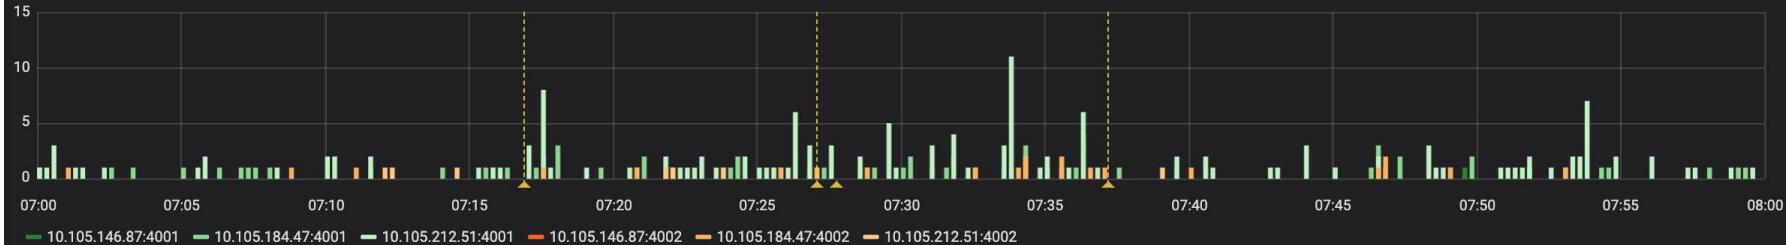


pod usage



etc

etc_server_proposals_pending

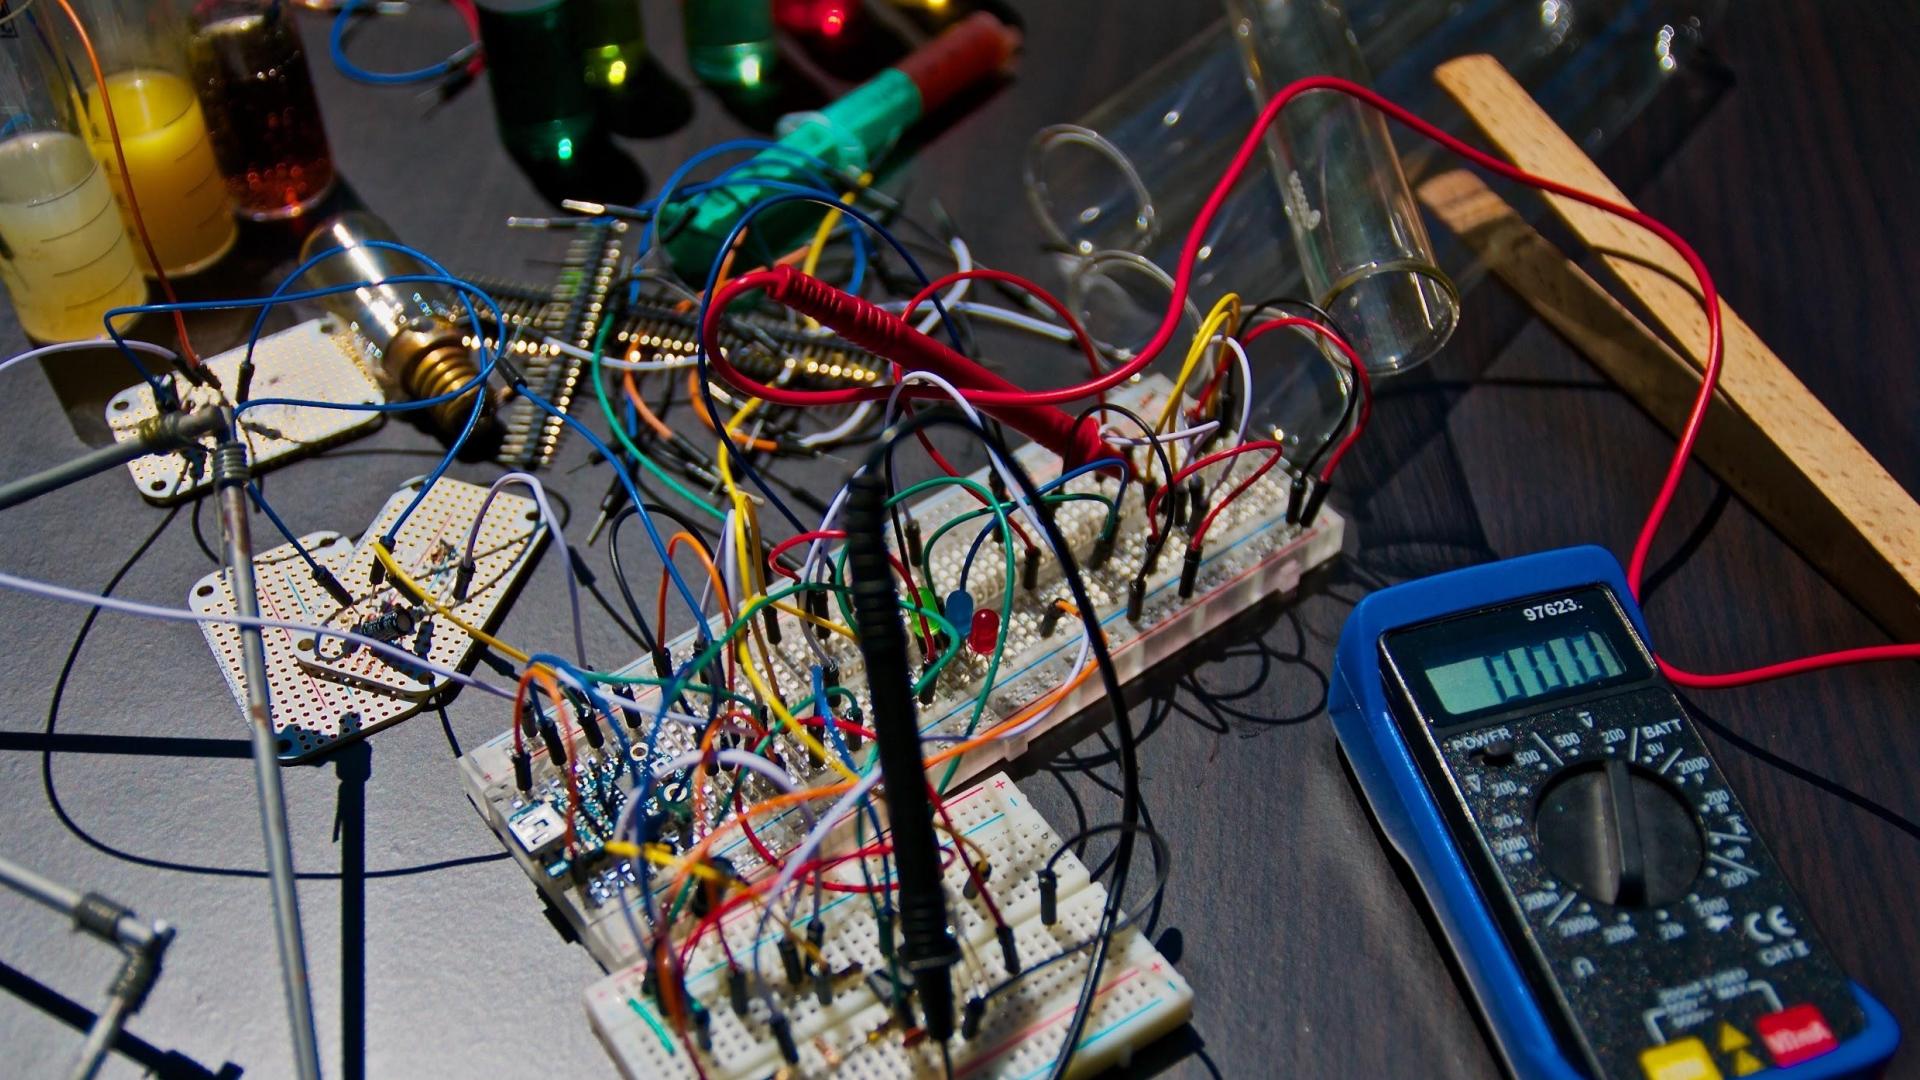


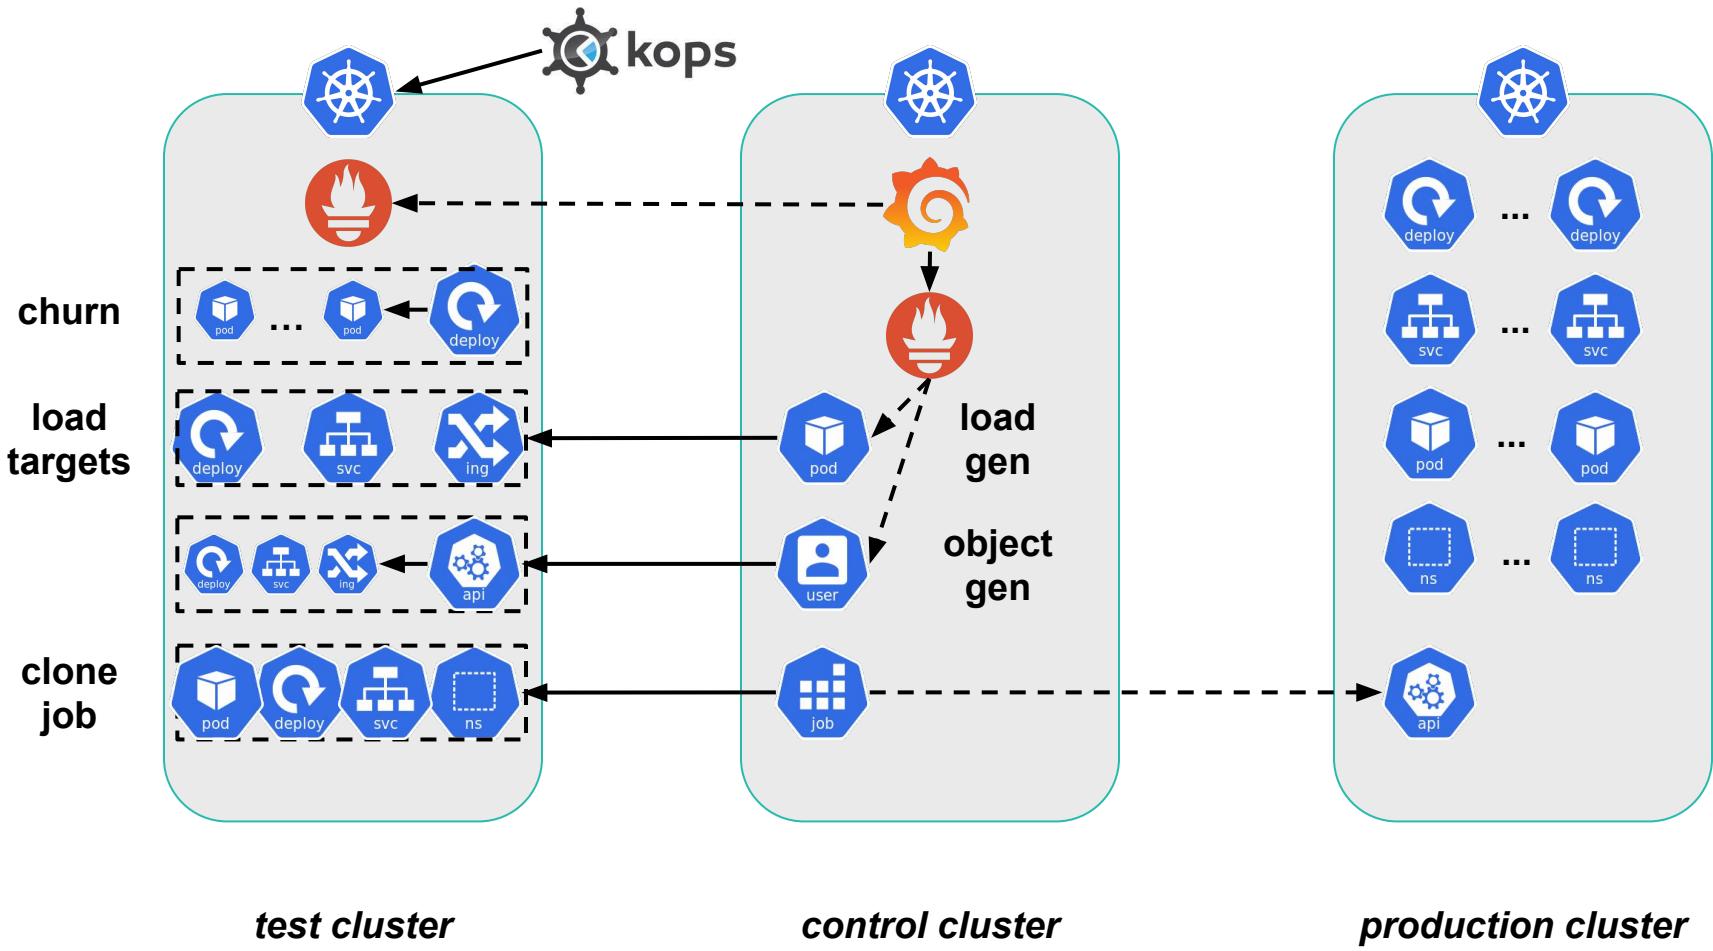
etc_server_proposals_committed per second

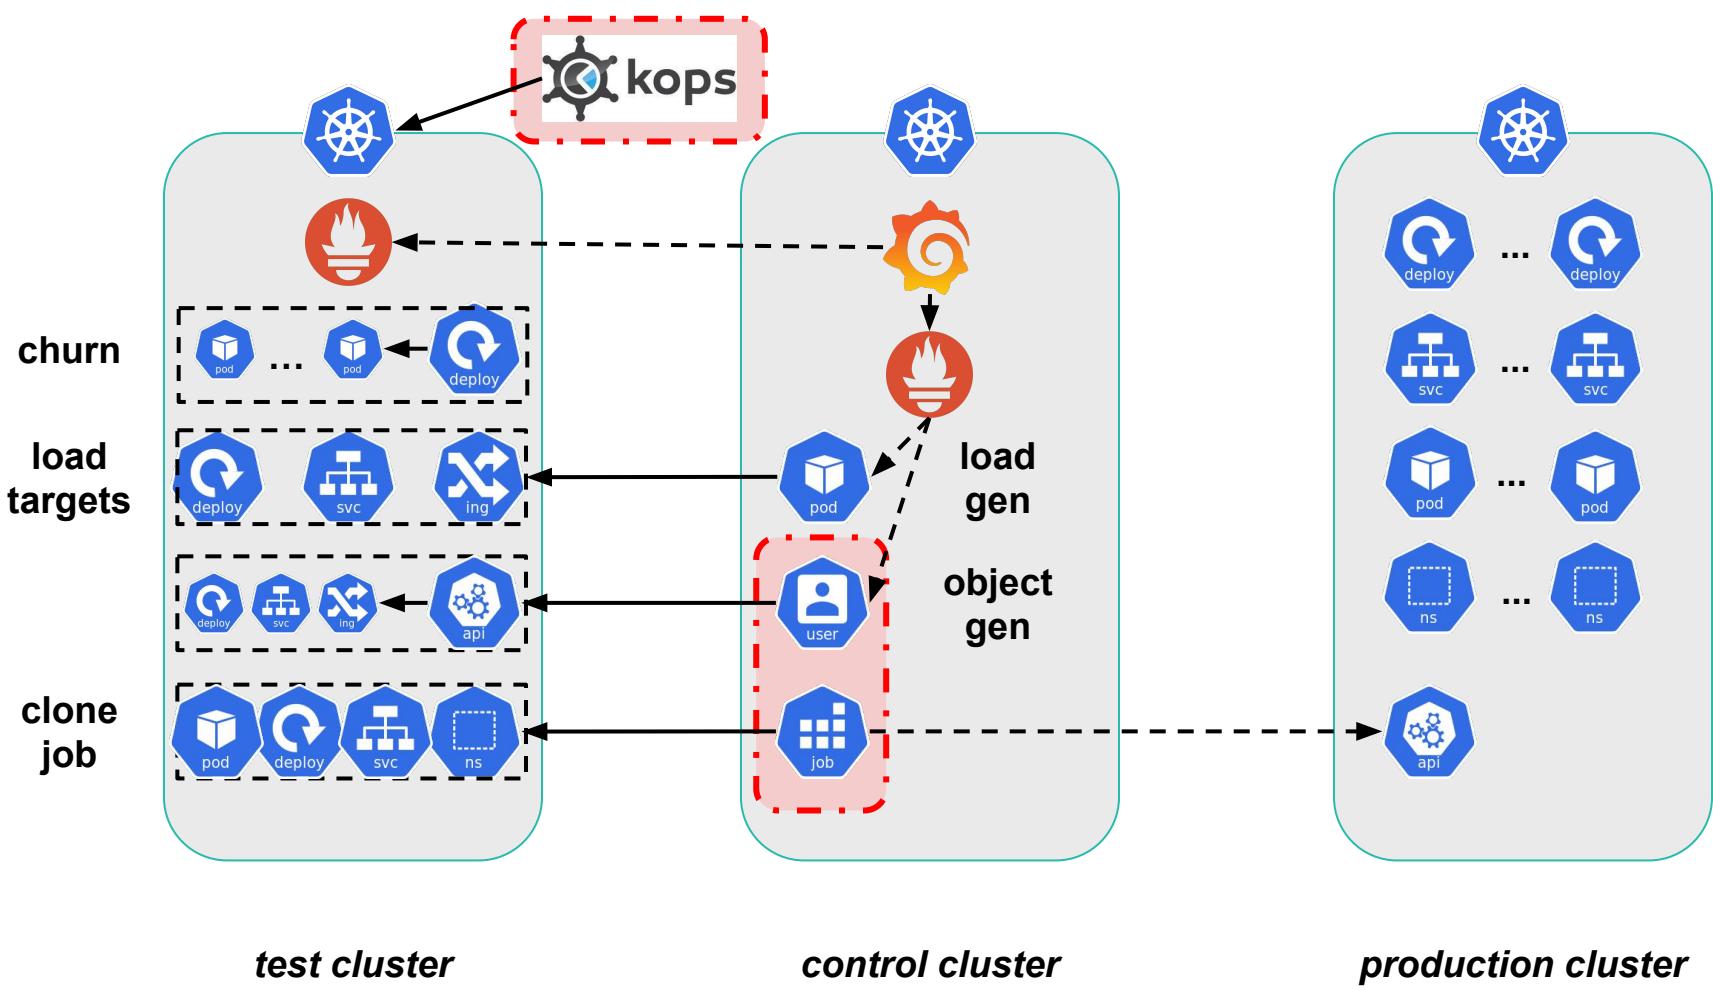


Quantile 0.99 etc_disk_wal_fsync_duration_seconds_bucket









Results and Learnings

Results and Learnings

Kubernetes itself is quite robust.

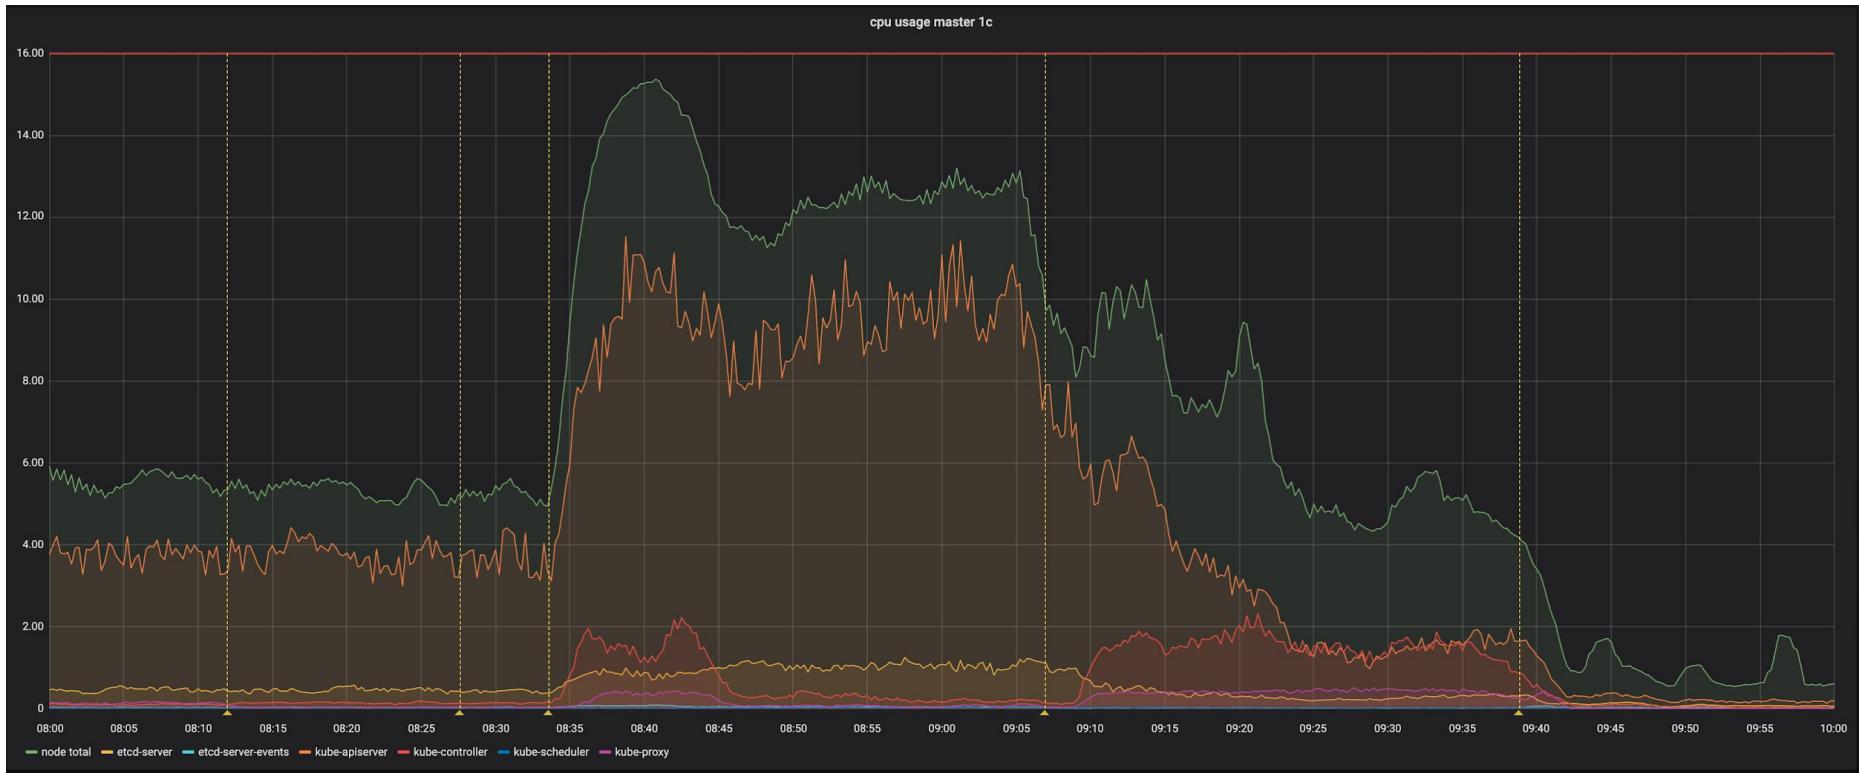
We could scale up to 4x the number of nodes and workloads of our current production cluster.

But...

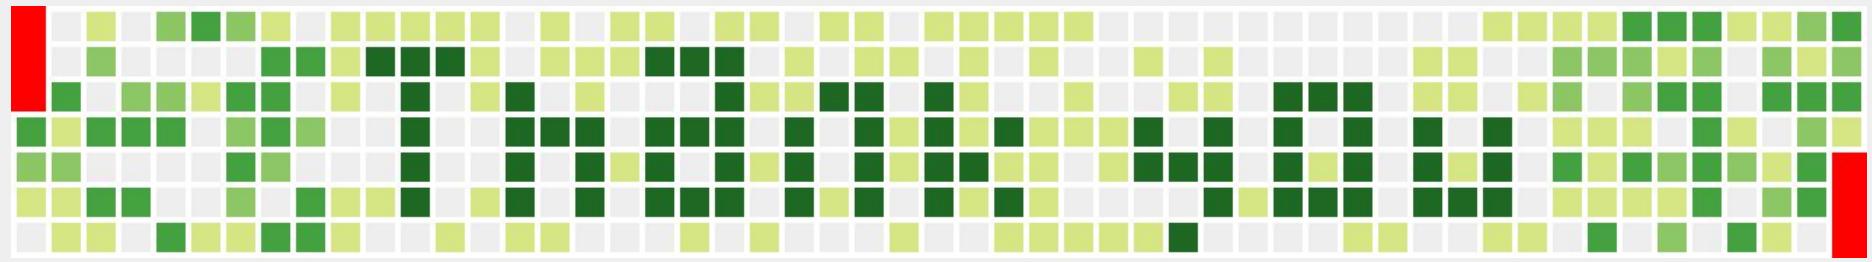
User experience starts to degenerate without performance tuning.

Some components for monitoring and add-on functionality needed resource tuning and different setup to cope with the large cluster.

Going from 500 to 900 nodes



<https://kubernetes.io/docs/setup/best-practices/cluster-large/>



Federico Hernandez
@recollar

Simone Sciarrati
@dezmodue