

Exploring Enhancements

Federico Hernandez

Principal Software Engineer
Meltwater



30000+ customers

1500+ employees

55 global offices

Founded 2001 in Oslo



9+ engineering offices (7+ countries)

330+ engineers in 50+ teams

Kubernetes in Production since 2018

2100 applications on Kubernetes
(50% production)

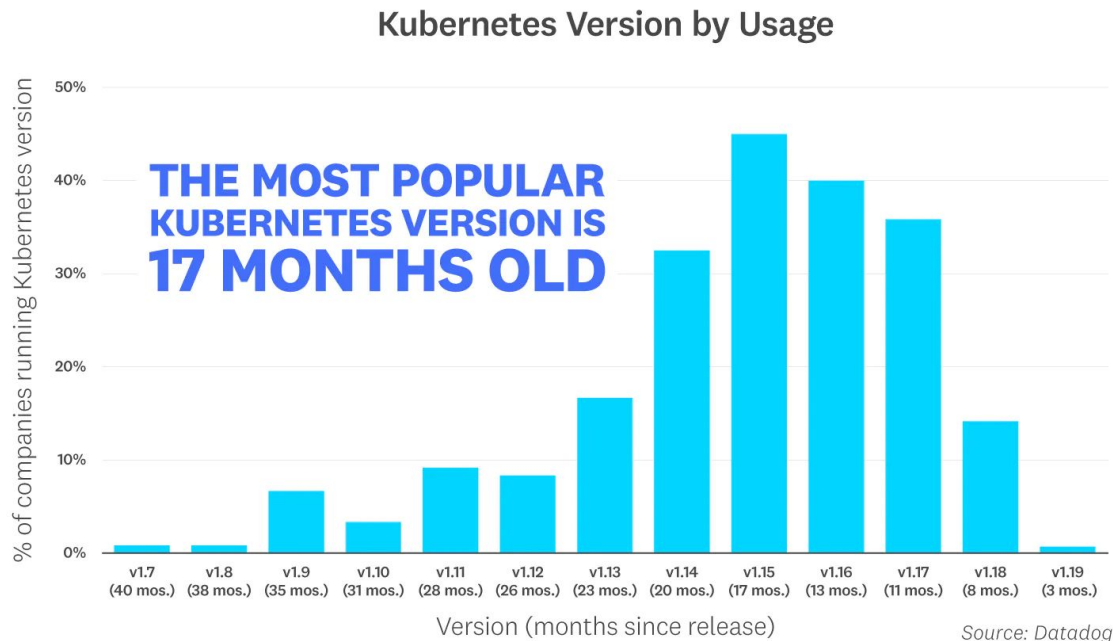
```
50 //fires the appear event when appropriate
51 var check = function() {
52     //is the element hidden?
53     if (!t.is(':visible')) {
54         //it became hidden
55         t.appeared = false;
56         return;
57     }
58
59     //is the element inside the visible window?
60     var a = w.scrollLeft();
61     var b = w.scrollTop();
62     var o = t.offset();
63     var x = o.left;
64     var y = o.top;
65
66     var ax = settings.acxX;
67     var ay = settings.acyY;
68     var th = t.height();
69     var wh = w.height();
70     var tw = t.width();
71     var ww = w.width();
72
73     if (y + th + ay >= b &&
74         y <= b + wh + ay &&
75         x + tw + ax >= a &&
76         x <= a + ww + ax) {
77         //trigger the custom event
78         if (!t.appeared) t.trigger('appear', settings.data);
79     } else {
80         //it scrolled out of view
81         t.appeared = false;
82     }
83 };
84
85 //create a modified fn with some additional logic
86 var modifiedFn = function() {
87     //mark the element as visible
88     t.appeared = true;
89
90     //is this supposed to happen only once?
91     if (settings.one) {
92         //remove the check
93         w.unbind('scroll', check);
94         var i = $.inArray(check, $.fn.appear.checks);
95         if (i >= 0) $.fn.appear.checks.splice(i, 1);
96     }
97
98     //trigger the original fn
99     fn.apply(this, arguments);
100
101     //bind the modified fn to the element
102     //if (settings.one) t.one('appear', settings.data, modifiedFn);
103     t.on('appear', settings.data, modifiedFn);
104 };
105
```

Survey

<https://sli.do>

#37090

Datadog: 11 facts about real-world container use



<https://www.datadoghq.com/container-report/>

Enhancements from 1.15 to 1.19

1.16 has 31 enhancements
(September 2019)

- 8 graduating to stable
- 8 move to beta
- 15 move to alpha

1.17 has 22 enhancements
(December 2019)

- 14 graduating to stable
- 4 move to beta
- 4 move to alpha

1.18 has 38 enhancements
(March 2020)

- 15 graduating to stable
- 11 move to beta
- 12 move to alpha

1.19 has 34 enhancements
(August 2020)

- 10 graduating to stable
- 15 move to beta
- 9 move to alpha

Plus 100s of other changes, improvements, bug fixes

<https://k8s.devstats.cncf.io/>

Enhancements from 1.15 to 1.19

1.16 has 31 enhancements
(September 2019)

- 8 graduating to stable
- 8 move to beta
- 15 move to alpha

1.17 has 22 enhancements
(December 2019)

- 14 graduating to stable
- 4 move to beta
- 4 move to alpha


1.18 has 38 enhancements
(March 2020)

- 15 graduating to stable
- 11 move to beta
- 12 move to alpha

1.19 has 34 enhancements
(August 2020)

- 10 graduating to stable
- 15 move to beta
- 9 move to alpha

Plus 100s of other changes, improvements, bug fixes



<https://k8s.devstats.cncf.io/>

Feature gates for Kubernetes

```
--feature-gates="...,EphemeralContainers=true,ServiceTopology=true,..."
```

Kops example:

```
kubeProxy:  
  metricsBindAddress: 0.0.0.0  
  featureGates:  
    EphemeralContainers: "true"  
    HPAScaleToZero: "true"  
    ImmutableEphemeralVolumes: "true"  
    ServiceTopology: "true"  
    EndpointSliceProxying: "true"
```

<https://kubernetes.io/docs/reference/command-line-tools-reference/feature-gates/>

Feature gates for Kubernetes

EndpointSlice	false	Alpha	1.16	1.16
EndpointSlice	false	Beta	1.17	
EndpointSlice	true	Beta	1.18	
EndpointSliceProxying	false	Alpha	1.18	1.18
EndpointSliceProxying	true	Beta	1.19	
PodPriority	false	Alpha	1.8	1.10
PodPriority	true	Beta	1.11	1.13
PodPriority	true	GA	1.14	-


<https://kubernetes.io/docs/reference/command-line-tools-reference/feature-gates/>

Exploring Enhancements

1. Ephemeral containers for kubectl debug
 2. Configurable pod process namespace sharing
 3. Even pod spreading across failure domains
 4. Topology aware routing of services
 5. A Startup Probe for slow starting pods
 6. Configurable scale velocity for HPA
 7. Changes in kubectl
-
- A. Graduation of Ingress to v1
 - B. Removed and Deprecated APIs

1. Ephemeral containers for kubectl debug

FEATURE STATE: v1.18 [alpha]




```
$ kubectl exec -ti foo -- sh
```

```
OCI runtime exec failed: exec failed:  
container_linux.go:349: starting container  
process caused "exec: \"sh\": executable  
file not found in $PATH": unknown command  
terminated with exit code 126
```

1. Ephemeral containers for kubectl debug

FEATURE STATE: v1.18 [alpha]




```
$ kubectl alpha debug  
-it pod-name  
--image=busybox  
--target=container-name
```

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-running-pod/#ephemeral-container>
<https://kubernetes.io/docs/concepts/workloads/pods/ephemeral-containers/>

2. Pod process namespace sharing

FEATURE STATE: v1.17 [stable]

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  shareProcessNamespace: true
  containers:
  - name: nginx
    image: nginx
  - name: shell
    image: busybox
```




```
$ kubectl attach -ti nginx -c shell
# ps aux
PID    USER     TIME   COMMAND
1      root      0:00   /pause
6      root      0:00   nginx: master process
33    101       0:00   nginx: worker process
34    root      0:00   sh
39    root      0:00   ps aux
```

<https://kubernetes.io/docs/tasks/configure-pod-container/share-process-namespace/>

3. PodTopologySpread

FEATURE STATE: v1.17 [alpha]


Affinity
Anti-Affinity
NodeSelector



3. PodTopologySpread

FEATURE STATE: v1.17 [alpha]

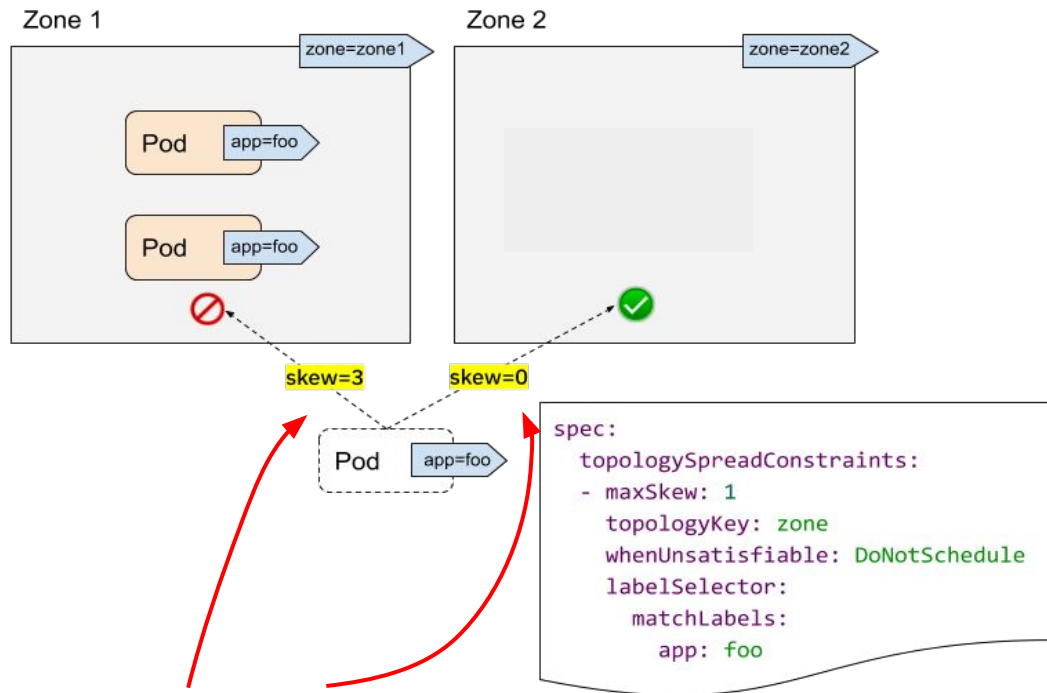
```
spec:
  topologySpreadConstraints:
  - maxSkew: 1
    topologyKey: zone
    whenUnsatisfiable: ScheduleAnyway
    labelSelector:
      matchLabels:
        app: bar
```



<https://kubernetes.io/blog/2020/05/introducing-podtopologyspread/>

<https://kubernetes.io/docs/concepts/workloads/pods/pod-topology-spread-constraints/>

3. PodTopologySpread



skew = Pods **number** matched in **current** topology - **min** Pods matches in a topology

<https://kubernetes.io/blog/2020/05/introducing-podtopologyspread/>


<https://kubernetes.io/docs/concepts/workloads/pods/pod-topology-spread-constraints/>

4. ServiceTopology

FEATURE STATE: v1.17 [alpha]

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
  topologyKeys:
    - "kubernetes.io/hostname"
    - "topology.kubernetes.io/zone"
    - "*"

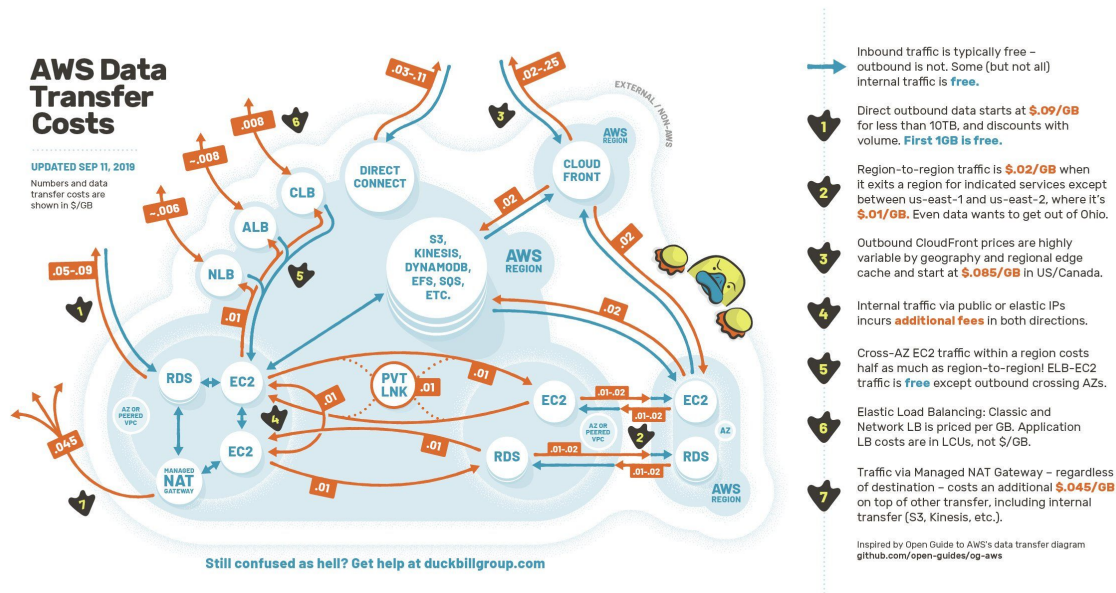
```



<https://kubernetes.io/docs/concepts/services-networking/service-topology/>

4. ServiceTopology


FEATURE STATE: v1.17 [alpha]




<https://kubernetes.io/docs/concepts/services-networking/service-topology/>

4. ServiceTopology

1 deployment + 1 HPA



3 deployment + 3 HPA




<https://kubernetes.io/docs/concepts/services-networking/service-topology/>

5. A Startup Probe for slow starting pods

FEATURE STATE: v1.16 [alpha], v1.18 [beta]

startupProbe:

```
httpGet:  
  path: /healthz  
  port: liveness-port  
failureThreshold: 30  
periodSeconds: 10
```



livenessProbe:

```
httpGet:  
  path: /healthz  
  port: liveness-port  
failureThreshold: 2  
periodSeconds: 10
```

readinessProbe:

```
httpGet:  
  path: /readyz  
  port: readiness-port  
failureThreshold: 1  
periodSeconds: 10
```

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/#define-startup-probes>

<https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/#container-probes>

5. A Startup Probe for slow starting pods

FEATURE STATE: v1.16 [alpha], v1.18 [beta]

↑	@@ -30,18 +30,23 @@ spec:		
30	httpGet:	30	httpGet:
31	path: /_status	31	path: /_status
32	port: 8080	32	port: 8080
33 -	initialDelaySeconds: 180	33 +	periodSeconds: 10
34 -	periodSeconds: 100		
35	failureThreshold: 3	34	failureThreshold: 3
36	successThreshold: 1	35	successThreshold: 1
37	livenessProbe:	36	livenessProbe:
38	httpGet:	37	httpGet:
39	path: /_status	38	path: /_status
40	port: 8080	39	port: 8080
41 -	initialDelaySeconds: 180	40 +	periodSeconds: 10
42 -	periodSeconds: 100		
43	failureThreshold: 3	41	failureThreshold: 3
44	successThreshold: 1	42	successThreshold: 1
		43 +	startupProbe:
		44 +	httpGet:
		45 +	path: /_status
		46 +	port: 8080
		47 +	periodSeconds: 5
		48 +	failureThreshold: 60
		49 +	successThreshold: 1
45	resources:	50	resources:
46	requests:	51	requests:
47	cpu: 1500m	52	cpu: 1500m
..... ↓			

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/#define-startup-probes>

<https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/#container-probes>

6. Configurable scale velocity for HPA

FEATURE STATE: v1.18 [v2beta2 of HPA]

```
minReplicas: 1
maxReplicas: 10
metrics:
- type: Resource
  resource:
    name: cpu
    target:
      type: Utilization
      averageUtilization: 50
```

```
desiredReplicas = currentReplicas * ( currentMetricValue / desiredMetricValue )
```

7. Configurable scale velocity for HPA

FEATURE STATE: v1.18 [v2beta2 of HPA]

```
behavior:
  scaleDown:
    stabilizationWindowSeconds: 300
  policies:
    - type: Pods
      value: 4
      periodSeconds: 60
    - type: Percent
      value: 10
      periodSeconds: 60
```

```
behavior:
  scaleUp:
    policies:
      - type: percent
        value: 200%
```

1 → 2 → 4 → 8

<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/#support-for-configurable-scaling-behavior>

7. Changes in kubectl

SINCE VERSION: v1.18

Remove kubectl run generators #87077

🔗 Merged

k8s-ci-robot merged 1 commit into `kubernetes:master` from `soltysch:remove_run_generators`

- **kubectl run** now only creates pods. Use **kubectl create** for other objects instead. But... :-/
- **kubectl diff** graduated to stable
- **kubectl get deployment bar -o yaml** contains now managed fields... :-/
<https://github.com/kubernetes/kubernetes/issues/90066>
- **kubectl get --export** query parameter removed but there is a plugin for it... :-/
- **kubectl apply --dry-run=client|server|none** is no longer boolean
- **kubectl logs --prefix** prefix each log line with pod and container name :-)
- **kubectl alpha debug** for ephemeral debug containers
-
-

Appendix A: Graduation of Ingress to v1

FEATURE STATE: v.18 [v1beta1], v1.19 [stable]

Since these features are new in Kubernetes 1.18, each Ingress controller implementation will need some time to develop support for these new features.

```
apiVersion: "networking.k8s.io/v1beta1"
kind: "Ingress"
metadata:
  name: "example-ingress"
spec:
  ingressClassName: "external-lb"
  rules:
  - host: "*.example.com"
    http:
      paths:
      - path: "/example"
        pathType: "Prefix"
      backend:
        serviceName: "example-service"
        servicePort: 80
```

<https://kubernetes.io/blog/2020/04/02/improvements-to-the-ingress-api-in-kubernetes-1.18/>
<https://v1-18.docs.kubernetes.io/docs/concepts/services-networking/ingress/>

Appendix B: Deprecated and Removed APIs

SINCE VERSION: v1.16

Deployment in the **extensions/v1beta1**, **apps/v1beta1**, and **apps/v1beta2** API versions is no longer served

- Migrate to use the **apps/v1** API version, available since v1.9. Existing persisted data can be retrieved/updated via the new version.
- Notable changes:
 - `spec.rollbackTo` is removed
 - `spec.selector` is now required and immutable after creation; use the existing template labels as the selector for seamless upgrades
 - `spec.progressDeadlineSeconds` now defaults to 600 seconds (the default in `extensions/v1beta1` was no deadline)
 - `spec.revisionHistoryLimit` now defaults to 10 (the default in `apps/v1beta1` was 2, the default in `extensions/v1beta1` was to retain all)
 - `maxSurge` and `maxUnavailable` now default to 25% (the default in `extensions/v1beta1` was 1)

The **v1.22** release will stop serving the following deprecated API versions in favor of newer and more stable API versions:

- Ingress in the **extensions/v1beta1** API version will no longer be served
 - Migrate to use the **networking.k8s.io/v1beta1** API version, available since v1.14. Existing persisted data can be retrieved/updated via the new version.

<https://kubernetes.io/blog/2019/07/18/api-deprecations-in-1-16/>

Appendix B: Deprecated and Removed APIs

Don't Panic: Kubernetes and Docker

Wednesday, December 02, 2020

Authors: Jorge Castro, Duffie Cooley, Kat Cosgrove, Justin Garrison, Noah Kantrowitz, Bob Killen, Rey Lejano, Dan "POP" Papandrea, Jeffrey Sica, Davanum "Dims" Srinivas

Kubernetes is [deprecating Docker](#) as a container runtime after v1.20.

You do not need to panic. It's not as dramatic as it sounds.

tl;dr Docker as an underlying runtime is being deprecated in favor of [ContainerD](#) and the [CRI](#) Interface(CRI) created for Kubernetes. Docker-produced images will continue to be supported by Kubernetes, as they always have.

The **v1.22** release will stop serving the following

- Ingress in the **extensions/v1beta1** API version
 - Migrate to use the **networking.k8s.io/v1** API version. Data can be retrieved/updated via the new version.

is no longer served

data can be retrieved/updated via the new

using template labels as the selector for

Can I still use Docker in Kubernetes 1.20?

Yes, the only thing changing in 1.20 is a single warning log printed at [kubelet](#) startup if using Docker as the runtime.

When will dockershim be removed?

Given the impact of this change, we are using an extended deprecation timeline. It will not be removed before Kubernetes 1.22, meaning the earliest release without dockershim would be 1.23 in late 2021. We will be working closely with vendors and other ecosystem groups to ensure a smooth transition and will evaluate things as the situation evolves.

<https://kubernetes.io/blog/2019/07/18/api-deprecations-in-1-16/>

Further information

Kubernetes release announcements:

- 1.16: <https://kubernetes.io/blog/2019/09/18/kubernetes-1-16-release-announcement/>
- 1.17: <https://kubernetes.io/blog/2019/12/09/kubernetes-1-17-release-announcement/>
- 1.18: <https://kubernetes.io/blog/2020/03/25/kubernetes-1-18-release-announcement/>
- 1.19: <https://kubernetes.io/blog/2020/08/26/kubernetes-release-1.19-accentuate-the-paw-sitive/>




Sysdig - What's new in Kubernetes:

- 1.16: <https://sysdig.com/blog/whats-new-kubernetes-1-16>
- 1.17: <https://sysdig.com/blog/whats-new-kubernetes-1-17>
- 1.18: <https://sysdig.com/blog/whats-new-kubernetes-1-18>
- 1.19: <https://sysdig.com/blog/whats-new-kubernetes-1-19>




(as well as the *lengthy* CHANGELOG like

<https://github.com/kubernetes/kubernetes/blob/master/CHANGELOG/CHANGELOG-1.19.md>)



Federico Hernandez
@recollir

We are looking for a **Junior Platform Engineer** (GOT)



<https://emp.jobylon.com/jobs/69272>